



# Modernizing your Database with Azure SQL

Azure SQL Team, Microsoft





# Please silence cell phones



# Explore everything PASS has to offer

Free Online Resources  
Newsletters  
[PASS.org](http://PASS.org)



24HOURS  
OF PASS

Free online  
webinar events



PASS  
LOCAL  
GROUPS



PASS  
SQLSATURDAY

Local user groups  
around the world

Free 1-day local  
training events



PASS  
MARATHON

Online special  
interest user groups



PASS  
VIRTUAL  
GROUPS



PASS  
VOLUNTEERS

Business analytics  
training

Get involved



# Azure SQL Team

Microsoft



<https://azure.microsoft.com/en-us/services/sql-database/>



@azuresql #AzureSQL

Sanjay Mishra  
Denzil Ribeiro  
Shreya Verma  
Mladen Andzic  
Andreas Wolter  
Davide Mauri  
Rohit Nayak  
Dimitri Furman  
Jovan Popovic  
Joe Sack  
Rie Irish



# Getting Started



# Agenda

Time	Topic	Speaker
08:40 - 08:50	Azure SQL: What, Why and How	Sanjay Mishra
08:50 - 09:30	Azure SQL VM: The Simplest way to run SQL in Azure	Shreya Verma
09:30 - 10:00	Azure SQL PaaS services: maximize your potential. What PaaS provides and how?	Sanjay Mishra
10:00 - 10:15	Refreshment Break	
10:15 - 10:45	Azure SQL MI: What and How? Migration to MI.	Mladen Andzic

# Agenda

Time	Topic	Speaker
10:45 - 11:15	Networking in Azure SQL DB and MI	Rohit Nayak
11:15 - 12:00	Security in Azure SQL DB and MI	Andreas Wolter
12:00 - 1:00	Lunch Break	
01:00 - 01:30	Azure SQL DB: Best cloud database for developers	Davide Mauri
01:30 - 02:00	Azure SQL DB: Best cloud database for Enterprise applications	Dimitri Furman, Davide Mauri

# Agenda

Time	Topic	Speaker
02:00 - 02:30	Azure SQL DB Hyperscale	Denzil Ribeiro
02:30 - 02:45	Refreshment Break	
02:45 - 03:25	Azure SQL DB: Best cloud database for DBAs	Dimitri Furman
03:25 - 03:55	Troubleshooting Common Performance and Scale problems	Denzil Ribeiro
03:55 - 04:25	Intelligent DB	Joe Sack

# Session Evaluations

Submit by 5pm Friday,  
November 15th to  
win prizes.

## 3 WAYS TO ACCESS



Go to [PASSsummit.com](http://PASSsummit.com)



Download the GuideBook App  
and search: PASS Summit 2019



Follow the QR code link on session  
signage



# Azure SQL

What, Why and How

Sanjay Mishra

Student of Life, for Life



# Making the most of it

- Learning Continues ...
- Sign up sheet
- Take a page out of your notebook

# Azure SQL

## SQL virtual machines

Best for migrations and applications requiring OS-level access



## Managed instances

Best for most lift-and-shift migrations to the cloud



## Databases

Best for modern cloud applications. Hyperscale and serverless options are available



### SQL virtual machine

- SQL Server and OS server access
- Expansive SQL And OS version support
- Automated manageability features for SQL Server

### Single instance

- SQL Server surface area (vast majority)
- Native virtual network support
- Fully managed service

### Instance pool

- Resource sharing between multiple instances to price optimize
- Enables migration of many small instances at scale
- Fully managed service

### Single database

- Hyperscale storage (up to 100TB)
- Serverless compute
- Fully managed service

### Elastic pool

- Resource sharing between multiple databases to price optimize
- Simplified performance management for multiple databases
- Fully managed service



# SQL Server in Azure Virtual Machine

Shreya Verma  
Program Manager, Microsoft



# SQL Server on Azure Virtual Machines

## Secure and compliant



Most compliant cloud



## Seamless cloud migration



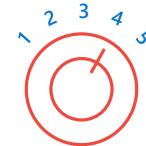
100% SQL Server compatibility and only cloud with pre-configured Developer edition

## Leading TCO



Save up to 43% vs.  
AWS EC2 with  
Azure Hybrid  
Benefit

## Flexibility and control with automation



Easier to maintain than EC2 with automated security patches and automated backup

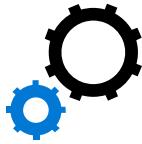


# SQL VM Resource Provider



# What is Resource Provider?

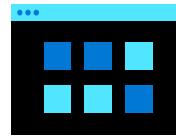
Resource Provider brings the functionality of Azure Marketplace images to SQL Server instances self-installed on Azure VMs.



## Comprehensive feature set

Self-installed VMs registered with RP now can access automation features in Azure Marketplace images

Leverage auto-backup, auto-patching, and automate Always On Availability Groups to avoid time-consuming admin and VM customization



## Dashboard view for VM awareness

Azure VMs are now discoverable on the new Azure SQL blade in Azure Marketplace



Easily manage your SQL VM and SQL PaaS deployments from one central location



## Simple license conversions

Self-installed VMs with RP can be easily converted to PAYG images



Save money by converting variable workloads with Software Assurance to PAYG images

# How to register with SQL VM RP?

- Register the resource provider to your subscription (one-time only) via the Azure Portal or Azure CLI

```
# Register the new SQL resource provider to your subscription  
az provider register --namespace Microsoft.SqlVirtualMachine
```

- Register single VM with SQL VM RP

```
# Register your existing SQL Server VM with the new resource provider  
az sql vm create -n <VMName> -g <ResourceGroupName> -l <VMLocation> --license-type <PAYG/AHUB>
```

- NEW Bulk register with SQL VM RP

```
# Register all existing SQL Server VM with the new resource provider  
Import-Module .\RegisterSqlVMs.psm1 Register-SqlVMs -SubscriptionList SubscriptionId1, SubscriptionId2
```

\*If you're unable to select Full or Lightweight RP, choose an Agentless approach. Note this option does not enable verification of Resource Provider.

DEMO

# Azure SQL VM – Resource Provider





# SQL Server Performance on Azure VM





# Performance Guidance for SQL Server on Azure VM

It is still just SQL Server....

Locked Pages in Memory (Windows)

Instant File Initialization (Windows)

Standard SQL Query Performance Best Practices and Tuning

SQL Linux Performance [Best Practices](#)

RHEL8 has XFS optimizations

Don't forget about File-Snapshot Backups in Azure

## IOPS Testing

- 1- Start with collecting peak IOPS requirement for Data, Log and Temp DB.
- 2- Choose the VM size that can scale to the total IOPS requirement
- 3- If workload is Temp DB heavy, locate Temp DB on Local SSD
- 4- Use RO cache enabled Premium SSD disks <4TB for Data Files
- 5- Use disk striping to achieve Data IOPS requirement
- 6- Provision Ultra Disk for Log file
- 7- Monitor SQL DB performance, write queues on Data Files becoming the bottleneck for throughput then move Data files to Ultra SSD.
- 8- <https://techcommunity.microsoft.com/t5/SQL-Server/Optimize-OLTP-Performance-with-SQL-Server-on-Azure-VM/ba-p/916794>

# SQL Server IaaS on Windows

GIGAOM

## SQL Server 2017 Enterprise (Windows) on Azure and AWS

- SQL Server on Azure includes Azure Hybrid Benefit
- SQL Server on AWS includes license mobility

### Performance

Transactions per second (TPS)

3.4x Faster

1088.76

aws

320.11

### Price/performance\*

3-year pricing / TPS

Azure

\$59.02

87% Less

AWS

\$445.65

\*Includes Azure Hybrid Benefit and Reserved Instances with three-year commitment.

\*\*Price-performance claims based on data from a study commissioned by Microsoft and conducted by GigaOm in October 2019. The study compared price performance between SQL Server 2017 Enterprise edition on Windows Server 2016 Datacenter edition in Azure E64s\_v3 instance type with 4x P30 1TB Storage Pool data (Read Only Cache) + 1x P20 0.5TB log (No Cache) and the SQL Server 2017 Enterprise edition on Windows Server 2016 Datacenter edition in AWS EC2 r4.16xlarge instance type with 1x 4TB gp2 data + 1x 1TB gp2 log. Benchmark data is taken from a GigaOm Analytic Field Test derived from a recognized industry standard, TPC Benchmark™ E (TPC-E). The Field Test does not implement the full TPC-E benchmark and as such is not comparable to any published TPC-E benchmarks. The Field Test is based on a mixture of read-only and update intensive transactions that simulate activities found in complex OLTP application environments. Price-performance is calculated by GigaOm as the cost of running the cloud platform continuously for three years divided by transactions per second throughput. Prices are based on publicly available US pricing in West US for SQL Server on Azure Virtual Machines and Northern California for AWS EC2 as of October 2019. Pricing incorporates three-year reservations for Azure and AWS compute pricing, and Azure Hybrid Benefit for SQL Server and Azure Hybrid Benefit for Windows Server and License Mobility for SQL Server in AWS, excluding Software Assurance costs. Price-performance results are based upon the configurations detailed in the GigaOm Analytic Field Test. Actual results and prices may vary based on configuration and region.





# SQL Server High Availability on Azure VM



# SQL Server HADR technologies for Azure IaaS



[Always On Availability Groups](#)



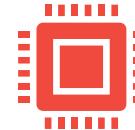
[Failover Cluster Instances](#)



[Log Shipping](#)



[SQL Server Backup and Restore with Azure Blob Storage Service](#)



[Database Mirroring](#) -  
Deprecated  
in SQL  
Server 2016

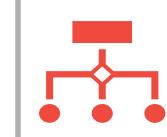
# NEW SQL Server FCI with Azure Premium File Share



Ease of management



File shares are fully managed by Azure.



Lower the work on your virtual

**machines**  
Input or output is offloaded to your managed file share



Burstable input or output capacity



Zonal Redundancy

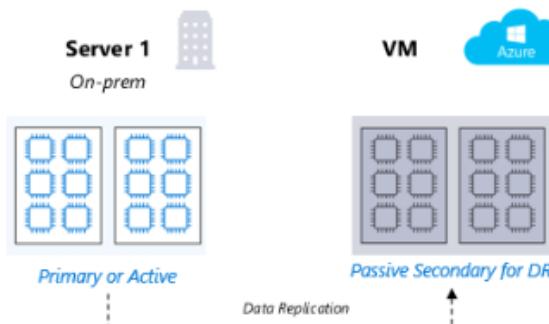
Premium scale factor	
IOPS	1 * provisioned GiB. (max 100,000 IOPS).
Burst IOPS	3 * Baseline IOPS. (max 100,000 IOPS)
Egress	60 MiB/s + 0.06 * provisioned GiB
Ingress	40 MiB/s + 0.04 * provisioned GiB

Example 10TB Share	
IOPS	10,000
Burst IOPS	30,000
Egress	660 MiB/s
Ingress	440 MiB/s

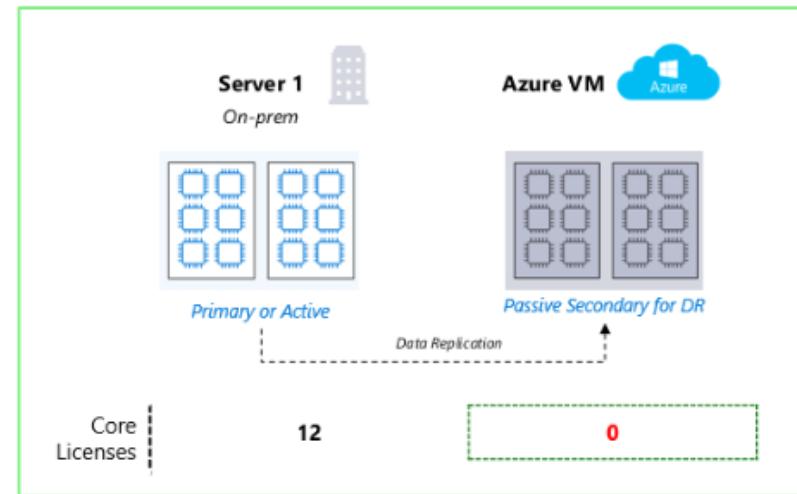
[Step-by-step tutorial to configure SQL Server FCI with Premium File Shares](#)

# **NEW** Free Failover servers for disaster recovery in Azure

## Under Old Licensing Rules



## With New SA Benefit





# Thank You

Shreya Verma



@shreyavermakale



# Can you please PASS the PaaS !

Azure SQL PaaS Services

Sanjay Mishra ([sanjaymi@microsoft.com](mailto:sanjaymi@microsoft.com))  
Student of Life, for Life

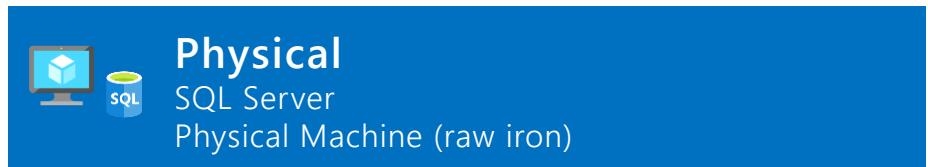
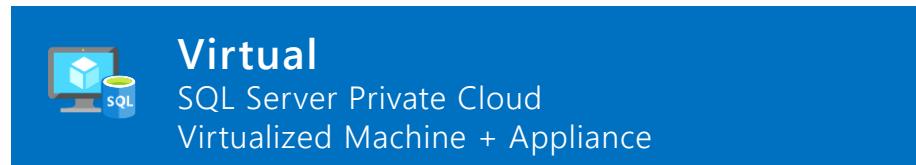
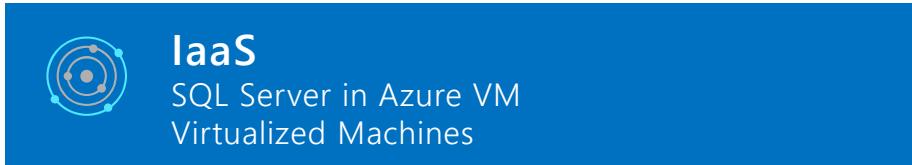


# Data platform continuum

Shared. Lower cost



Dedicated. Higher cost



Higher administration

Lower administration

# Opportunity for modernization

Managed by customer

Managed by Microsoft

AI

**On-premises costs** tend to be driven by hardware and data center management costs

**Infrastructure-as-a-Service** reduces cost categories related to data center and compute

**Platform-as-a-Service** off-loads customers' most administrative tasks to Azure, further improving efficiency with machine-learning capabilities for performance and security

- **Managed Instance:** instance-level deployment for lift-shift existing apps to Azure, fully backward compatible
- **Single database:** database-level deployment for new apps

## On-premises

Applications
Data
High availability /DR/Backups
Database Provision/ Patch/Scaling
O/S provision /patching
Virtualization
Hardware
Datacenter Management

SQL Server 2019

## Infrastructure (as a Service)

Applications
Data
High availability /DR/Backups
Database Provision/ Patch/Scaling
O/S
Virtualization
Hardware
Datacenter Management

Azure SQL VMs

## Platform (as a Service)

Intelligent performance/security
Applications
Data
High Availability/ DR/Backups
Database Provision/ Patch/Scaling
O/S
Virtualization
Hardware
Datacenter Management

Azure SQL Database

# Azure SQL

## SQL virtual machines

Best for migrations and applications requiring OS-level access



### SQL virtual machine

- SQL Server and OS server access
- Expansive SQL and OS version support
- Automated manageability features for SQL Server

## Managed instances

Best for most lift-and-shift migrations to the cloud



### Single instance

- SQL Server surface area (vast majority)
- Native virtual network support
- Fully managed service

### Instance pool

- Resource sharing between multiple instances to price optimize
- Enables migration of many small instances at scale
- Fully managed service

## Databases

Best for modern cloud applications. Hyperscale and serverless options are available



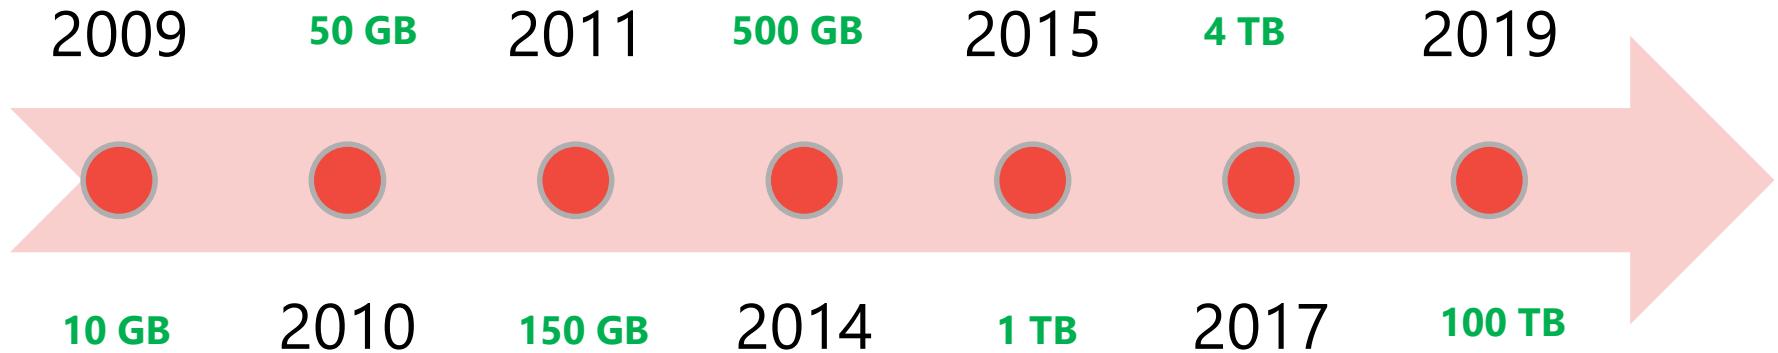
### Single database

- Hyperscale storage (up to 100TB)
- Serverless compute
- Fully managed service

### Elastic pool

- Resource sharing between multiple databases to price optimize
- Simplified performance management for multiple databases
- Fully managed service

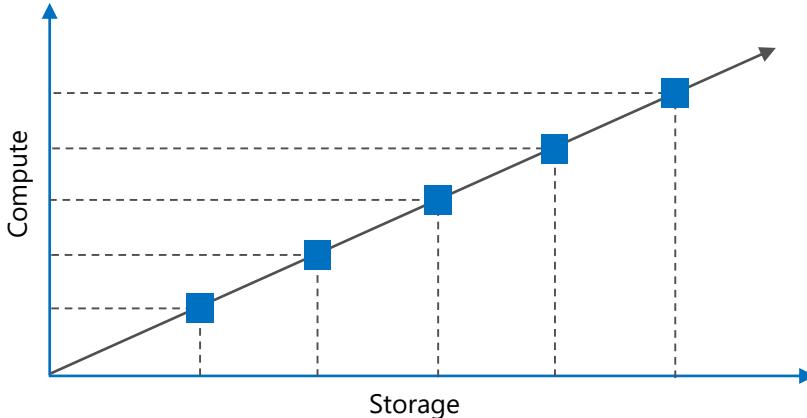
# Azure SQL DB Journey



# Flexible compute & storage options

## DTU model

Simple, preconfigured



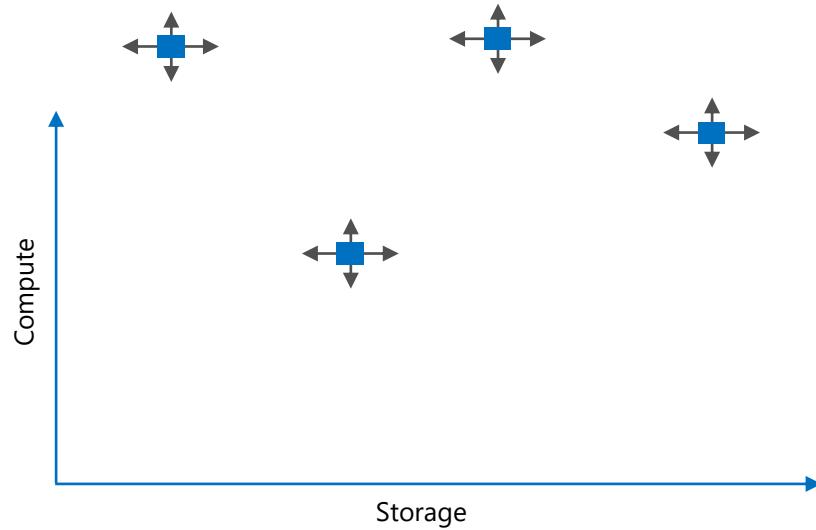
Pre-packaged, bundled unit that represents the database power

Designed for predictable performance, but somewhat inflexible and limited in options

DTU sizing offers simplicity of choice

## vCore model

Independent scalability



This model allows you to independently choose compute and storage resources. It also allows you to use Azure Hybrid Benefit for SQL Server to gain cost savings.

Best for customers who value flexibility, control and transparency

# Flexible compute, storage & performance options

Service tier	 General purpose	 Business critical	 Hyperscale		
<b>Best for</b>	Offers budget oriented balanced compute and storage options	OLTP applications with high transaction rate and low IO latency.	Most business workloads. Auto-scaling storage size up to 100 TB, fluid vertical and horizontal compute scaling, fast database restore		
<b>Deployment option</b>	Single / Elastic Pools	Managed Instance	Single / Elastic Pools	Managed Instance	Single
<b>Compute tiers</b>	Gen4: 1 to 24 vCore Gen5: 2 to 80 vCore	Gen4: 4 to 24 vCore Gen5: 2 to 80 vCore	Gen4: 1 to 24 vCore Gen5: 2 to 80 vCore	Gen4: 4 to 24 vCore Gen5: 4 to 80 vCore	Gen4: 1 to 24 vCore Gen5: 2 to 80 vCore
<b>Storage</b>	<b>Premium remote</b>		<b>Local SSD</b>		<b>Local SSD Cache</b>
	32GB – 4TB per instance	32GB – 8TB per instance	32GB – 4TB per instance	32GB – 4TB per instance	Scale from 5GB to 100TB of storage in 10GB increments

# Azure SQL Database — Everything built-in

## Intelligent performance



Realize automatic performance improvements from continuous assessment and innovation

## Scales on the fly



Change service tiers, performance levels, and storage dynamically without downtime

## Business continuity



Easily manage and monitor business critical functions for reliable operations

## Works in your environment



Develop your app and connect to SQL Database with the tools and platforms you prefer

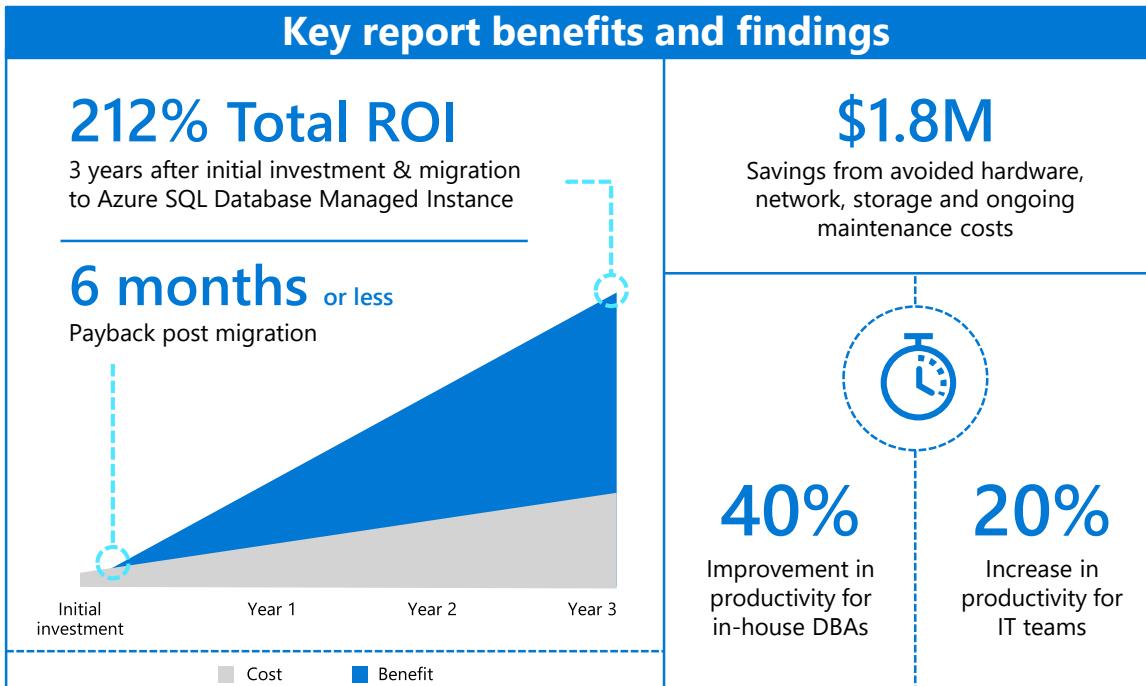
## Advanced threat protection



Build security-enhanced apps with built-in protection and industry-leading compliance

# The Total Economic Impact of Azure SQL Database Managed Instance

Microsoft commissioned Forrester Consulting to conduct a Total Economic Impact™ study to examine potential cost savings and business benefits enterprises would achieve from migrating on-premises workloads to Azure SQL Database Managed Instance.



“ Azure SQL Database Managed Instance is part of our strategic mandate to move all of our application, database, and services footprint to the Azure cloud. We can quickly integrate and are more nimble and more efficient as a result.

*Head of development,  
technology company*

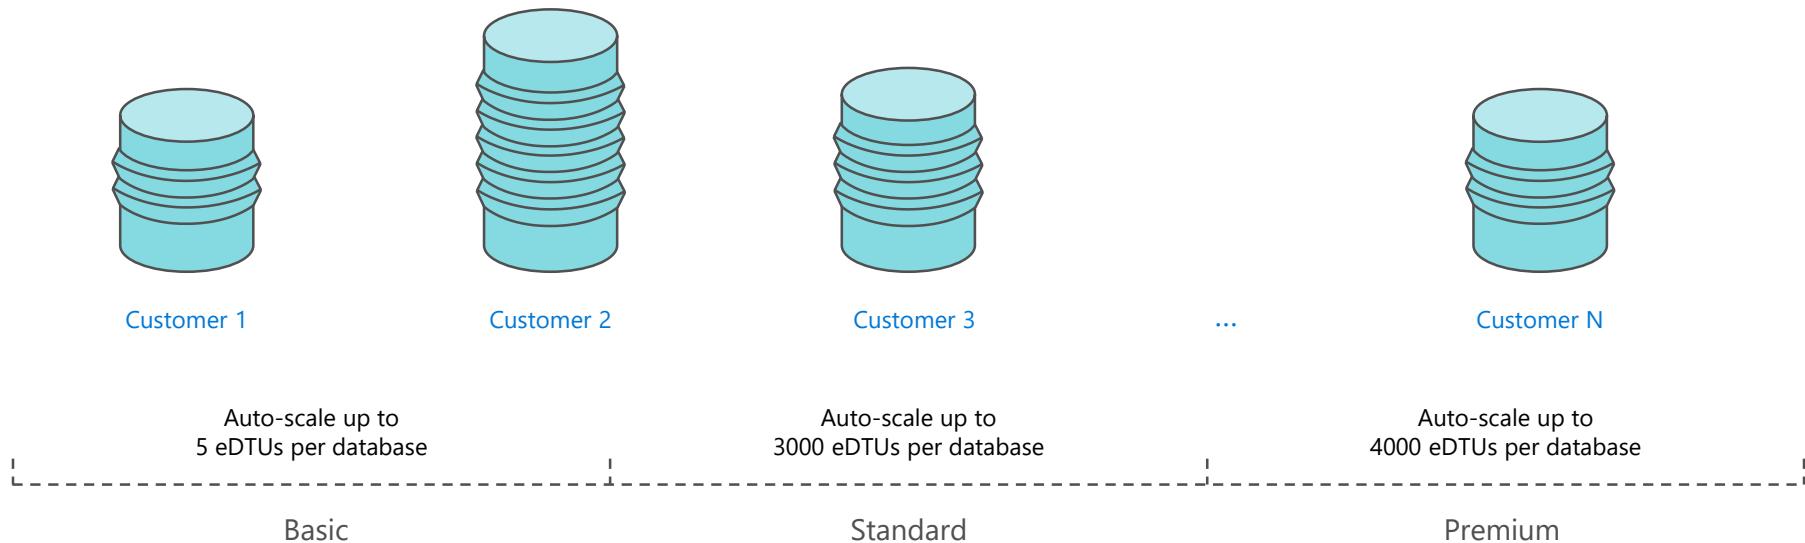
Download the full [Total Economic Impact™ of Azure SQL Database Managed Instance report](#)



# Elastic database pool service tiers

Buy a fixed number of eDTUs or vCores, share compute across many databases

## ELASTIC DATABASE POOLS



# Elastic database jobs overview

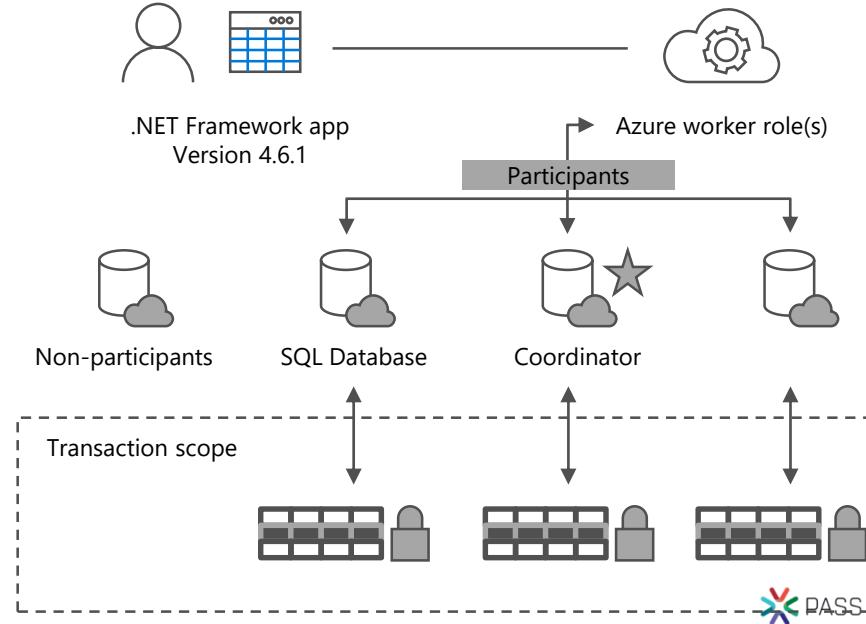
- Elastic database jobs enables you to run T-SQL scripts (jobs) against a set of databases
- Execute T-SQL scripts reliably with automatic retry and at scale
- Track job execution state



# Elastic database transactions

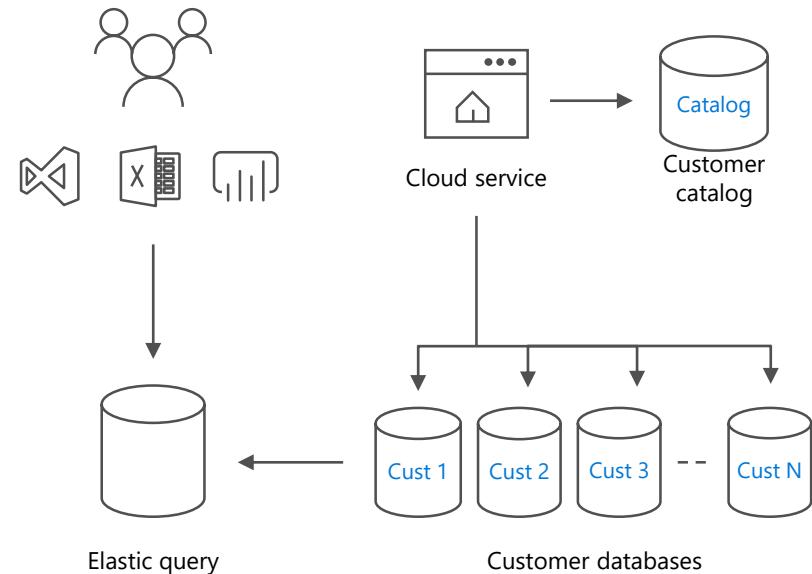
- Distributed transactions: perform operations across several databases with transactional properties in Azure SQL Database
- Ensure that changes are made to all databases or not at all
- Use the same .NET APIs that are used on-premises today

## Distributed transactions on Azure SQL Database



# Querying across many databases

- Connect to a single Azure SQL Database instance using familiar Azure SQL Database connection strings
- Simplify setup with Transact-SQL Data Definition Language (DDL)
- Transparently query many databases from a single database
- Familiar programming experience with Transact-SQL, ADO .NET LINQ, Entity Framework (EF), and others
- Use familiar development and business-intelligence (BI) tools to query across many databases
- Designed for interactive querying



# Serverless supports common industry scenarios of sporadic or unpredictable usage



## Line of business apps

Expense reporting and employee tracking apps  
Procurement systems



## E-commerce

Opening new marketplaces, marketing campaigns, sales promotions



## Content management systems

Updating and publishing web content  
Content clearinghouses that pull select content by third parties



## Dev/test workloads

Handling unpredictable workload needs

# Azure SQL Serverless

Copy    Restore    Export    Set server firewall    Delete    Connect with...    Feedback

Resource group (change) : [REDACTED]  
Status : Paused  
Location : West US  
Subscription (change) : [REDACTED]  
Subscription ID : [REDACTED]  
Tags (change) : Click here to add tags

Server name : [REDACTED]  
Connection strings : Show database connection strings  
Pricing tier : General Purpose: Serverless, Gen5, 16 vCores  
Auto-pause delay : 1 hour  
Earliest restore point : 2019-10-30 23:51 UTC

Show data for last: 1 hour   24 hours   7 days

Compute utilization

View: Max

Database not in use

Oct 30 8:37 PM   Oct 31   6 AM   12 PM   1:30 AM

CPU percentage (Max)  
smsserver/smsldb  
0 %

Data IO percentage (...)  
smsserver/smsldb  
0 %

Log IO percentage (Max)  
smsserver/smsldb  
0 %

App CPU billed

View: Sum

Database not billed

Oct 30 8:30 PM   Oct 31   6 AM   12 PM

App CPU billed (Sum)  
smsserver/smsldb  
---

# Provisioned compute and serverless meet different needs

Optimize compute provisioning and billing for your workload

## Serverless databases...

Scale up or down to meet workload requirements, instead of pre-provisioning

Bill on a per-second basis

### Common scenarios

Workloads with unpredictable and intermittent usage patterns or performance requirements

Workloads where the requirements are unknown and you can delegate compute sizing to the service



## Databases with provisioned compute...

Provision compute resources upfront

Bill on an hourly basis

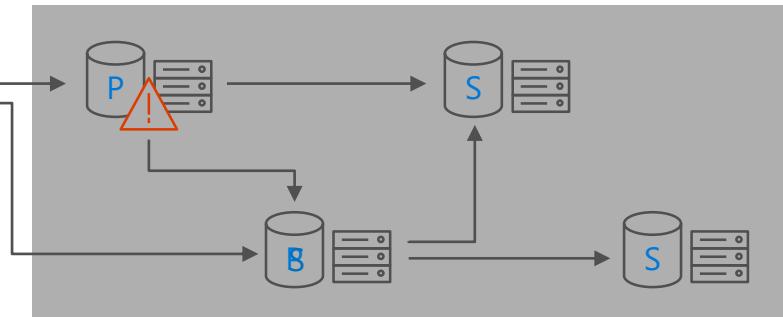
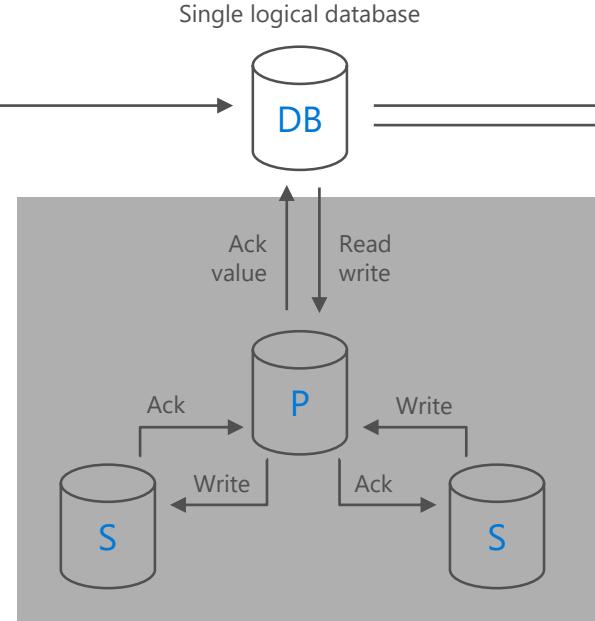
### Common scenarios

Workloads with regular and substantial compute utilization

Multiple databases with bursty usage patterns that can be consolidated into a single server and use *elastic pools* for better price optimization



# High-availability platform

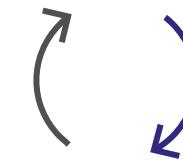


Recovery from machine failure  
Critical capabilities:

- ✓ Create new replica
- ✓ Synchronize data
- ✓ Stay consistent
- ✓ Detect failures
- ✓ Failover
- ✓ 99.995% SLA

# BCDR: Roles and responsibilities

- Azure SQL Database
  - Geo-distributed service
  - Customer metadata protection and recovery
  - Transparent high availability and data protection from local platform failures
  - Automatic geo-distributed backups
  - Automatic data synchronization of geo-replicated databases
  - Platform compliance testing and certification
  - Alert to impacted customers during regional failures
- Customer
  - Implementing retry logic in your application
  - Detecting user errors and initiating point-in-time restore
  - Planning, database prioritization, and region selection for disaster recovery
  - Initiating geo-restore to selected region
  - Application disaster recovery drills



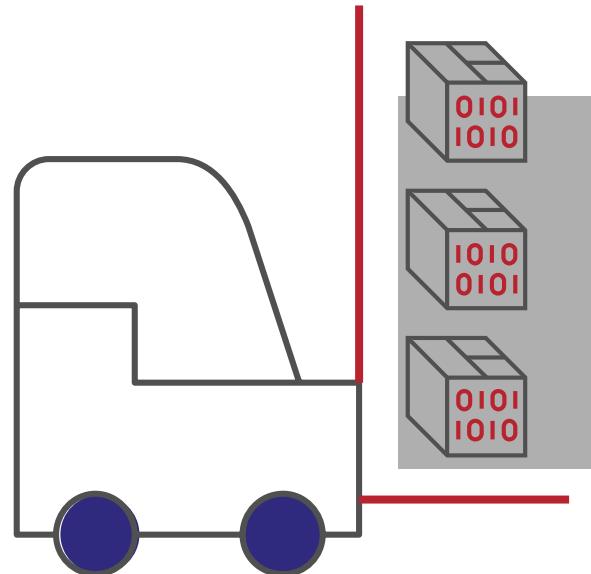
# Long-term data retention

Automatically created with LTR capability  
in Azure SQL Database

- Full database backups
- Store backups for up to 10 Years
- Read-access geo-redundant storage (RA-GRS)

Export a database

- Generate a BACPAC in external storage and hydrate as needed



# Point-in-time restore

## Automatic backups

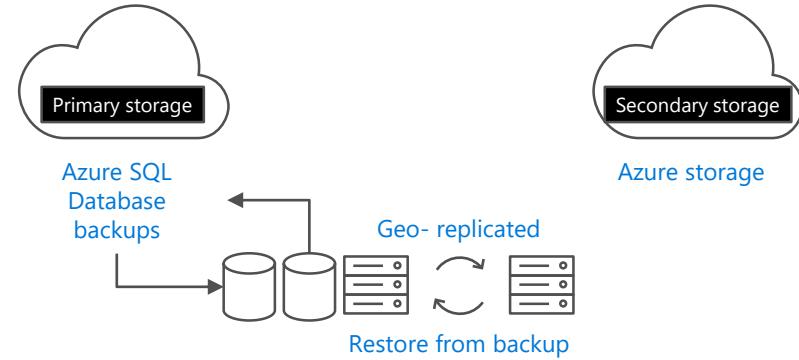
- Full backups weekly, differential backup daily, log backups every 5 minutes
- Daily and weekly backups automatically uploaded to geo-redundant Azure Storage

## Self-service restore

- Point-in-time up to a second granularity
- REST API, PowerShell, Azure CLI, Azure portal
- Creates a new database in the same logical server

## User-controlled retention policy

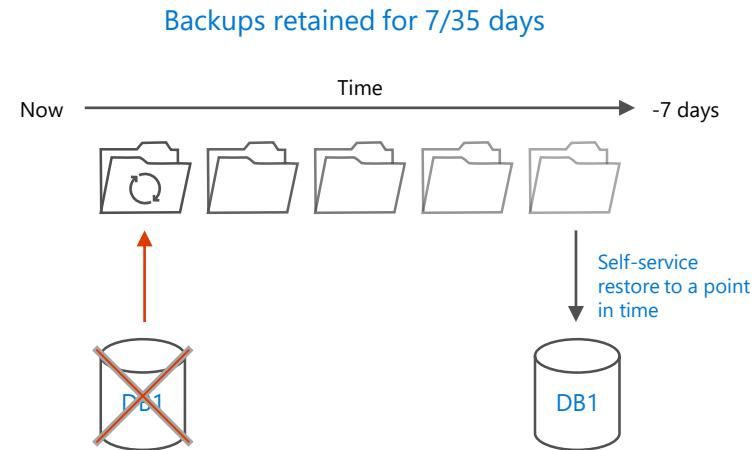
- 7 days default retention in all service tiers
- Up to 35 days of additional retention if required at additional cost
- Choice of storage tier for data sovereignty



# Database recovery

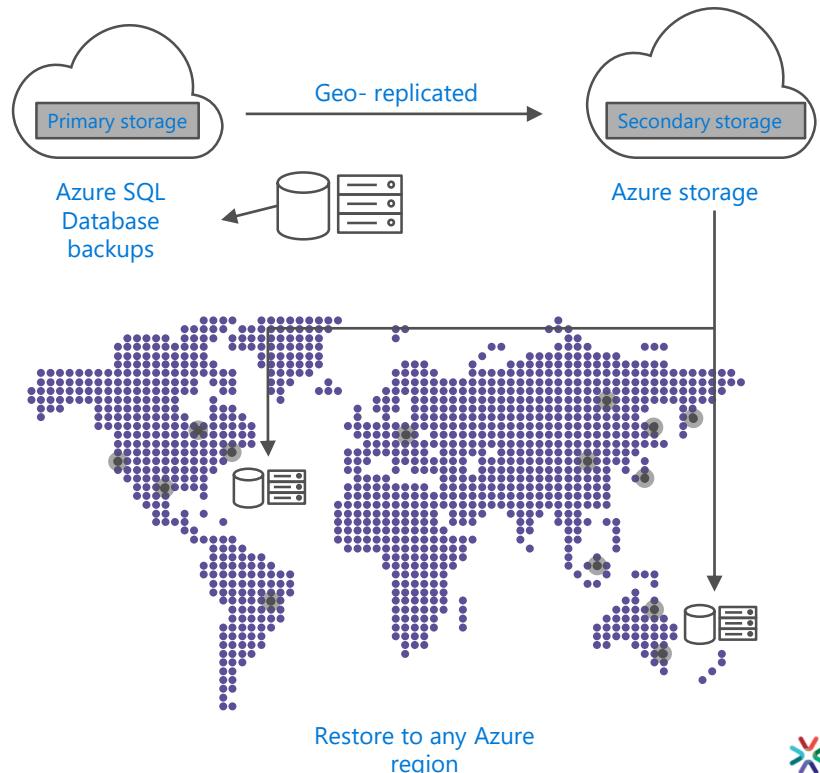
## Restoring a deleted database

- Restores the database to any point in time within the retention period
- Creates a new database on the server used by the original database
- You can choose to switch over to the restored database or use queries to recover data



# Geo-restore protects from disaster

- Self-service restore API
- Built on geo-redundant Azure Storage
- Restores last replicated backup to any Azure region as a new database
- No extra cost, no capacity guarantee
- Estimated recovery time <12h, RPO=1h



# Active geo-replication

Mission-critical business continuity on your terms, via programmatic APIs

**Service levels**

All

**Readable secondaries**

Up to 4 (16 possible)

**Regions available**

Any Azure region

**Replication**

Automatic, asynchronous

**Manageability tools**

REST API, PowerShell, Azure CLI, or Azure Portal

**Recovery time objective (RTO)**

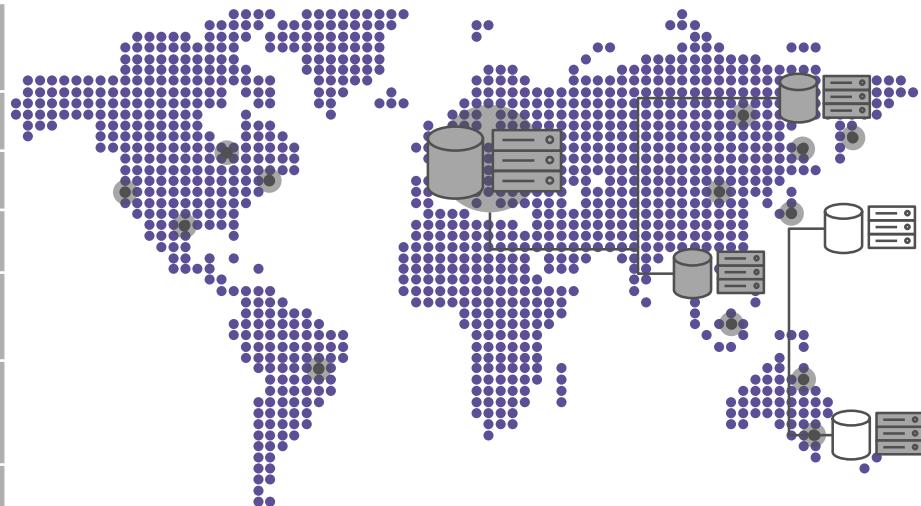
<30 sec

**Recovery point objective**

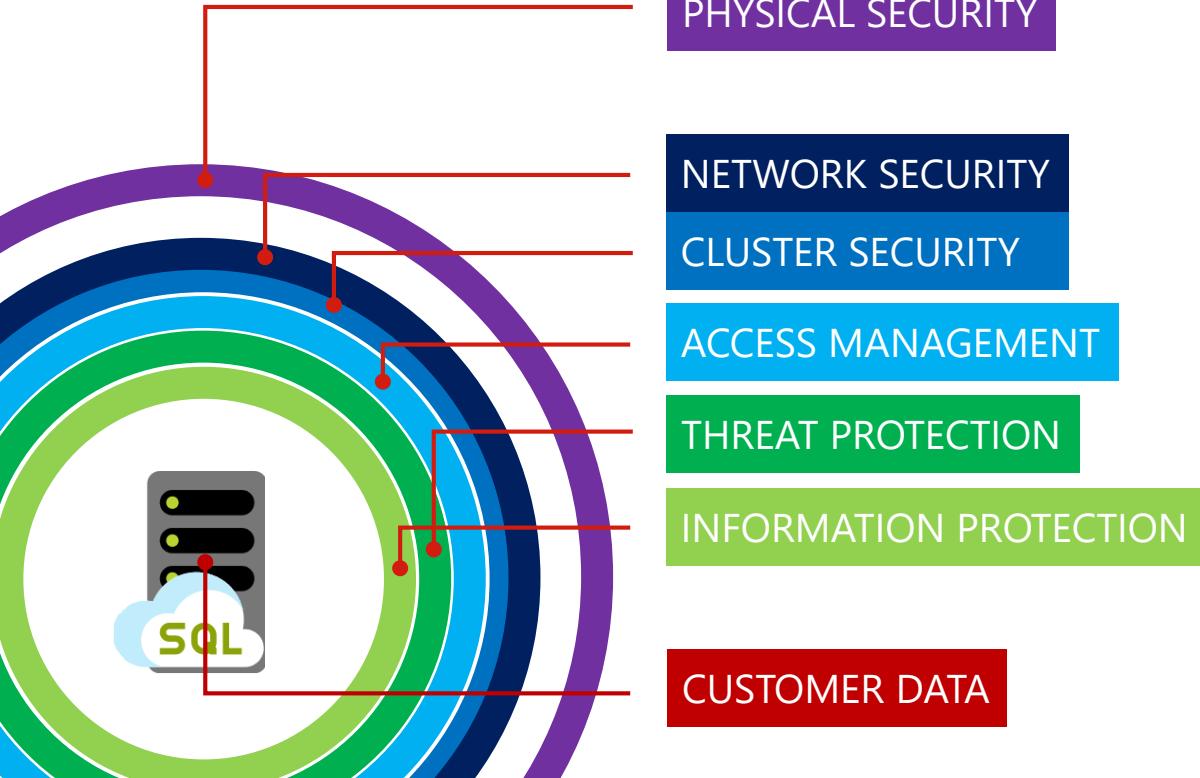
<5sec

**Failover**

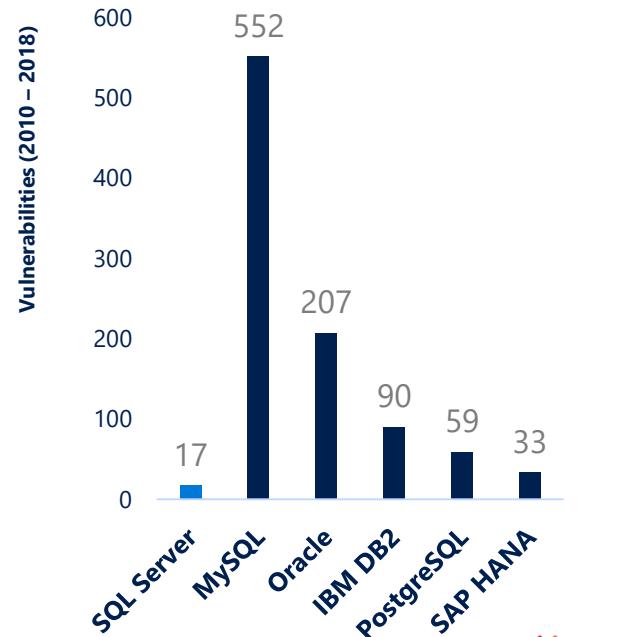
On demand



# Enterprise-grade security that is easy to use



Trusted: most secure over last 7 years



# Azure SQL: The database for all seasons

## Best for all Developers

- “Batteries Included”
- Awesome relational database with 25+ years for product development and 40+ years of research
- Wide range of development languages and tools
- Write code once, deploy anywhere (in any form factor)
- Control and abstraction on storage format (row store and column store)
- Features for modern apps (JSON, Graph, Spatial, memory-optimized tables, temporal tables, and so on)
- Encryption, Masking and Row-Level Security
- Blocking and Non-Blocking Transactions
- Wide range of prices and elasticity

# Azure SQL: The database for all seasons

Best for DBAs

Responsibilities  
of the DBA

Value to the organization

Strategy, Future directions

Designing data life cycle, from  
ingestion to consumption to archival

Performance Tuning, Capacity  
Planning

Securing organizational data assets

Monitoring resources

Building a backup / restore strategy

Infrastructure design

Installing, Upgrading, Patching

# Azure SQL: The database for all seasons

## Best for SaaS Applications

- SQL Server has grown in association with a variety of ISV applications
- Wide range of prices and elasticity:
- Flexible design and deployment choices for multi-tenant applications
- Enterprise-class performance, scale, reliability, security
- Intelligence Query Processing, and Intelligent insights



# Azure SQL: The database for all seasons

## Best for Enterprise Applications

- Enterprise-class built-in Business Continuity (high availability, disaster recovery) backed by financial SLA
- Security “Always a step ahead”
- Meets varieties of government and industry compliance requirements
- World’s best price / performance
- Well-integrated with overall data ecosystem





# Azure SQL Database

Managed Instance

Jovan Popovic, Mladen Andzic  
Program Managers, Microsoft



# Azure SQL

## SQL virtual machines

Best for migrations and applications requiring OS-level access



### SQL virtual machine

- SQL Server and OS server access
- Expansive SQL and OS version support
- Automated manageability features for SQL Server

## Managed instances

Best for most lift-and-shift migrations to the cloud



### Single instance

- SQL Server surface area (vast majority)
- Native virtual network support
- Fully managed service

### Instance pool

- Resource sharing between multiple instances to price optimize
- Enables migration of many small instances at scale
- Fully managed service

## Databases

Best for modern cloud applications. Hyperscale and serverless options are available



### Single database

- Hyperscale storage (up to 100TB)
- Serverless compute
- Fully managed service

### Elastic pool

- Resource sharing between multiple databases to price optimize
- Simplified performance management for multiple databases
- Fully managed service

# Why Azure SQL Managed Instance?

## Instance-level features

- Cross-DB queries
- Linked servers
- CLR
- SQL Agent
- Restore
- Service Broker
- Server-level objects
- Database mail

## Network isolation

- Deployed in VNet
- Private IP address

## Still full PaaS

- High availability
- Automatic backups
- Patching
- Intelligence

CLR?

SQL Agent?

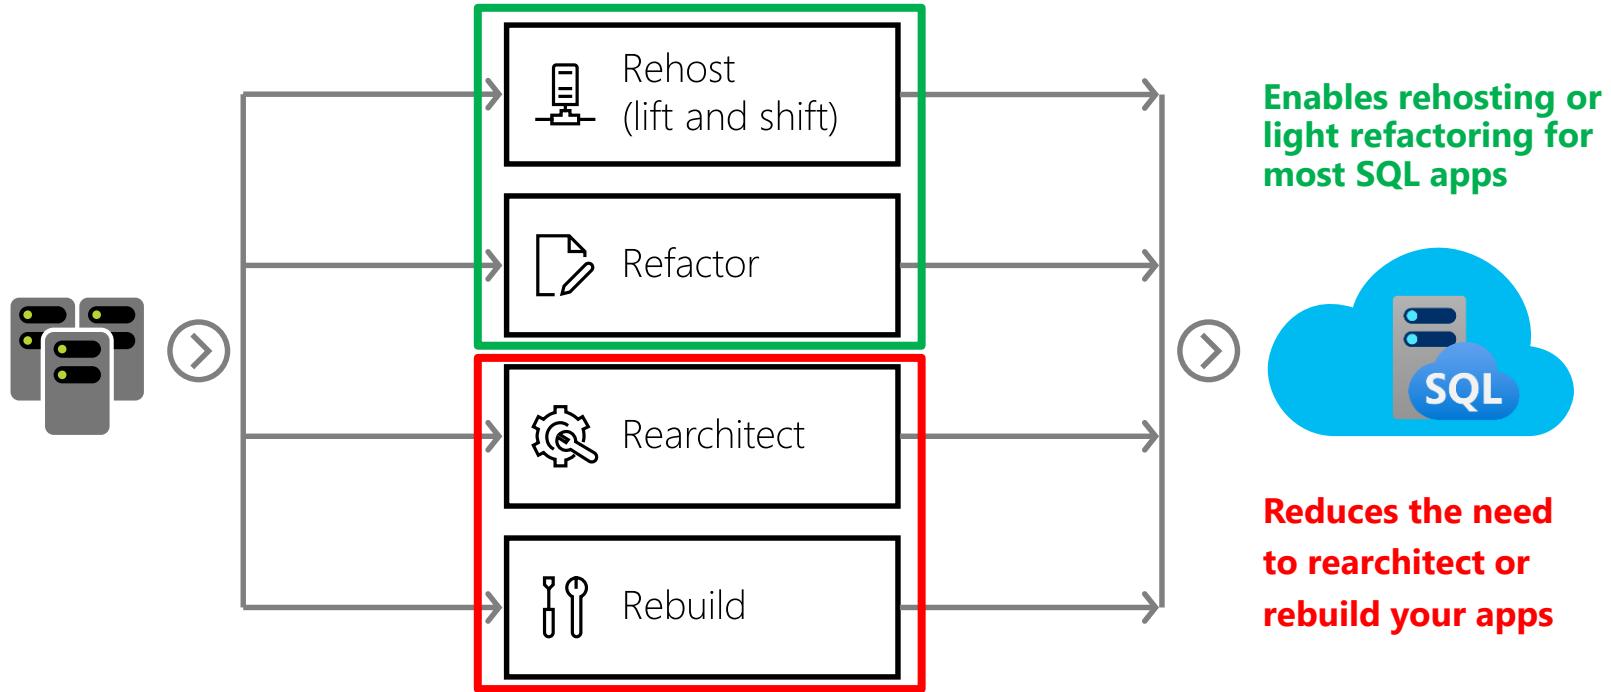
RESTORE?



**Migration to cloud** may require some changes in the solution

# Migration to Cloud

➤ Migrate\*



D E M O

# Migration Assessment

Data Migration Assistant (DMA)



# Proof of Concept – Best Practice

## Start simple

Default settings

Public endpoint

Native backup / restore

General Purpose service tier

## Evolve and optimize

Customized settings

Customized security

Online migration

Consider Business Critical

Auto-failover groups

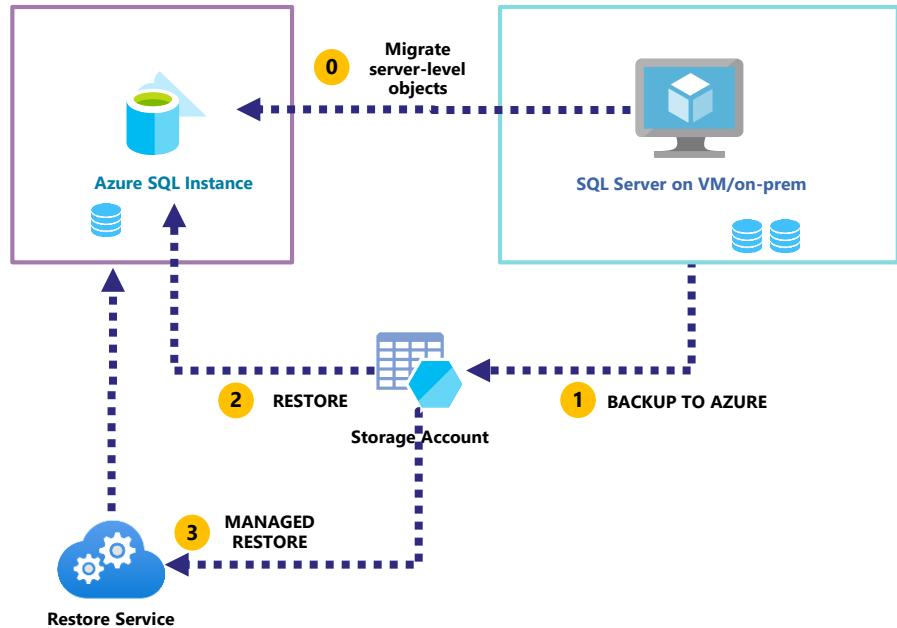
# Easy Database Migration

## Offline

- Native backup/restore

## Online

- Data Migration Service
- Log shipping



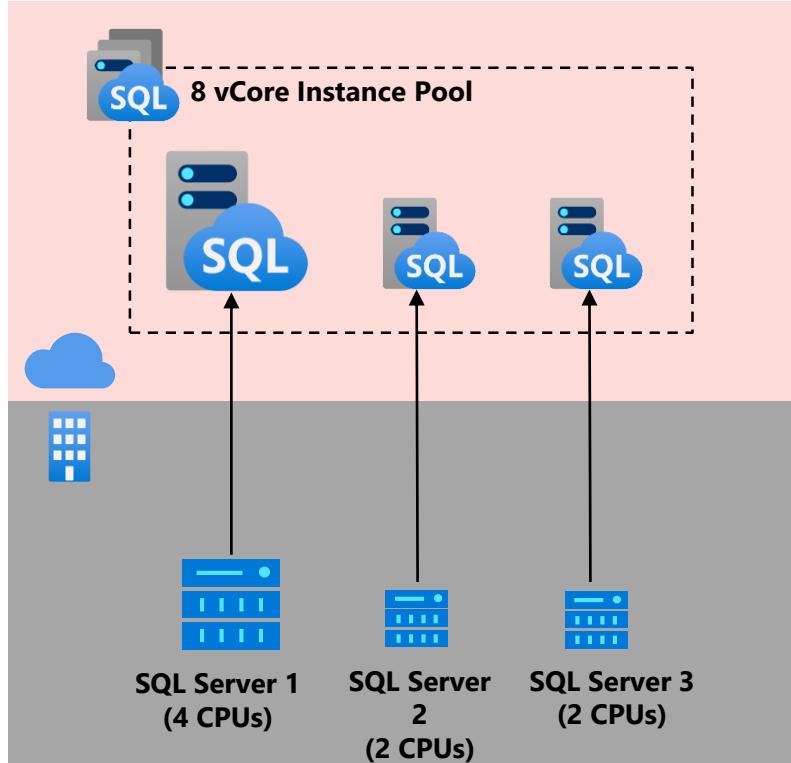
DEMO

# Migration

## Database restore



# Instance Pools (preview)



Migrate multiple **small** SQL instances together to a fully managed instance pool

Add instances starting from 2 vCores up to your pool's limit.

Higher "DBs per vCore" density

Benefit from fast provisioning and scaling operations at the instance level

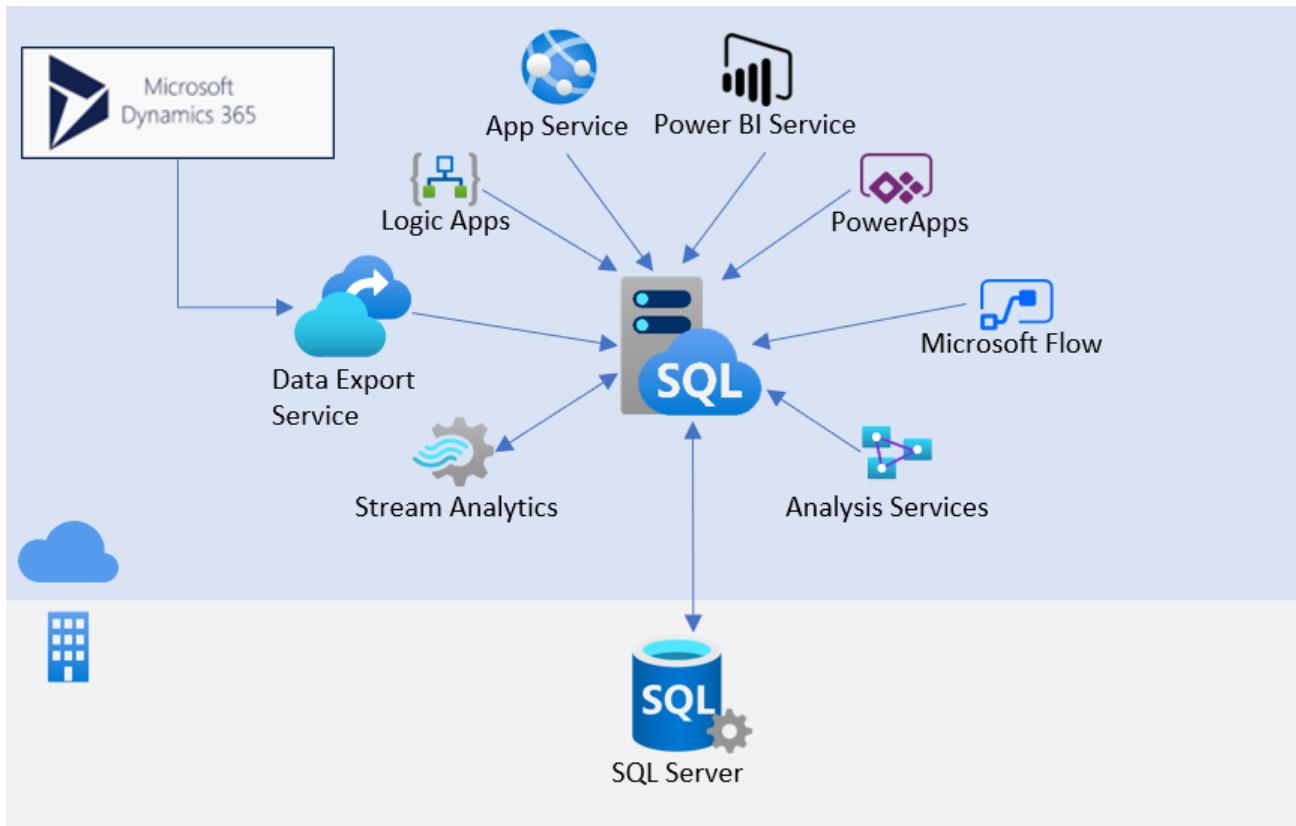


# Modernization

...because migration is just an appetizer

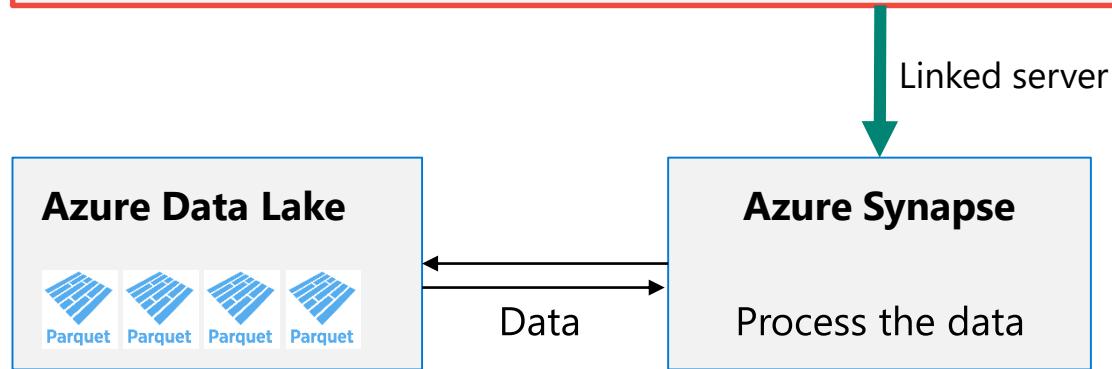


# Integration with Cloud Services



# Query big data via Azure Synapse<sup>NEW</sup>

```
SELECT TOP 100 *
FROM      OPENROWSET(
                      BULK 'https://<storage>/path/to/files/*.parquet',
                      FORMAT = 'Parquet'
) AS [r]
```



# Monitor and Manage Cost

Search (Ctrl+)

«

Save Save as Delete view Cost by resource Share Refresh Export Settings Quickstart tutorial New support request

How satisfied are you with cost analysis? →

ACTUAL COST (USD)

\$146.6K

FORECAST: CHART VIEW ON

--

BUDGET: MI-PERF-N-SCALE

\$10,000 /mo

The chart displays the daily cost accumulation. A horizontal dashed red line at \$10,000 indicates the monthly budget. The cost starts near zero on Oct 1 and rises steadily, reaching approximately \$146.6K by Oct 21, well above the budget limit.

Oct 12

Meter subcategory	Cost (\$)
managed instance general purpose - compute gen5	\$23,455.29
managed instance general purpose - compute gen4	\$18,879.96
managed instance general purpose - storage	\$2,784.64
managed instance business critical - compute gen5	\$2,166.07
all	\$1,377.72
managed instance business critical - compute gen4	\$1,625.43
dv3/dsv3 series windows	\$697.75
premium ssd managed disks	\$748.39
premium page blob	\$512.68
single/elastic pool business critical-compute gen5	\$416.64
Others	\$3,778.92
Monthly budget	\$10,000.00

This chart provides a detailed look at the cost distribution for the current month. It shows various Azure services contributing to the budget. A horizontal dashed red line marks the \$10,000 monthly budget. The total cost for the period shown is exactly \$10,000.

Service name

A gauge chart illustrating the contribution of different services to the total cost. The yellow arc represents the largest portion of the cost.

sql database  
storage  
virtual machines

\$130.2K  
\$4,654.25

Location

A gauge chart showing costs by location. The purple segment represents the highest cost location.

eu west  
us east  
us west central

\$19,579.03  
\$16,368.36

Resource group name

A gauge chart showing costs by resource group. The teal segment represents the highest cost resource group.

clperftesting weu ra  
srki test  
srbozovi test

\$10,188.51  
\$7,692.15



# Learn more

Friday, Nov 8, 9:30 AM

SQL Managed Instance: a Fully  
Managed SQL Server in the Cloud





# Networking for Azure SQL DB and Managed Instance (MI)

Rohit Nayak, Sr. Program Manager,  
Microsoft



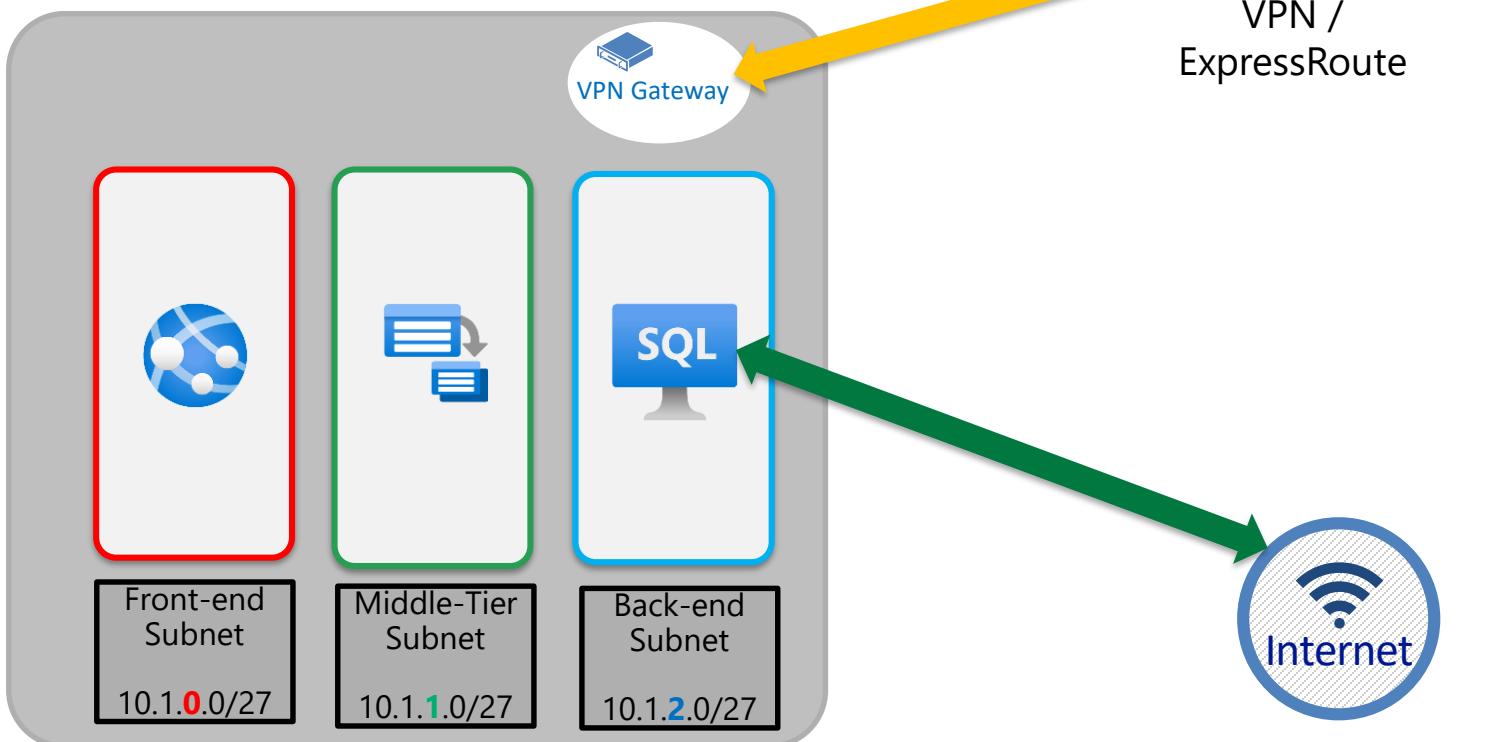
# Agenda

1. Azure Networking 101
2. Connectivity Architecture – SQL DB & Managed Instance (MI)
3. Network Access Controls – SQL DB

# Azure Networking 101

## Virtual Network

10.1.0.0/16 -> [10.1.0.0 : 10.1.255.255]



DEMO

# Quick tour of Azure Networking

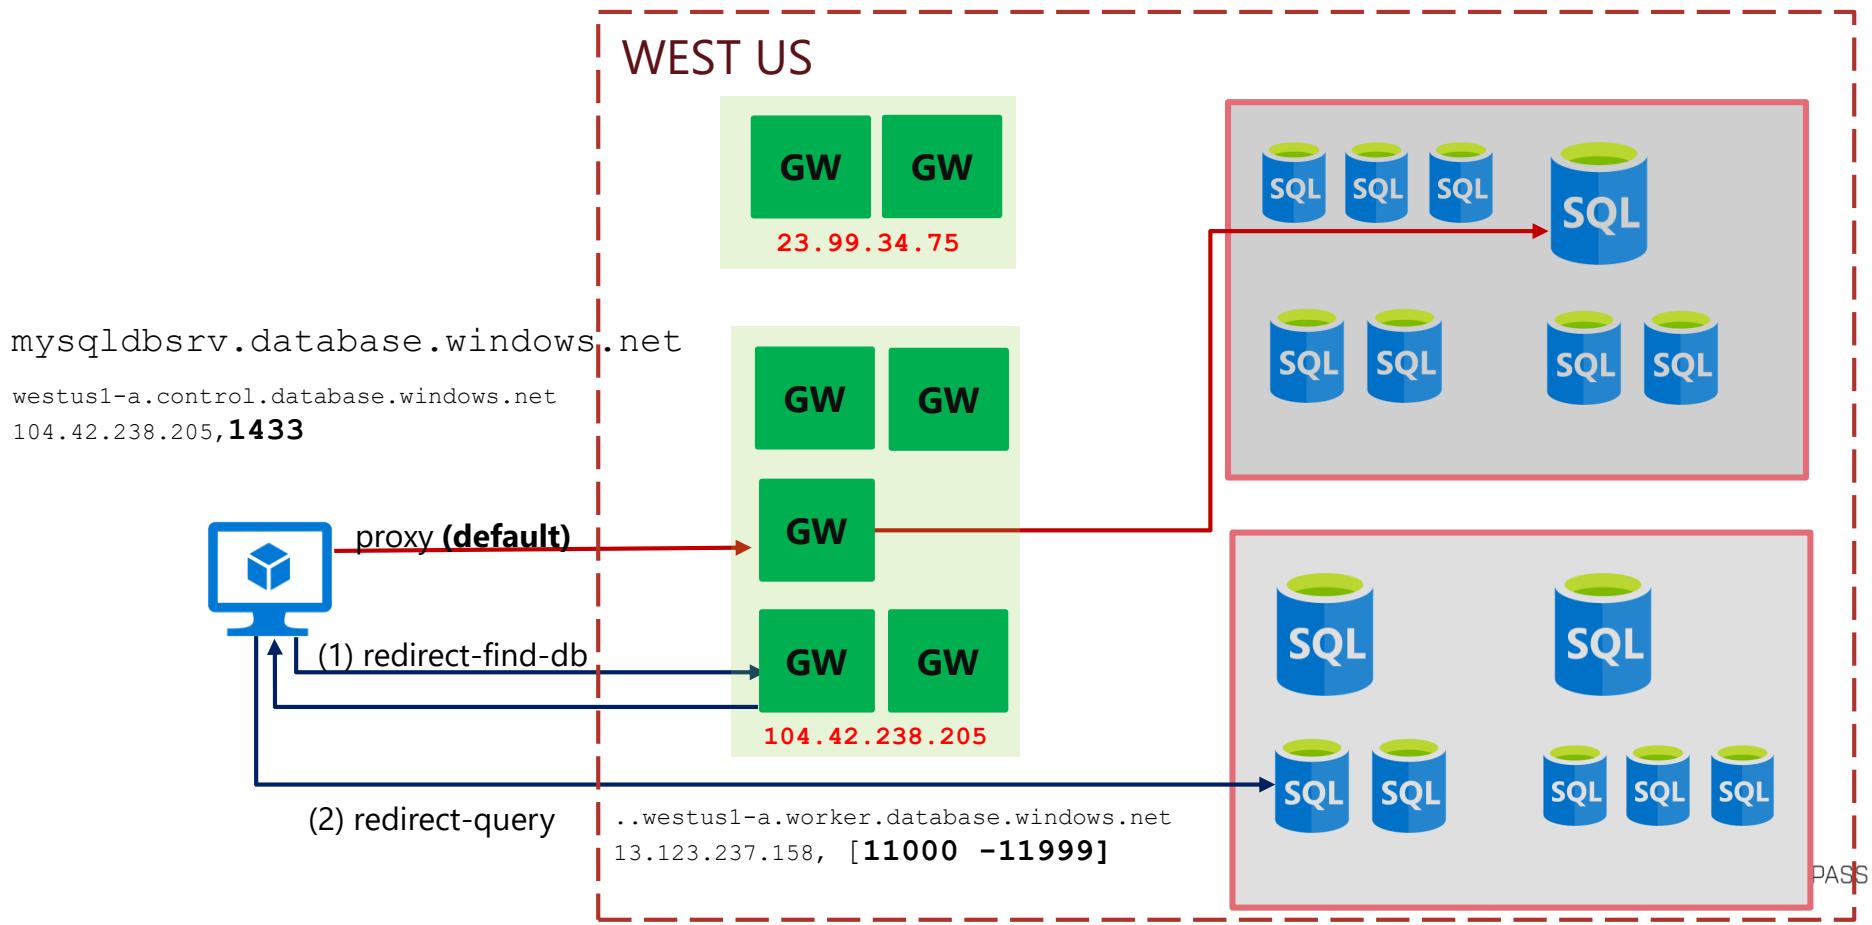




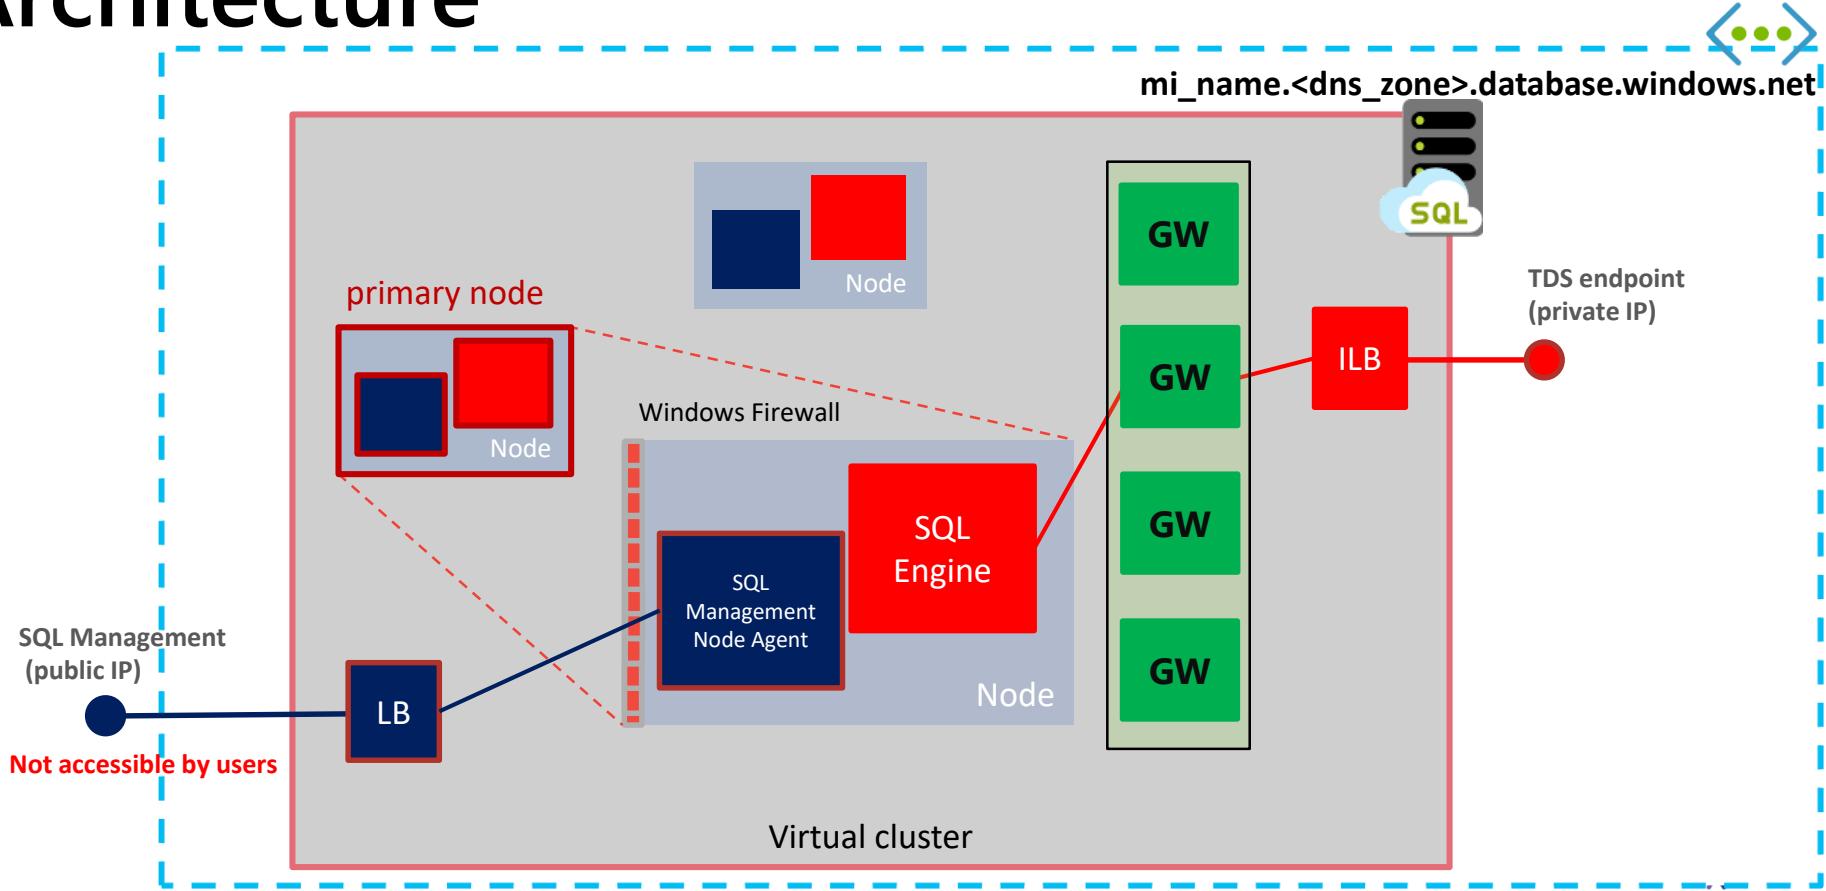
# Connectivity Architecture – SQL DB vs MI



# Sql Database Connectivity Architecture



# Sql Managed Instance Connectivity Architecture





# Network Access Controls – SQL DB



# SQL Database network access controls

Connections from the IPs specified below provides access to all the databases in rndatadbdrv.

Allow access to Azure services  ON  OFF

Allow from inside Azure

Client IP address 131.107.159.236

RULE NAME	START IP	END IP	...
Corpnet	131.107.0.0	131.107.255.255	...

Connections from the VNET/Subnet specified below provides access to all databases in rndatadbdrv.

Virtual networks		+ Add existing virtual network	+ Create new virtual network
RULE NAME	VIRTUAL NETWORK	SUBNET	ADDRESS RANGE
 newVnetRule2	rndatavnet	vmsubnet	10.0.0.32/27

<https://aka.ms/sqlnetworkaccesscontrols>

IP based Firewall rules  
Allow from your laptop or specific Azure VM

VNET Firewall rules  
Allow from specific subnet over private IP



# Private Link

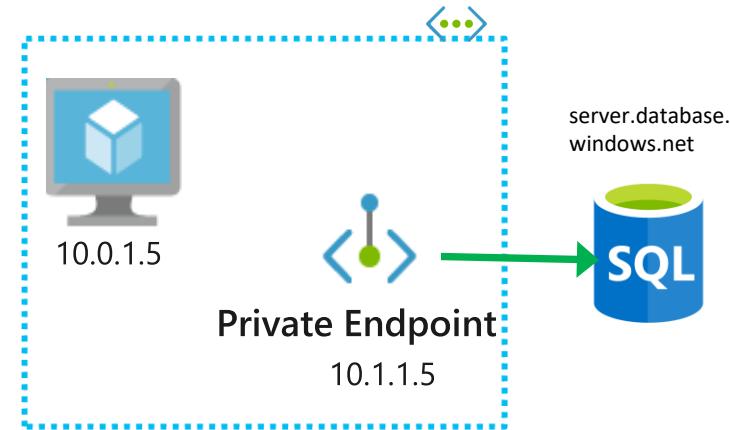
## Private connectivity to PaaS Service from within customers VNET

Benefits:-

- Private endpoint for Azure SQL (server level)
- On-premises connectivity via Express Route/VPN
- Data exfiltration protection
- Can be part of network monitoring strategy using a Network Virtual Appliance (NVA)

Currently in public preview for:

- Azure SQL Database (singleton databases)
- Azure SQL Data Warehouse
- Azure Storage



VIDEO

# Network Access Controls Demo



# Resources

- SQL Database Connectivity Architecture [ <https://docs.microsoft.com/en-us/azure/sql-database/sql-database-connectivity-architecture>]
- SQL Managed Instance Connectivity Architecture[<https://docs.microsoft.com/en-us/azure/sql-database/sql-database-managed-instance-connectivity-architecture>]
- Connect to Azure SQL Database V12 via Redirection  
[<https://techcommunity.microsoft.com/t5/DataCAT/Connect-to-Azure-SQL-Database-V12-via-Redirection/ba-p/305362>]
- Overview of Network Access Controls[<http://aka.ms/sqlnetworkaccesscontrols>]
- Impact of removing Allow Access to Azure Services [<https://docs.microsoft.com/en-us/azure/sql-database/sql-database-networkaccess-overview#allow-azure-services> ]
- IP Based Firewall [<https://docs.microsoft.com/en-us/azure/sql-database/sql-database-firewall-configure>]
- Vnet Firewall Rules [<https://docs.microsoft.com/en-us/azure/sql-database/sql-database-vnet-service-endpoint-rule-overview> ]



# Thank You

## Rohit Nayak



@sqlrohit



rohit.nayak@microsoft.com

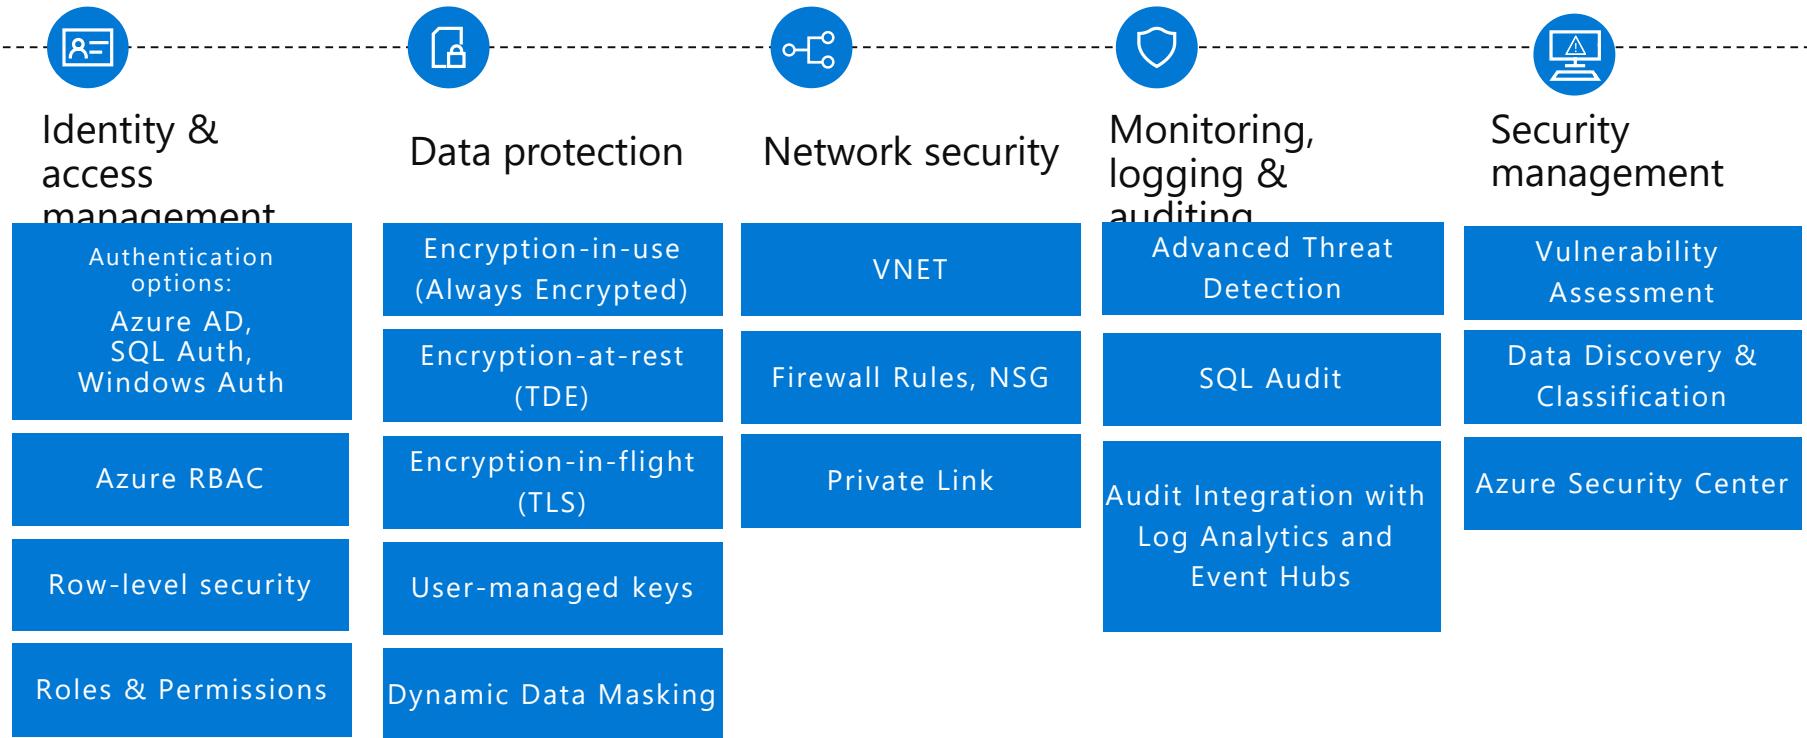


# Security for Azure SQL

Andreas Wolter  
Program Manager SQL Security

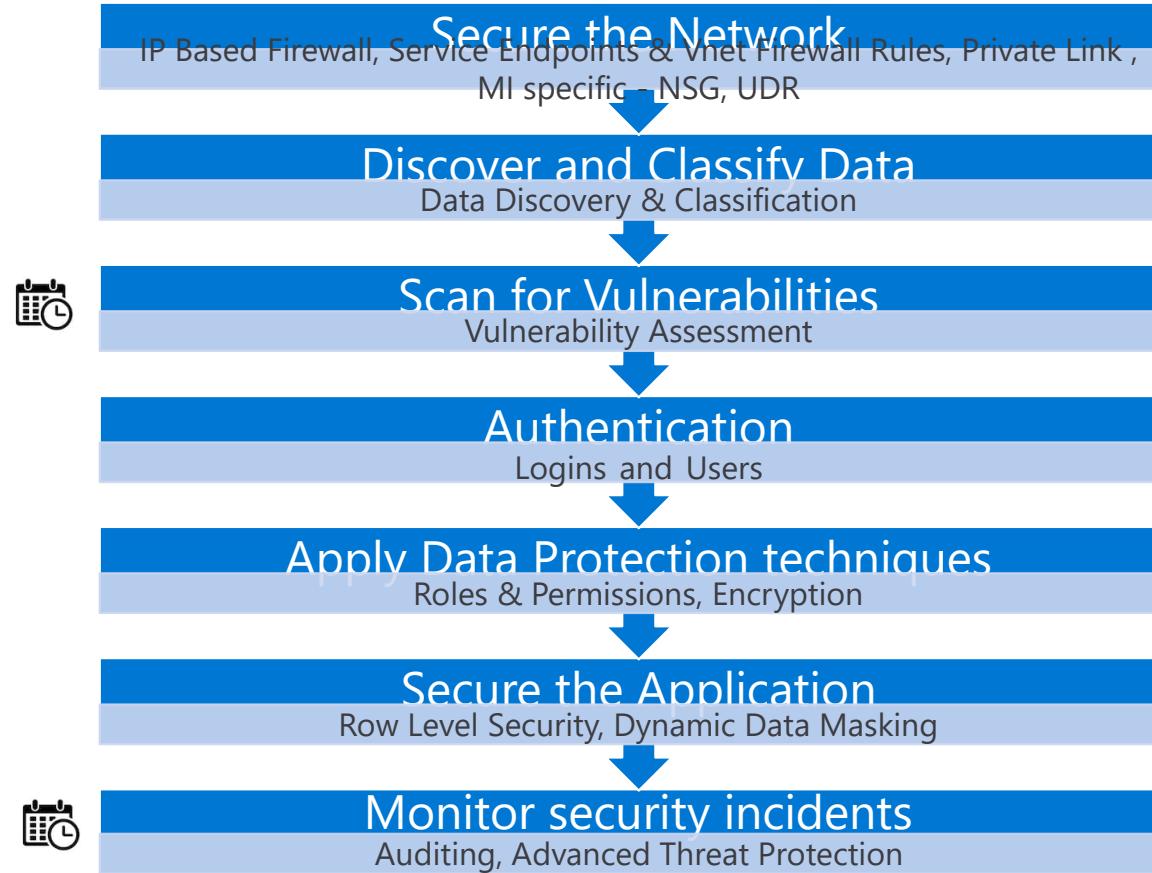


# Enterprise-class Security Controls



+ Partner Solutions: Imperva SecureSphere, IBM Guardium

# Securing your databases in Azure SQL



Every step may need to  
be repeated in case of  
schema changes, new  
users or other needs!



= Run continuously /on schedule



# Check your assets



# Where to start

- You need to know your assets before you know what to protect and how
- Azure SQL supports you in this process:
  - Use Data Discovery to learn about your data
  - Use Vulnerability Assessment (VA) to check your current configuration and set baselines
    - this should be run on a scheduled basis

DEMO

# Data Discovery & Classification, Vulnerability Assessment

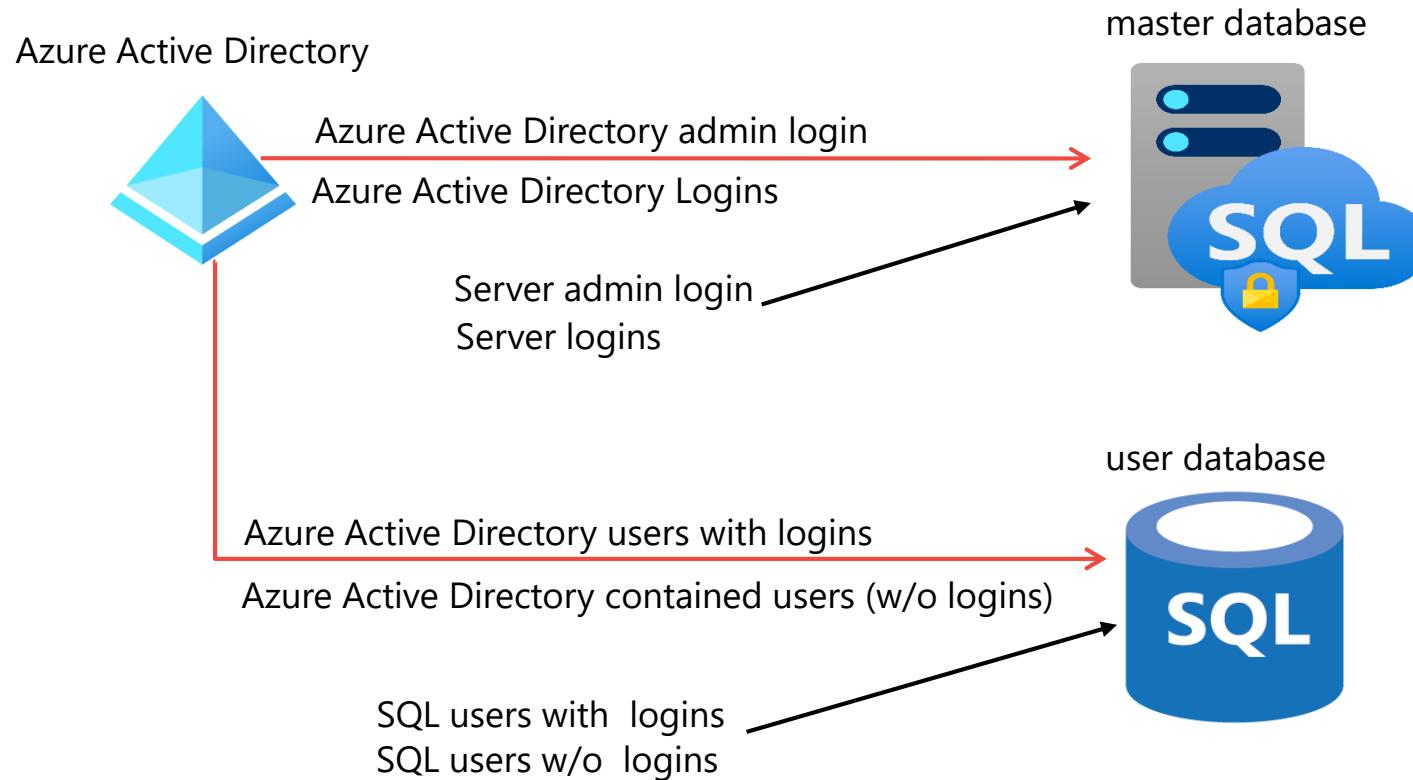




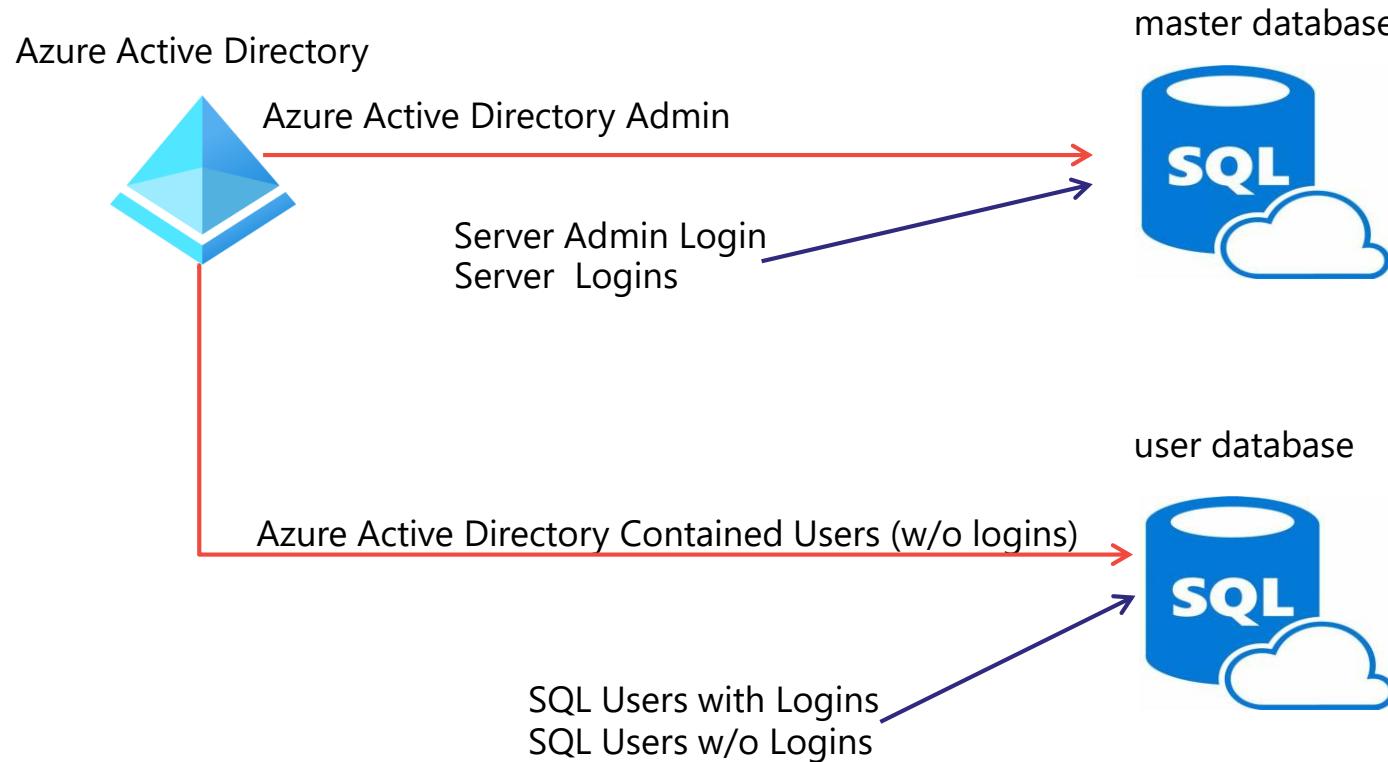
# Authentication



# GA of Azure AD server principals (Azure AD logins) for Managed Instance



# Azure AD for SQL DB





# Data Protection techniques



# Data Protection techniques

- Permissions and Schemas
  - Use the Principle of Least Privilege when granting permissions
    - Azure SQL comes with up to 240 permissions depending on the SKU
  - Inside the databases, schemas can be used to group objects with similar security requirements
- Roles
  - Roles group permissions together and make security manageable
  - There are built-in roles and you can create custom roles that have permissions that suit your needs
- Add Users to Roles, Grant Permissions to Roles



DEMO

# Permissions and Roles



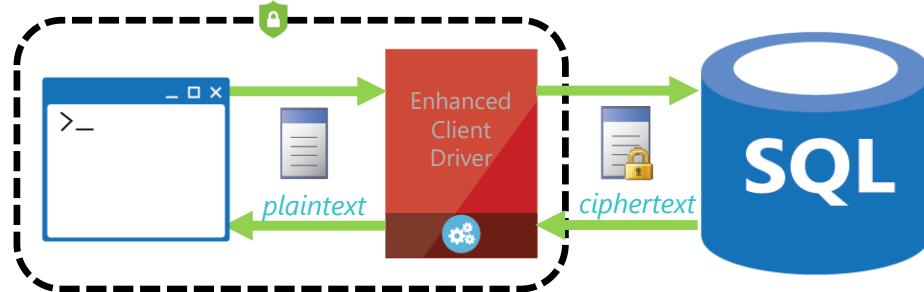
# Encryption

- TDE – protects at rest (by default in Azure since 2017)
- Encryption further strengthens security and can prevent data exfiltration
- Use Always Encrypted to protect data in use from high privileged users by applying Separation of Duties between data owner and data manager (like a DBA)

# Always Encrypted

Protects sensitive data **in use** from high-privileged yet unauthorized SQL users both on-premises and in the cloud

In SQL Server 2016/17 and Azure SQL DB



## Client side Encryption

Client-side encryption of sensitive data using keys that are **never** given to the database system

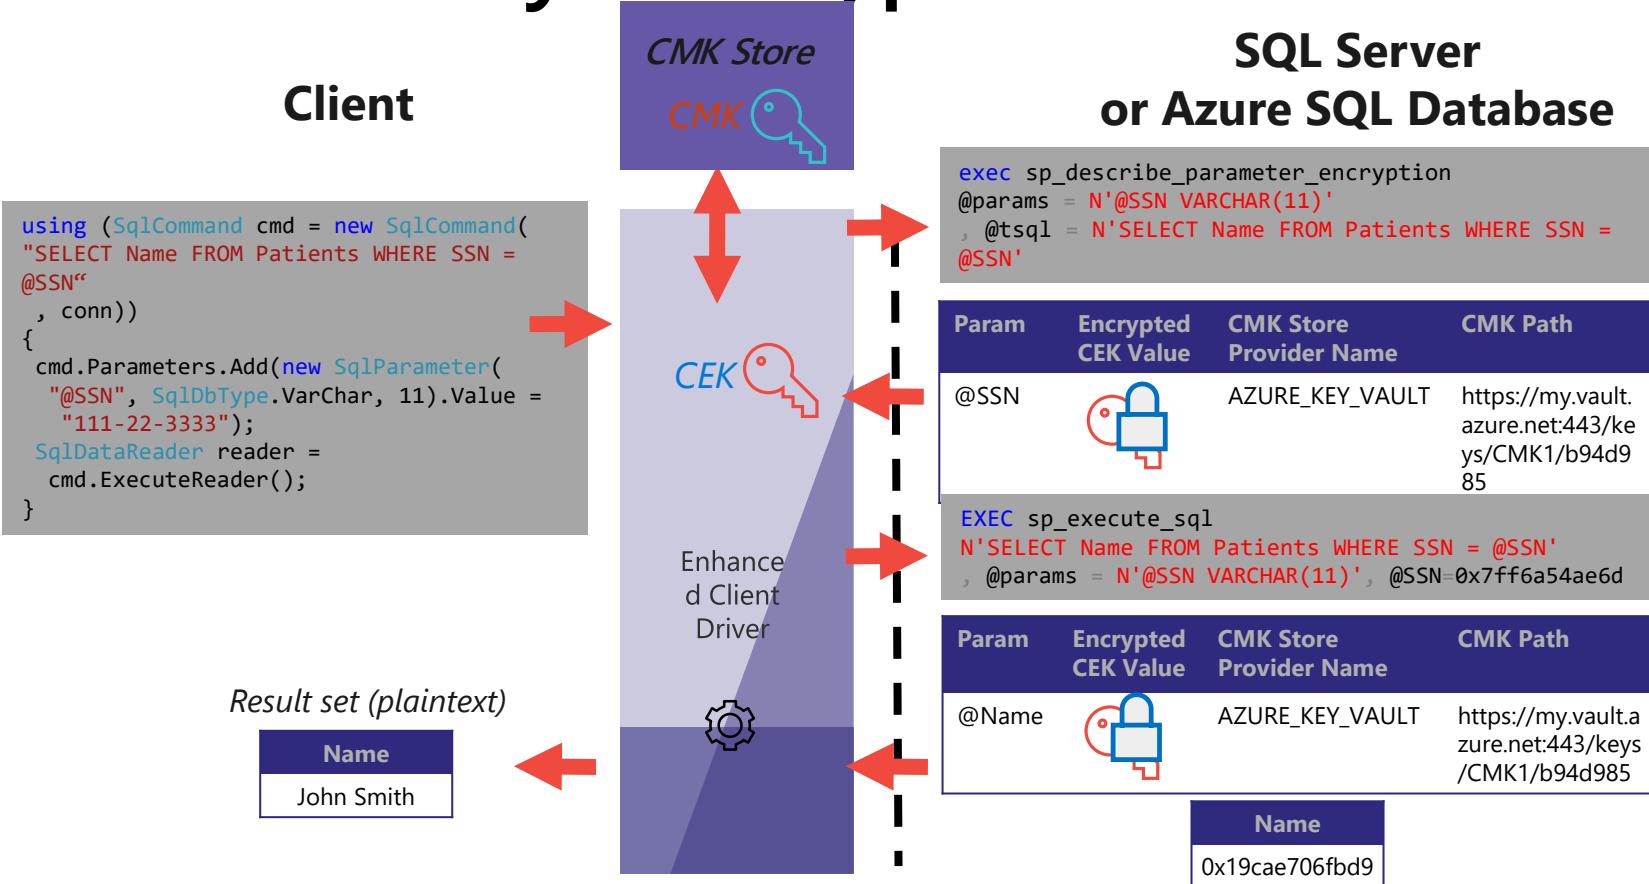
## Encryption Transparency

Client driver transparently encrypts query parameters and decrypts encrypted results

## Queries on Encrypted Data

Support for equality comparison, including join, group by and distinct operators via deterministic encryption

# How Always Encrypted works





# Secure the Application



# Application security features

- intended for scenarios where the queries that a user can execute are controlled by a middle-tier application.
- users with ad-hoc query access to the database may be able to infer the existence of filtered data, using side channels.
- Dynamic Data Masking
  - Obfuscates sensitive data in the application, using built-in or custom masking rules
- Row Level Security use case:
  - limit users' access to only certain rows of data in a database

# How Row Level Security works

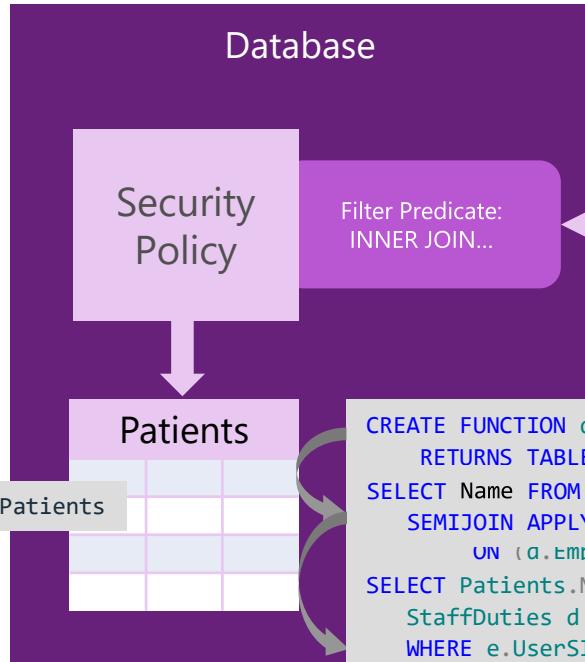
- ③ Security Agent runs a security filter from Patients table, applying policy predicate binding the predicate to the Patients table

Nurse  
(employee ID=100986)



Application

```
SELECT Name FROM Patients
```



Policy Manager



```
CREATE FUNCTION dbo.fn_securitypredicate(@wing int)
RETURNS TABLE WITH SCHEMABINDING AS
SELECT Name FROM Patients
SEMIJOIN APPLY dbo.fn_securitypredicate(patients.Wing);
    ON (a.EmpId = e.EmpId)
SELECT Patients.Name FROM Patients,
StaffDuties d INNER JOIN Employees e ON (d.EmpId = e.EmpId)
WHERE e.UserId = SUSER_SID() AND Patients.wing = d.Wing;
WITH (STATE = UN)
```

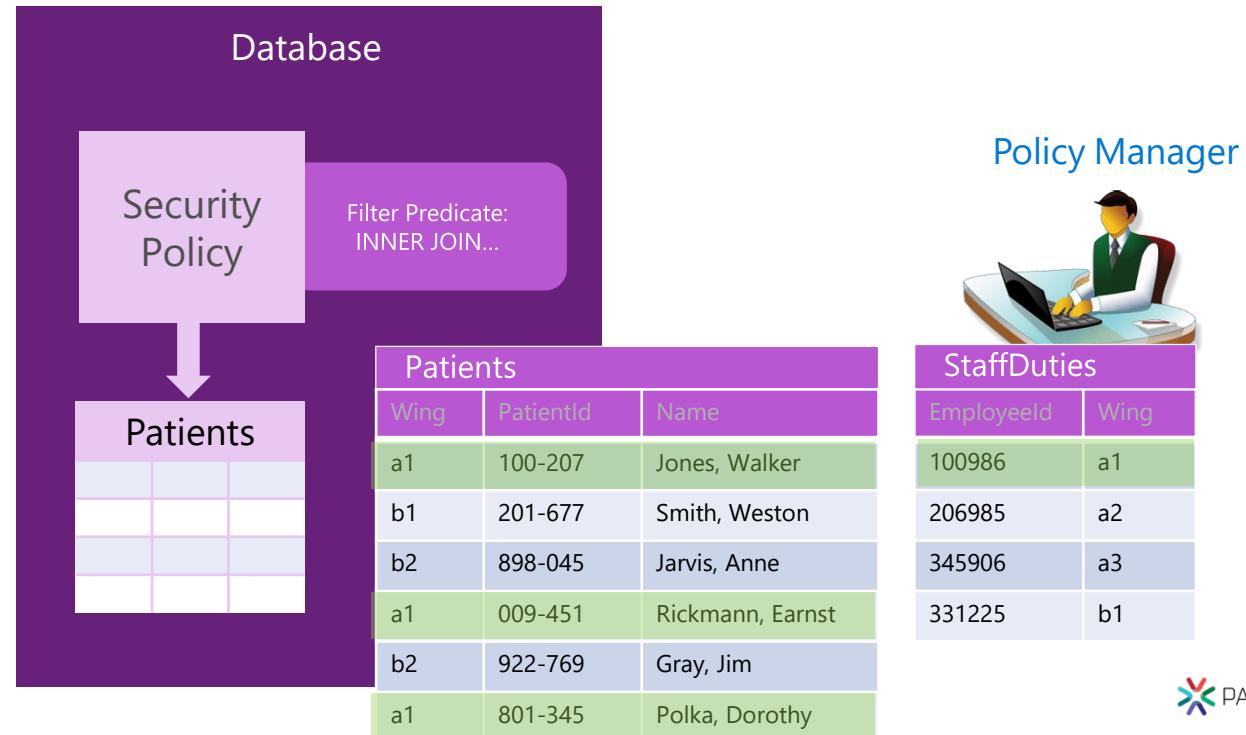
# How Row Level Security works

4) Filtered rows are returned and nurse sees only the patients on her wing

Nurse  
(employee Id=100986)



Application



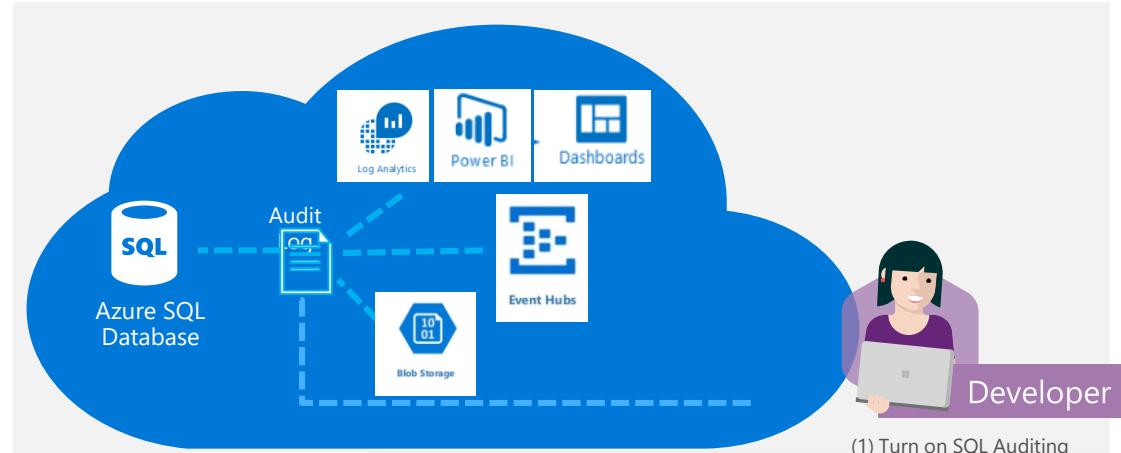




# Monitor security incidents



# SQL Auditing with Log Analytics and Event Hubs



The screenshot shows the Azure Log Analytics interface with a search bar containing the query: 'where Category == "SQLSecurityAuditEvents" | project TimeGenerated, server\_principal\_name\_s, statement\_s, affected\_rows\_d, SeverityLevel | sort by TimeGenerated asc'. The results table displays 62 rows of audit log data. The columns are: TYPE, LOGICALSERVERNAME\_S, CATEGORY, and statement\_s. The statement\_s column contains several lines of T-SQL code, including 'exec sp\_executesql N'SELECT @Name AS [Name], SCHEMA\_NAME(...)', 'exec sp\_executesql N'SELECT ISNULL(@PERMS\_BY\_NAME)QUOTEN...', and 'IF OBJECT\_ID (N'[sys].[database\_query\_store\_option]') IS NOT NULL BE...'. The developer is also visible in the background of the screenshot.

- ✓ Configurable via audit policy
- ✓ SQL audit logs can reside in
  - Azure Storage account
  - Azure Log Analytics
  - Azure Event Hubs
- ✓ Rich set of tools for
  - Investigating security alerts
  - Tracking access to sensitive data

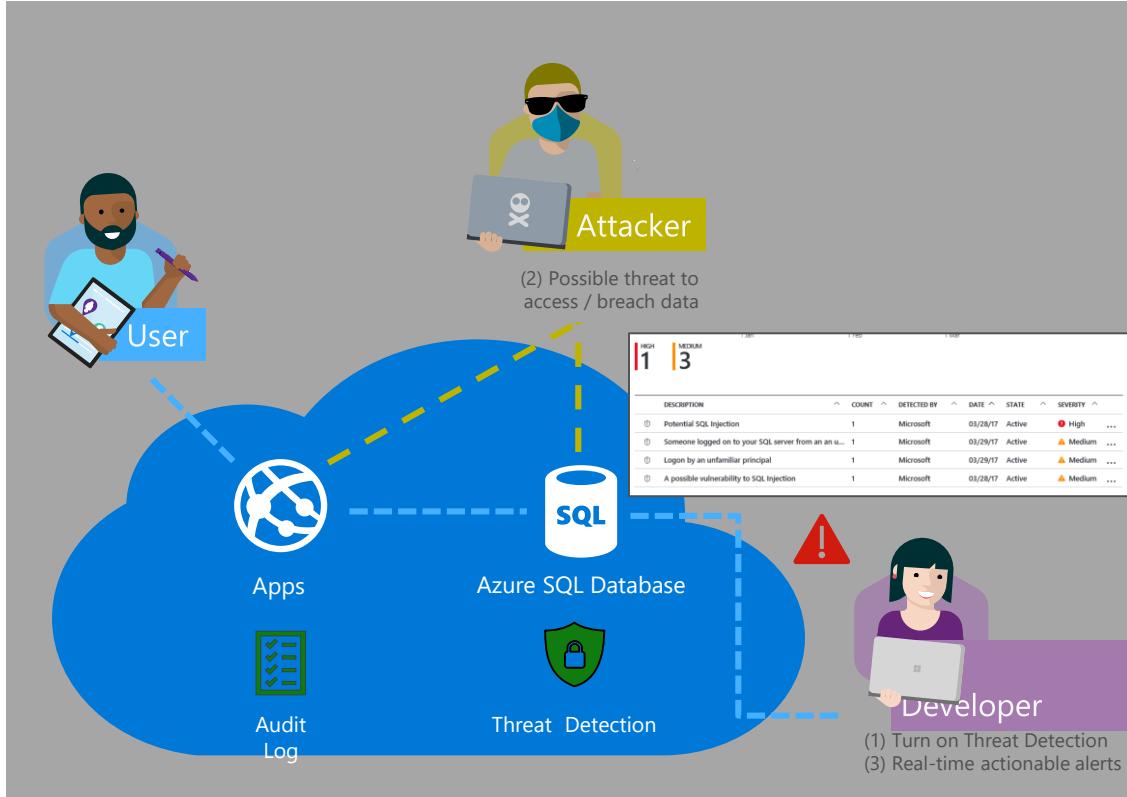
D E M O

# Auditing Sensitive Data Access



# Advanced Threat Protection

Detect unusual and harmful attempts to breach your database.



- ✓ Detects potential **SQL injection** attacks
- ✓ Detects **unusual access & data exfiltration** activities
- ✓ **Actionable** alerts to investigate & remediate
- ✓ **View alerts** for your entire Azure tenant using Azure Security Center

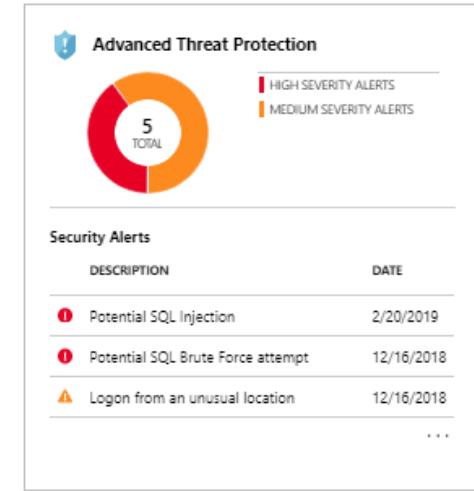
# SQL Threat Protection Suite

## Anomalous access patterns

1. Someone has logged from an unusual location - change in the access pattern from an unusual geographical location
2. An unfamiliar principal successfully logged- change in the access pattern using an unusual SQL user.
3. Someone is attempting to brute force SQL credentials abnormally high number of failed logins with different credentials.
4. Someone has logged from a potentially harmful application

## Potential SQL injection attacks

5. SQLi attempt ([SQL Server 2019](#)) - An application generated a faulty SQL statement, which may indicate a potential vulnerability of the application to SQL injection.
6. SQLi attack ([SQL Server 2019](#)) - Potential exploitation of application code vulnerability to SQL Injection, which may indicate a SQL Injection attack.



## Anomalous queries patterns

8. Unsecure commands ([SQL Server 2019](#)) - Someone has executed unsecure commands (e.g. xp\_cmdshell...)
9. Data exfiltration by volume ([coming at GA](#)) - someone has extracted anomalous amounts of data in an hour or using a single query



# Advanced Threat Protection





# Thank You

## Andreas Wolter



[@AndreasWolter](https://twitter.com/AndreasWolter)



[andreas.wolter@microsoft.com](mailto:andreas.wolter@microsoft.com)



# Azure SQL Database

Best cloud database for DEVs

Davide Mauri  
Senior Program Manager



# Azure SQL is the best cloud db for devs

## Really! Why?

- “Batteries included” philosophy.
  - Yep, that’s a [cit. from Python](#) 😊
  - Means that you have access to \*a lot\* of cool features, all in one place
    - Columnstore/Rowstore, Graph, JSON, XML(!), Spatial, Encryption, Analytics and, of course, Relational (with its declarative power)
- Based on the finely tuned SQL Server engine...just cloud-scale
  - Up to 100TB. Several tiers (General Purpose, Business Critical, Hyperscale, Serverless)

# Azure SQL is the best cloud db for devs



# Choose your own language

Azure SQL is the SQL Server in the cloud

Build an app using SQL Server

Get started with SQL Server on macOS, Linux, and Windows.

The screenshot shows a user interface for choosing a development environment. At the top, a purple bar says "Choose a language and operating system below". Below it are several cards representing different stacks:

- C#** (circled in red)
- Java**
- node**
- php**
- python**
- Ruby**
- Go**
- Perl**
- VB.NET**

Each card has a list of supported operating systems:

- C# supports macOS, RHEL, SLES, Ubuntu, and Windows.
- Java supports macOS, RHEL, SLES, and Ubuntu.
- node supports macOS, RHEL, SLES, Ubuntu, and Windows.
- php supports macOS, RHEL, SLES, Ubuntu, and Windows.
- python supports macOS, RHEL, SLES, Ubuntu, and Windows.
- Ruby supports macOS, RHEL, SLES, Ubuntu, and Windows.
- Go supports macOS, RHEL, SLES, Ubuntu, and Windows.
- Perl supports macOS, RHEL, SLES, Ubuntu, and Windows.
- VB.NET supports macOS, RHEL, SLES, Ubuntu, and Windows.

A blue box highlights the Ubuntu option under the Java stack.

<https://www.microsoft.com/en-us/sql-server/developer-get-started>

# Learning Paths

Learning for everyone

The screenshot displays a grid of nine learning modules for Azure SQL, each with a thumbnail icon, title, duration, rating, and level.

Module Title	Duration	Rating	Level
Secure your cloud data	6 hr 13 min	4.7 (519)	Beginner   Developer
Persist and retrieve relational data with Entity Framework Core	1 hr 4 min	4.7 (519)	Intermediate   Developer
Scale multiple Azure SQL Databases with SQL elastic pools	39 min	4.5 (535)	Beginner   Developer
See Azure in action	30 min	4.6 (497)	Beginner   Business User
Analyze images in real-time with machine learning, Azure IoT Hub and Azure Stream Analytics	1 hr 53 min	4.7 (171)	Beginner   Developer
Develop and configure an ASP.NET application that queries an Azure SQL database	50 min	4.7 (132)	Beginner   Solutions Architect
Secure your Azure SQL Database	1 hr 7 min	4.6 (932)	Intermediate   Administrator
Design for performance and scalability in Azure	48 min	4.7 (929)	Intermediate   Solutions Architect
Design for availability and recoverability in Azure	59 min	4.8 (841)	Intermediate   Solutions Architect

<https://aka.ms/learn-azure-sql>

The screenshot shows the homepage of the sqlworkshops website, which is a collection of SQL Server and Azure SQL Labs and Workshops.

**sqlworkshops**  
SQL Server Workshops

Project maintained by [microsoft](#) Hosted on GitHub Pages — Theme by [matgraham](#)

**Microsoft**

**SQL Server and Azure SQL Labs and Workshops**  
(<https://aka.ms/sqlworkshops>)

This site is a map of learning content produced by and curated by the SQL Server and Azure SQL teams in Microsoft Engineering. These materials are meant to be instructor-led, but you can work through the materials on a test system on your own if desired. Labs are shorter and Workshops are more comprehensive. You can view all materials directly in this interface, or you can [view the raw github site for this content here](#).

See the license information at the bottom of this [README.md file](#)

Find a problem? Spot a bug? [Post an issue here](#) and we'll try and fix it.

**SQL Server Data Platform**

- » [Workshop: SQL Server Ground to Cloud](#)
- » [Workshop: SQL Server 2019 on OpenShift](#)
- » [Lab: SQL Server 2019](#)

**Azure SQL**

- » [Workshop: SQL Server Ground to Cloud](#)
- » [Lab: Microsoft Azure SQL Labs](#)

**Programming**

- » [Lab: Python for Data Professionals](#)
- » [Workshop: R for Data Professionals](#)

<https://aka.ms/sqlworkshops>



# CI/CD support in GitHub

## GitHub Actions

The screenshot shows the GitHub repository page for `Azure/sql-action`. The top navigation bar includes links for Code, Issues (0), Pull requests (0), Actions, Projects (0), Wiki, Security, and Insights. The repository statistics show 21 commits, 3 branches, 1 release, 4 contributors, and an MIT license. The commit history lists several recent changes, including updates to README.md, tests, lib, and src, as well as the introduction of Azure Sql Action and the initial commit. A prominent URL at the bottom is <https://github.com/Azure/sql-action>.

File / Commit Message	Description	Time
<code>N-Usha Update README.md</code>	Latest commit 51dca9d 3 days ago	
<code>_tests_</code>	Use common webclient and code refactor	24 days ago
<code>lib</code>	Use common webclient and code refactor	24 days ago
<code>src</code>	Use common webclient and code refactor	24 days ago
<code>.gitignore</code>	Introduce Azure Sql Action	last month
<code>CODE_OF_CONDUCT.md</code>	Initial commit	last month

<https://github.com/Azure/sql-action>

# Getting Started

Many sample databases ready to be used

- [WideWorldImporters](#)
  - .bak => Azure SQL MI
  - .bacpac => Azure SQL Database
- [SqlPackage.exe](#) to import .bacpac or Azure Portal/PowerShell/AZ CLI/REST API
  - [Azure SQL Database Import](#)

```
$importRequest = New-AzSqlDatabaseImport ` 
    -ResourceGroupName "dm-pass-2019" ` 
    -ServerName "dmsqlsrv-pass2019" ` 
    -DatabaseName "WideWorldImportersFull2" ` 
    -DatabaseMaxSizeBytes "$(100 * 1024 * 1024 * 1024)" ` 
    -StorageKeyType "StorageAccessKey" ` 
    -StorageKey $(Get-AzStorageAccountKey -ResourceGroupName "dm-pass-2019" -StorageAccountName "dmpass2019").Value[0] ` 
    -StorageUri "https://dmpass2019.blob.core.windows.net/bak/WideWorldImporters-Full.bacpac" ` 
    -Edition "BusinessCritical" ` 
    -ServiceObjectiveName "BC_Gen5_2" ` 
    -AdministratorLogin "<admin-login>" ` 
    -AdministratorLoginPassword $(ConvertTo-SecureString -String "<admin-password>" -AsPlainText -Force) 

Get-AzSqlDatabaseImportExportStatus -OperationStatusLink $importRequest.OperationStatusLink
```

```
storageKey=$(az storage account keys list -g dm-pass-2019 -n dmpass2019 --query "[0].value" -o 
az sql db import\ 
    -g dm-pass-2019 \
    -s dmsqlsrv-pass2019 \
    -n $sqlDatabase \
    -u $sqlLogin \
    -p $sqlPassword \
    --storage-key-type StorageAccessKey \
    --storage-uri https://dmpass2019.blob.core.windows.net/bak/WideWorldImporters-Full.bacpac \
    --storage-key $storageKey
```

# Connecting to Azure SQL

## Microsoft.Data.SqlClient

- The new future-proof library to connect and use Azure SQL
  - <https://devblogs.microsoft.com/dotnet/introducing-the-new-microsoftdatasqlclient/>
  - Supports both .NET and .NET Core
  - Open Source
    - <https://github.com/dotnet/SqlClient>
  - Install via NuGet
  - Drop-In replacements for System.Data.SqlClient

```
create-new-app.bat
```

```
1  mkdir myapp  
2  
3  dotnet create console  
4  
5  dotnet add package Microsoft.Data.SqlClient  
6  
7  [ ... ]  
8  
9  dotnet run  
10
```

```
C# Program.cs > ...
```

```
1  using Microsoft.Data.SqlClient;  
2  
3  namespace Azure.SQLDB.Samples.Connection  
4  {  
    0 references  
5      class Program  
6      {  
        0 references  
7          static void Main(string[] args)  
8          {  
            // Connection string can be found in the portal  
10             using(var conn = new SqlConnection("<connection string>"))  
11             {  
12                 conn.Open();  
13                 var cmd = new SqlCommand("SELECT Column1 FROM dbo.Table1", conn);  
14                 var result = cmd.ExecuteScalar();  
15                 conn.Close();  
16             }  
17         }  
18     }  
19 }
```

# Using Azure SQL from .NET

## Three main options

- `SqlClient` classes
  - You're in control of everything
- Full-featured ORM like Entity Framework
  - You want the tool to take care of all database interaction for you
- Lightweight ORM like MicroORM
  - You just want to execute queries as fast as possible and map it to a class
    - [Dapper](#), [Massive](#), [PetaPoco](#), [Simple.Data](#), [OrmLite](#)

# Using Azure SQL from .NET

## Dapper Sample

```
public class Dog
{
    public int? Age { get; set; }
    public Guid Id { get; set; }
    public string Name { get; set; }
    public float? Weight { get; set; }

    public int IgnoredProperty { get { return 1; } }
}

var guid = Guid.NewGuid();
var dog = connection.Query<Dog>("select Age = @Age, Id = @Id", new { Age = (int?)null, Id = guid });
```

<https://medium.com/dapper-net>

# Pushing compute to Azure SQL

## Use everything you're paying for

- Azure SQL offer a great amount of features to do complex data manipulation
  - And uses a lot of advanced optimization techniques
    - Batch (or vector) execution, using CPU [SIMD instructions](#)
    - Intelligent Query Execution (adaptive joins, memory grant feedbacks, etc..)
    - Data and Execution Plan (aka DAG) caching
    - Data pre-fetching
    - Data Distribution Statistics awareness
  - Keep your application as simple as possible and focused on business logic processing (eg: stored procedure)

# Pushing compute to Azure SQL

## Windowing Functions

- Allows complex data manipulation over subset of data (windows)
  - Ranking, Running Totals, Moving Averages, One-vs-Total Ratios,
  - Allows access to previous and next rows in the same (ordered) window
  - Learning how to master them can save *hours or days* of development
    - Simpler applications
    - Less maintenance, better performance scale

DEMO

# Windowing Functions



# Pushing compute to Azure SQL

## String Manipulation

- Historically a weak point of SQL Server / Azure SQL
  - Required data to be sent to the app and back: No more!
    - Less chatty app = less roundtrips = better performances + save money
  - STRING\_SPLIT, STRING\_AGG, CONCAT, CONCAT\_WS, TRANSLATE

# Keeping Data Consistent

## Transaction

- Azure SQL uses a row-version approach to keep consistency
  - READ COMMITTED SNAPSHOT ISOLATION LEVEL enabled by default
  - Reads will not be blocked by writes!
  - Reads will not block writes!
  - What if you \*need\* locking instead?
    - Usual Isolation Levels are always available
      - READCOMMITTEDLOCK hint

# Going Fast

## Storage Engine Options

- Azure SQL / SQL Server support both Column-Store and Row-Store
  - Column-Store: ideal for read-intensive app that scans a lot of data
  - Row-Store: ideal for highly transactional app with a lot of small reads and writes
    - Or sequential (range-scan) reads
  - OLTP and HTAP scenarios 100% supported
    - Technically also DWH scenarios supported but Azure DW is a better choice

DEMO

# Columnstore



# Azure SQL is really NewSQL

## Going Beyond Relational

- Full support to a lot of non-strictly relational features
  - Graph database support
    - EDGE, NODE tables, dedicated constructs (MATCH) and functions  
`(SHORTEST_PATH => Transitive Closure)`
  - JSON Support
    - FOR JSON, OPENJSON, JSON\_VALUE, JSON\_MODIFY
  - Spatial Support
    - Geography and Geometry data types (Open Geospatial Consortium)
    - Full Globe, Polygons, Spatial Indexes, WKT/WKB support, Functions

DEMO

# Graph Support



# A Serverless World

## Azure Integration

- Easy to use Azure SQL from other Azure services
  - With Azure Functions it is easy to create complete serverless solutions
- Azure SQL Serverless SKU
  - Automatically scales based on workload demands
  - Automatic Pause-Resume
    - Resume latency: ~1 minute

# Can Azure SQL DB Do <>?

Some myths are still around...

- Can Azure SQL be used to create really, really, scalable solution?
  - Like loading 1 Billion of rows per day
  - No Bulk-load
    - With bulk load you could do even more
  - 10K inserts per seconds?
  - Keep a scoreboard updated in real-time?

<https://github.com/Azure-Samples/streaming-at-scale/>

DEMO

# Can SQL do...?



# Can Azure SQL DB Do <>?

Some myths are still around...

- Can Azure SQL be used to create real-time solution?
  - Like loading 1 Billion of rows per day?
  - No Bulk-load
    - With bulk load you could do even more?
  - 10K inserts per seconds?
  - Keep a scoreboard updated in real-time?





Azure SQL is the best cloud db for devs  
Right?



# Thank You

## Davide Mauri



@mauridb



davide.mauri@microsoft.com



# Azure SQL Database

Best cloud database for enterprise apps

Davide Mauri, Dimitri Furman  
Senior Program Managers  
Microsoft



# Agenda

1. Robustness and resiliency
2. Scalability and performance
3. Workload monitoring
4. Database design

# Best cloud db for enterprise apps

Azure SQL Database is built on SQL Server

SQL Server:

- 40+ years of R&D
- 25+ years of real-world usage and optimizations
  - Fine-tuned to run critical enterprise applications

Azure SQL: leverages all the enterprise experience of the last 25 years

- It is the only cloud database doing that

# Enterprise Apps

## What is an Enterprise app?

- Enterprise applications are about the display, manipulation, and storage of large amounts of often complex data and the support or automation of business processes with that data
  - Martin Fowler
- Although there is no single, widely accepted list of enterprise software characteristics, they generally include performance, scalability, and robustness.

[https://en.wikipedia.org/wiki/Enterprise\\_software](https://en.wikipedia.org/wiki/Enterprise_software)



# Robustness and Resiliency

In the cloud, be prepared to handle disconnections

- Connection Resiliency
  - Best practice: every mission critical app, right from the start
  - DIY
    - For .NET: Use System.Data.SqlClient >= 4.6.1 or Microsoft.Data.SqlClient
    - [Handle Azure SQL DB Connectivity Issues](#)
    - [Resiliency in Java Apps](#)
  - Or using existing retry framework
    - [Polly](#)

DEMO

# Connection Resiliency



# Scalability and Performance

## Network latency

- Every time your app talks to the database, it is affected by network latency. For chatty apps:
  - Scalability will be harder and more expensive
  - Performance won't be as good as it could be
  - Same best practices can be applied to on-prem
    - (write once deploy everywhere)
- So, Avoid Chatty Apps
  - Push data processing logic to database
  - Avoid sending unnecessary data on the wire (do not use SELECT \*)
  - Send sets of data (arrays, lists) and process them as a whole
    - Table Valued Parameters
    - JSON

# Scalability and Performance

## Storage latency

- Every transaction commit takes 1-10 milliseconds
  - Avoid unnecessarily frequent commits
- Latency depends on storage type

Basic, Standard, General Purpose	remote storage	5-10 milliseconds
Premium, Business Critical	local storage	1-2 milliseconds
Hyperscale	local/remote storage	1-10 milliseconds

- Scaling database up or down
  - Changes IOPS/throughput
  - Does not change latency

# Query Monitoring

Know what's happening or what happened

- What's going on \*right now\*
  - [https://github.com/amachanic/sp\\_whoisactive](https://github.com/amachanic/sp_whoisactive)
- DMV
  - Dynamic Management Views
- Query Store
  - Query performance flight recorder
- Extended Events
  - X-Rays into Azure SQL

# Import external data

## What if you need to import CSV or JSON data?

- Azure SQL can access Azure blob storage
  - But \*not\* files on user local disk or on the network
  - OPENROWSET, BULK IMPORT (including BCP native/character formats)
  - An easy way is to integrate Azure SQL with Azure Functions
    - <https://medium.com/@mauridb/automatic-import-of-csv-data-using-azure-functions-and-azure-sql-63e1070963cf>

DEMO

# Importing JSON/CSV



# Database design for enterprise apps

## Azure SQL Database vs. SQL Server: What is different

### Instance-level features

- Cross-database joins
- SQL Agent
- Linked servers
- Resource Governor
- Database Mail
- ...

### Authentication

### Database size limit

### Database options and features

- Recovery model
- Compatibility level
- Files/filegroups
- RCSI
- Auto-tuning
- MAXDOP
- Latest SQL Server features

# Database design for enterprise apps

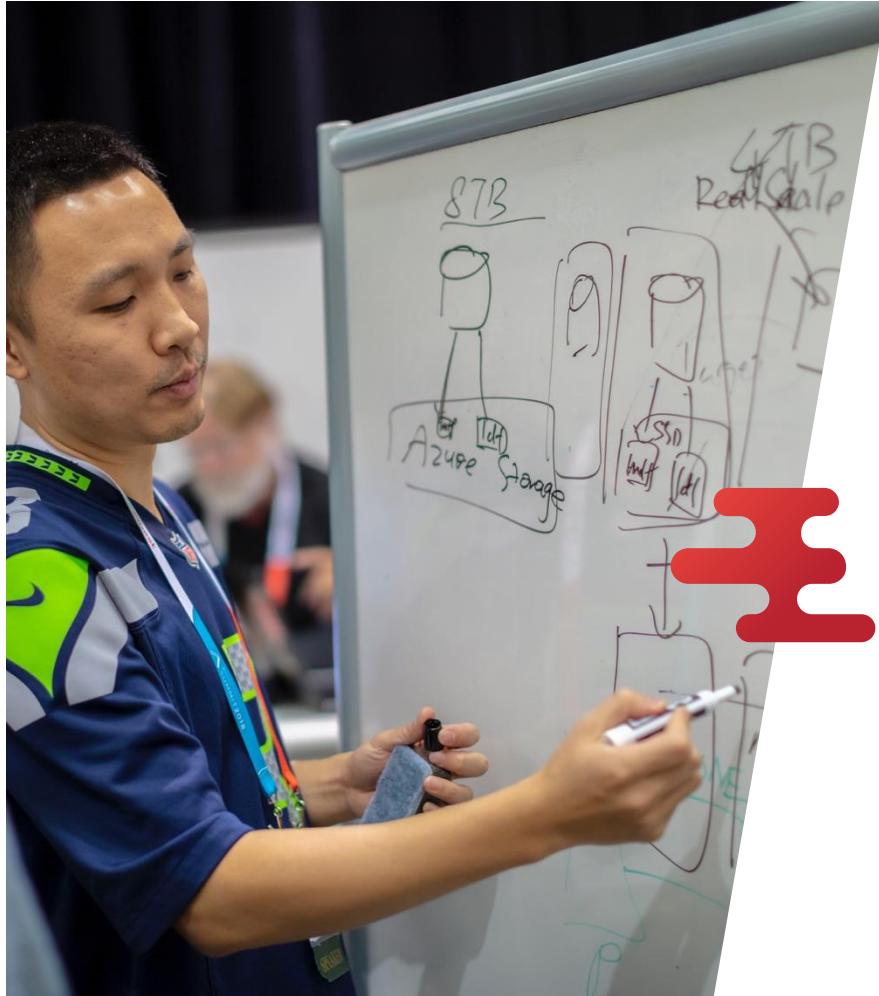
## Azure SQL Database vs. SQL Server: What does not change

Database design best practices

SQL Server database engine

Performance tuning methods

- Indexing
- Query rewrite
- Statistics and query plans
- Locking, blocking, deadlocks
- Hints



# Thank You

Dimitri Furman,  
Davide Mauri





# Azure SQL Database Hyperscale

Denzil Ribeiro  
Principal Program Manager, Microsoft



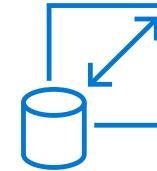
# Challenges with Large Databases

## Size of data



- Provisioning more storage to expand the database can be painful
- Operations take a LONG time (days in some cases)
- Ongoing operations degrade database performance
- Even in cloud environments, need to shard to fit in a tier

## Scaling Compute



- Logistics of moving to larger box
- Economics of sizing for max peaks

# Azure SQL Database Hyperscale?



Cloud native

Architected for cloud



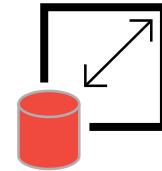
Storage

Scalable new tiered storage architecture



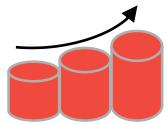
Seamless compatibility

Fully compatible with Azure SQL Database



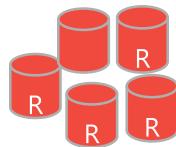
No limits

Scale compute and storage separately  
Scale up in constant time.



Large database

Support for 100TB+  
Infinite Log



Read scale with replicas

Supports 4 Read replicas

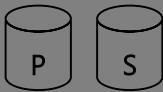


Performance

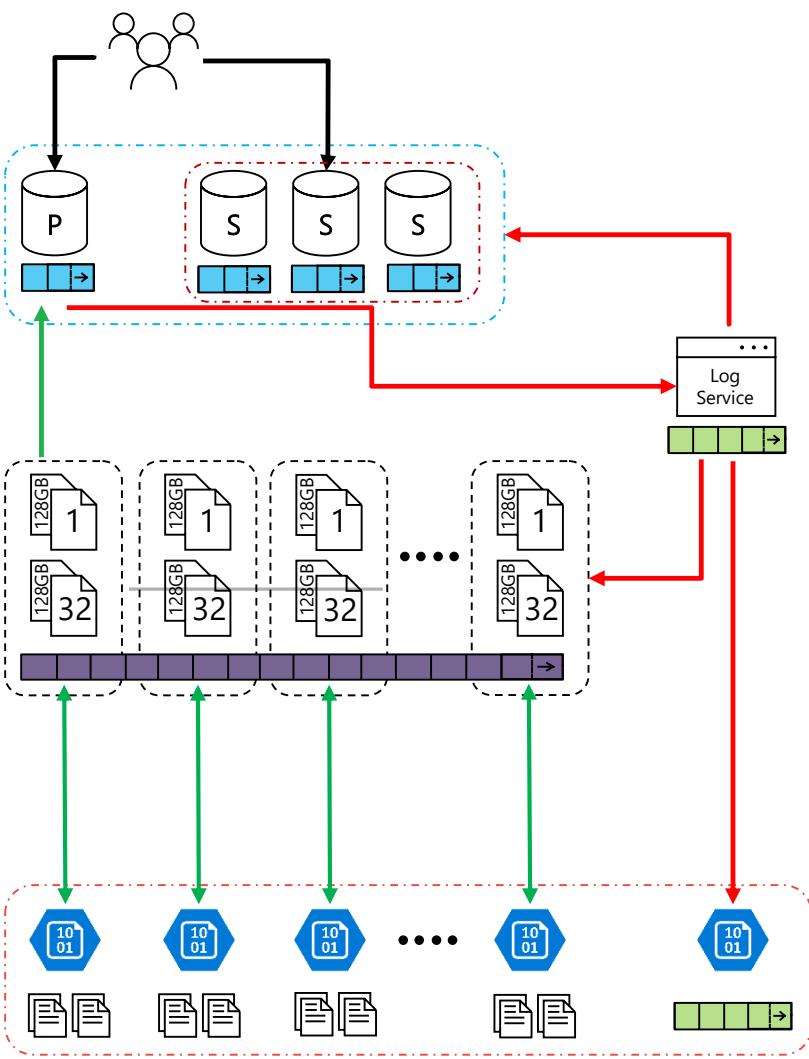
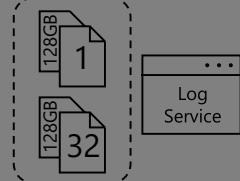
Vldb operations without headaches

- Restore in constant time
- No impact of backups on compute
- No impact of checkpoints on compute
- Accelerated database recovery

## Compute

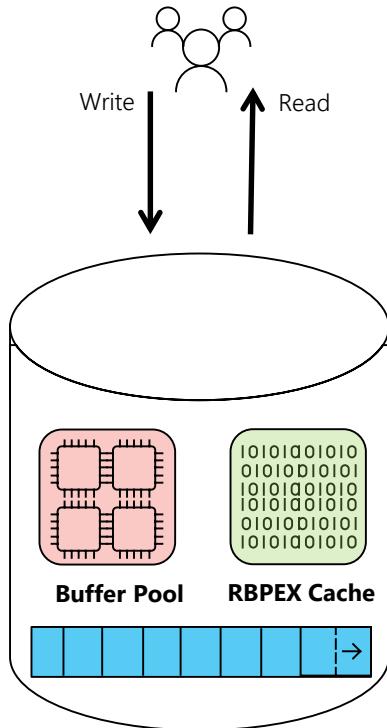


## Storage



- Features that make this possible:
  - Resilient Buffer Pool Extension (RBPEx)
  - Snapshot-based backups
  - Accelerated Database Recovery (ADR)
- External log service
- Tiered storage
- Completely separated compute and storage
- No single point of failure
- Storage encrypted at rest

# Primary

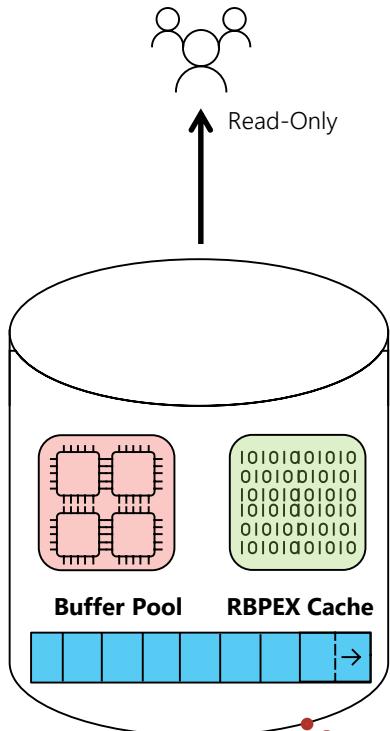


- Configurable vCores 2 – 80
- Memory and directly attached SSD proportional to number of cores
  - Memory dependent on hardware generation
  - **Non-covering** RBPEX
  - 1-2 ms data access latency
- Read-Write Workloads
- Seamless end user interactions

Compute



# Readable Secondary



- Currently same size as the Primary
- Up to 4 readable secondary replicas
- Read-only workloads
  - ApplicationIntent = READONLY
  - Round robin distribution
- Hot failover target

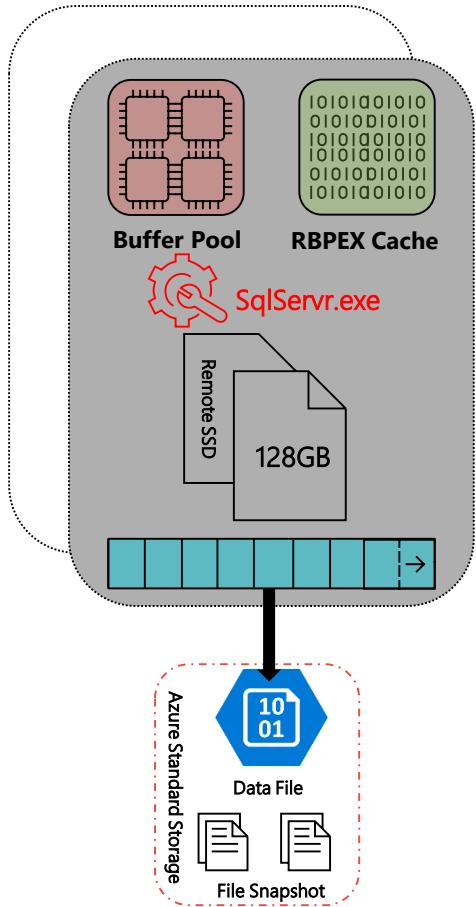
Replica priced at  
AHUB pricing

Higher SLA with >  
0 replicas

Compute



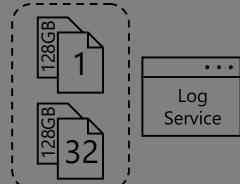
# Page Server



- “Remote” SSD - 128GB per Page Server
- **Covering** RBPEX
  - $\geq 2$ ms data access latency
- Each page server has its data file on Azure Standard storage
- Offload operations from compute
  - Checkpoints executed on Page Servers
  - File snapshots for backup operations are isolated from compute
- Built-in high availability for page servers and storage
- Allocations are 10 GB increments

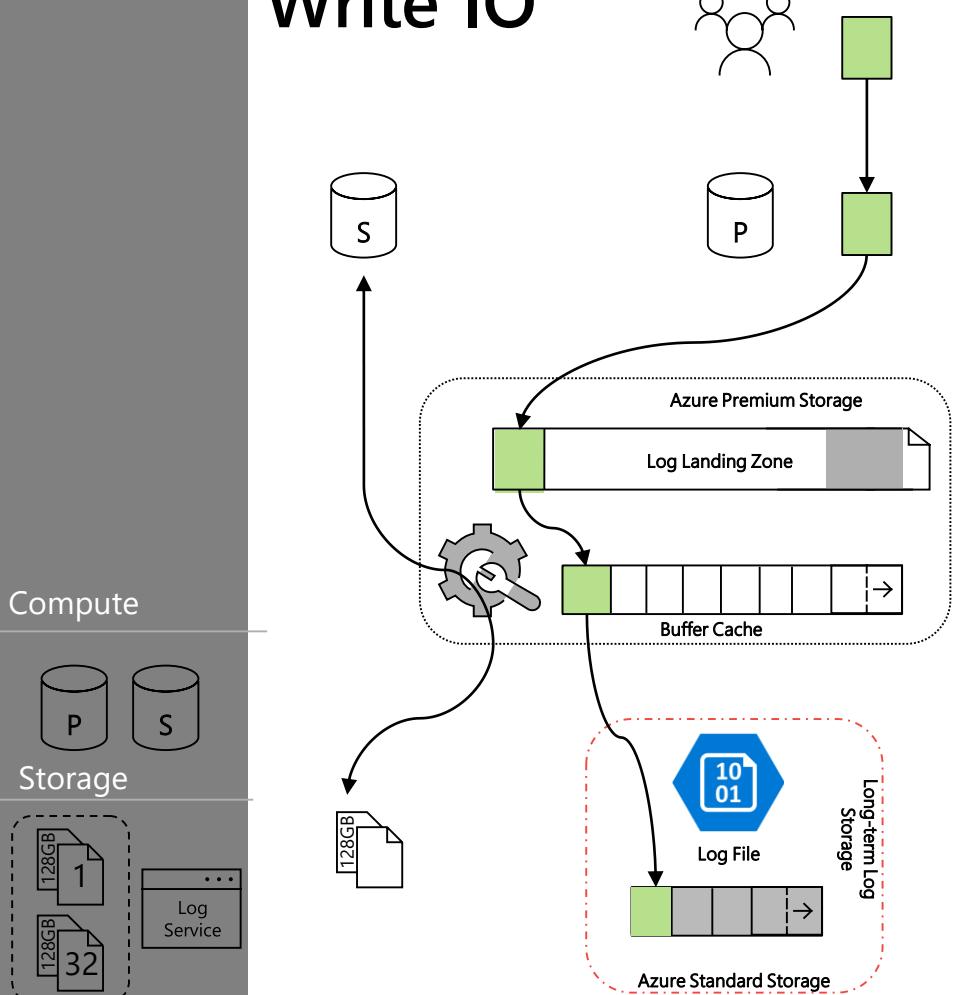
# Log Service

Compute  
Storage



- Log Service consists of 3 components
  - Landing Zone
  - The actual Log Service
  - Long-Term Log Storage
- Availability and Resiliency built-in for Storage
- New Transactions are hardened to the Log Landing Zone (Tx Commit)
- Log broker/ Cache will stage Log blocks
  - Apply Log blocks to secondaries / Page servers.
- Data is offloaded to Long-Term log storage

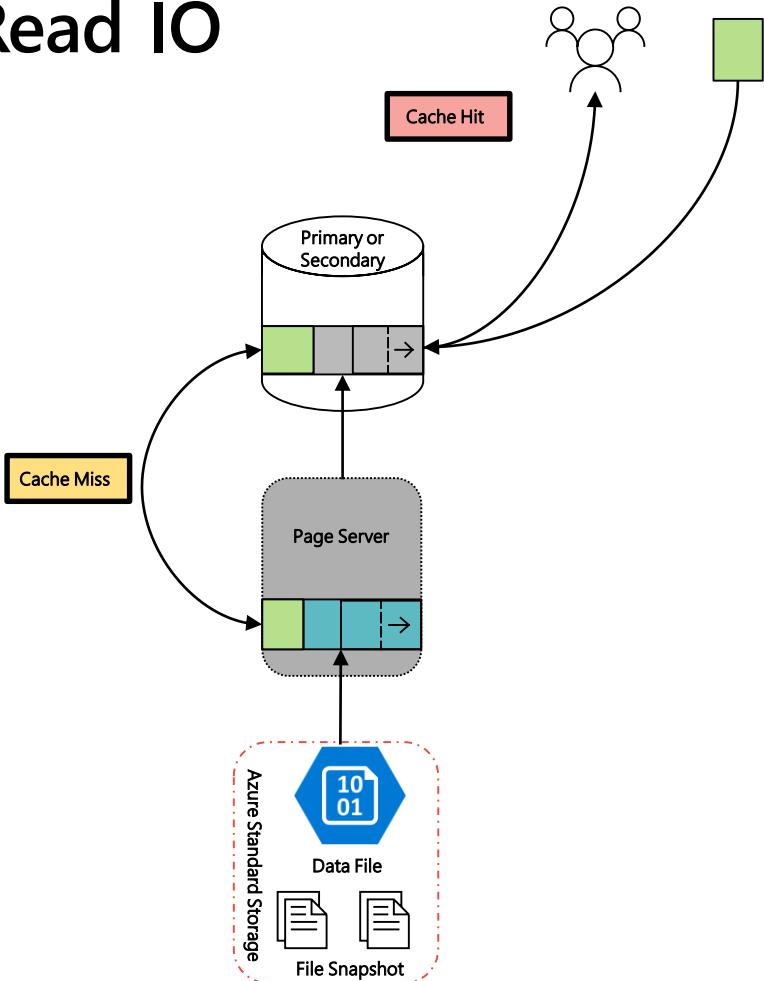
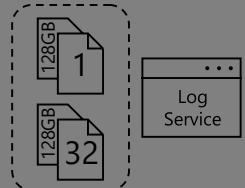
# Write IO



- A write Tx is generated by a user or application to the primary
- Transaction is written to the log landing zone
  - **Tx Commit** on hardening of Tx to Log
- Log service stages the new Tx in its cache
- Log service will asynchronously and simultaneously apply the Tx to the secondaries and the page servers
- Tx Log is destaged from the buffer cache and shipped to the long-term log storage

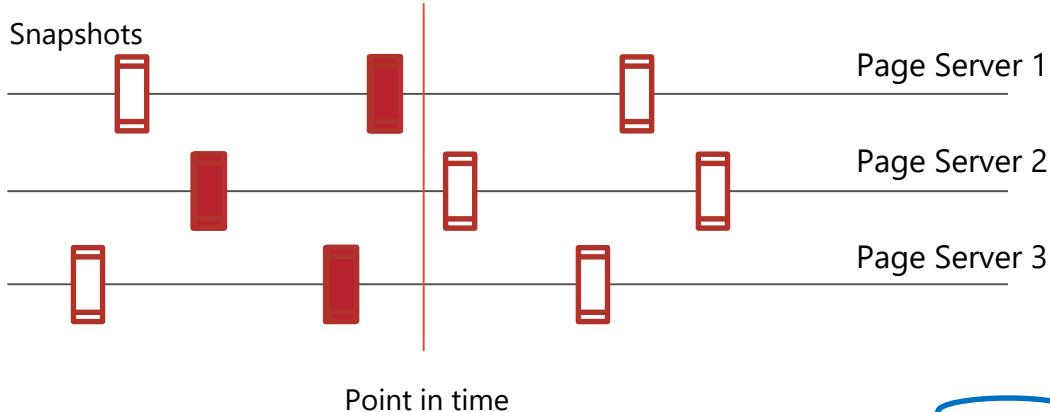
# Read IO

Compute  
Storage



- From a Cold Start – Data Files stored on Azure Standard Storage
- Data files will populate the covering RBPEX on the Page Servers
- Page servers will populate the non-covering RBPEX on the Primary and Secondary compute
- Query to return data pages is issued by a user or an application
- Cache Hit – Data is in the local RBPEX 1-2ms latency
- Cache Miss – Data has not on local RBPEX, to be retrieved from Page Server > 2ms data latency

# Backup & Restore

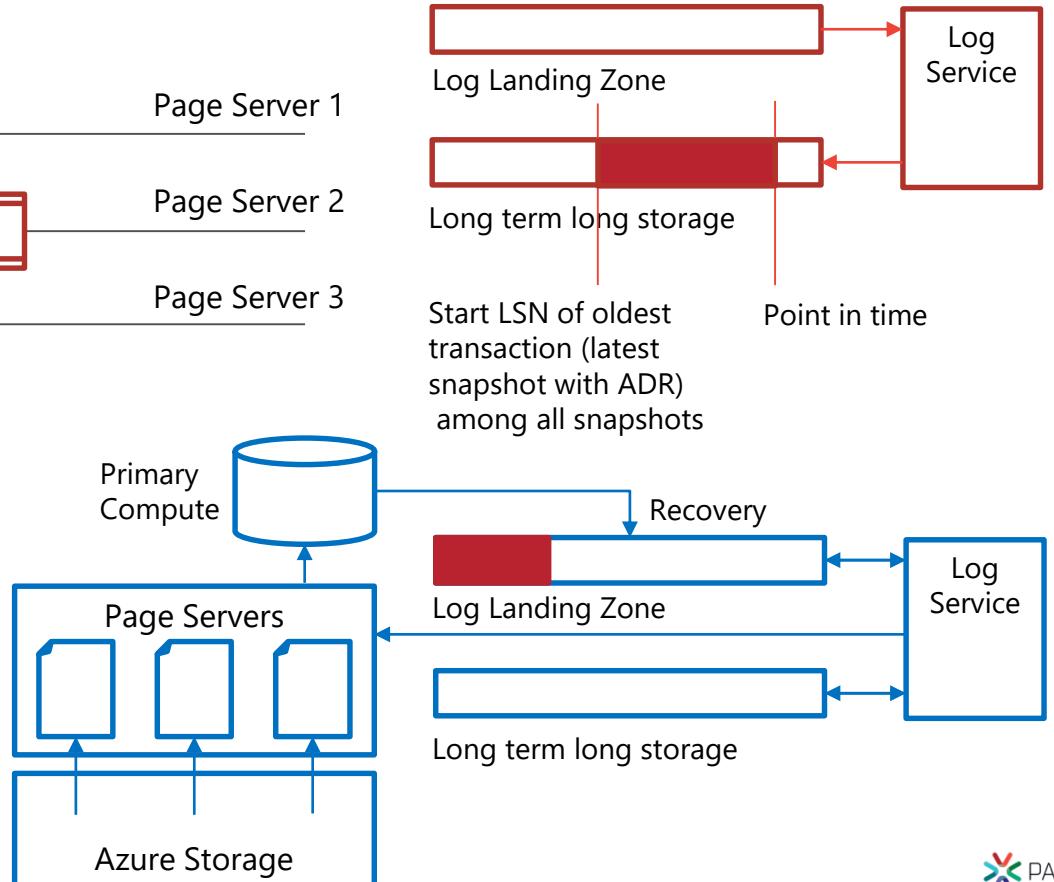


Backup has no impact to apps

Parallelized copy (metadata)

Constant time PITR

Accelerated database recovery



DEMO

# Read Scale



# Hyperscale benefits

Performance

- Better performance than Standard/General Purpose overall, for **all** databases
- Offloading backups/checkpoints to page servers
- Rapid scale up/down
- Read scale out with Primary + 4 replicas
- Snapshot based restores (fast for any database size)
- Workloads requiring very low latency commit may work better on Premium/Business Critical today

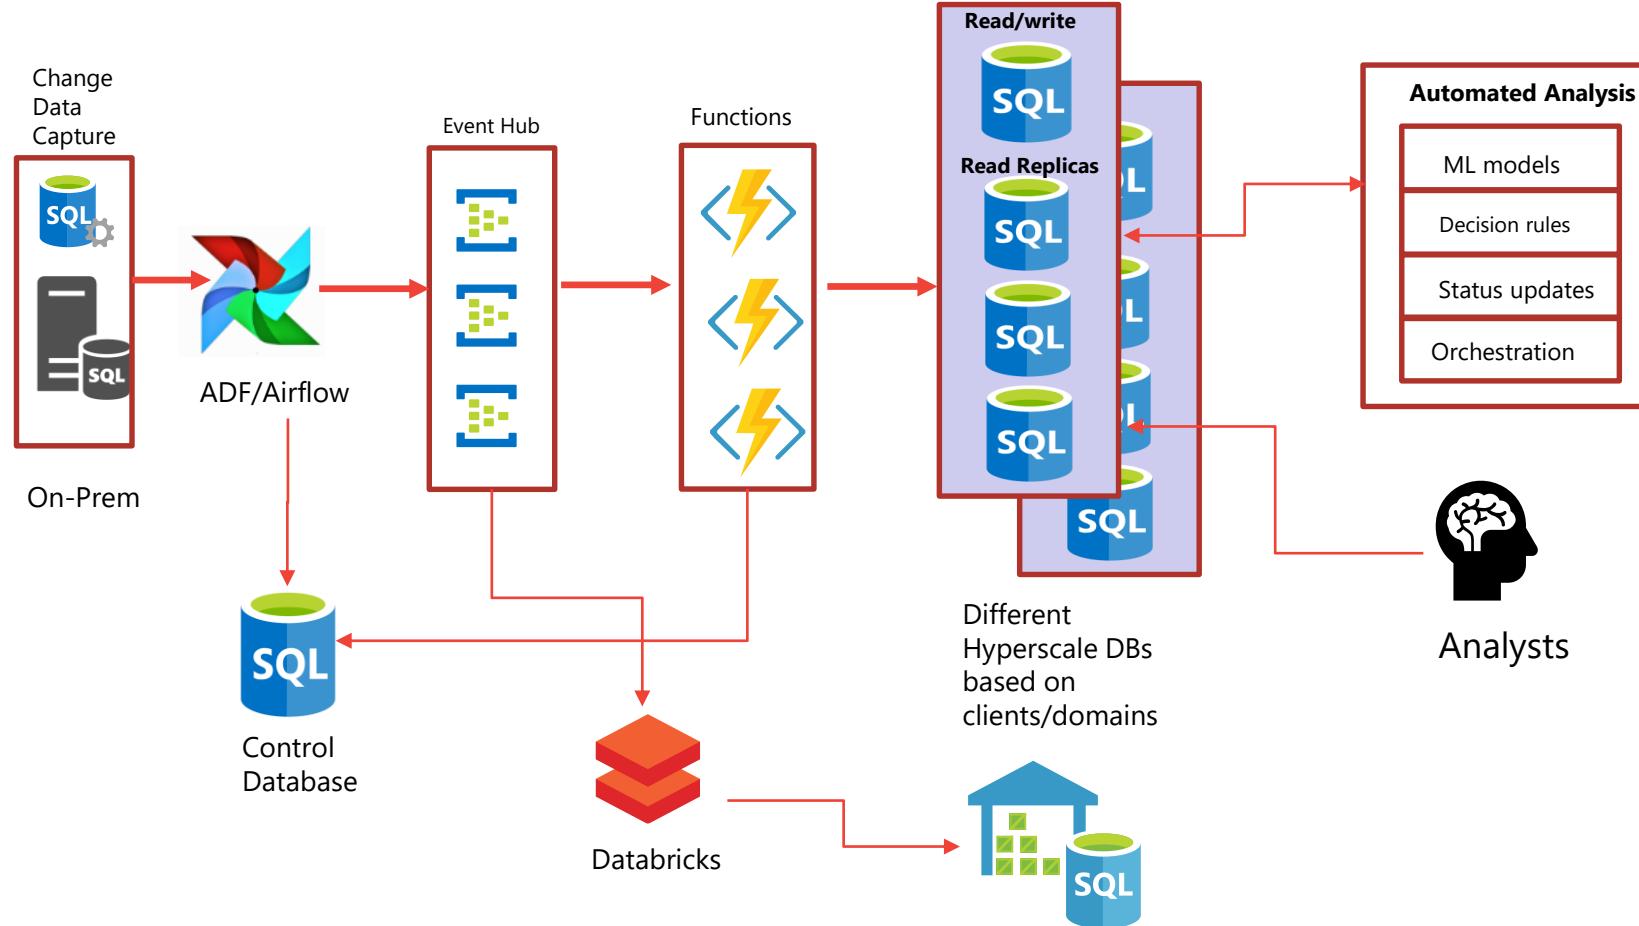
Size

- Large databases up to 100 TB
- Auto-scaling size
- Storage decoupled from compute

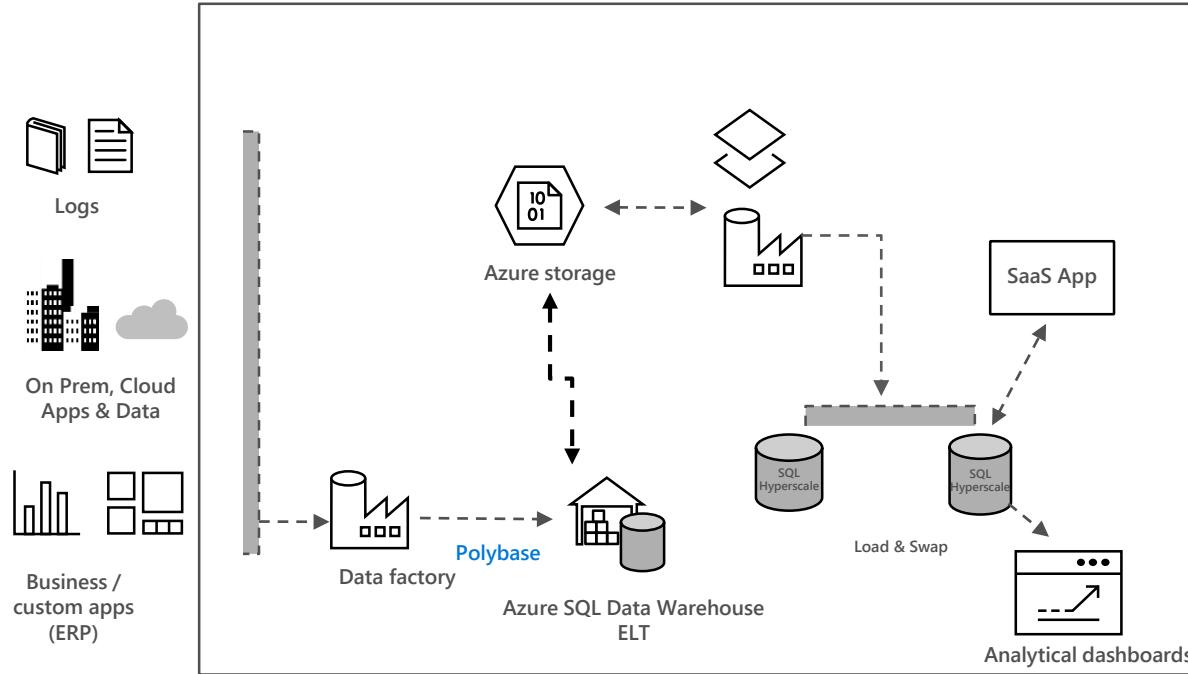
Workload Patterns

- OLTP
- HTAP
- Operational data stores/marts
- IOT
- PRS migrations

# Scenario: Near-Realtime Fraud Detection



# Scenario: Hyperscale for Data Marts



Several customers use Hhyperscale as ODS/data mart for reporting.

Patterns:

- Load and swap (2 separate HS database), allows different tiers on each.
- Batch/ETL loads with ADF or Spark.
- Database sizes in 5-20 TB range currently

# Migration – Moving Data to Hyperscale

## [Azure Data Migration Service](#)

### Bulk load/Bulk copy

- Using [Azure Data Factory](#)
- Using [Spark connector for SQL](#)
- Using Sqlbulkcopy Sample: [smart bulkcopy](#)
- Bulk load into heaps for larger tables, create indexes later
- If using clustered Columnstore, bulk load directly into the table with CCI

### Transactional replication

- Use concurrent snapshots
- Requires primary keys

# Migration from existing Azure SQL DB

Migration from another service tier is just a service objective change

\*Currently\* cannot go back to prior tier

Existing DBs moved to Hyperscale

- Existing files > 128 GB, land on 1 TB page servers
- Existing files <=128 GB land on 128 GB page servers (more efficient)

If performance is an issue in testing

**Option 1:**

Create new Hyperscale DB, Load data.

**Option 2:**

- Add files at the source, Rebuild indexes to redistribute data, shrink
- Once files are < 128GB each then SLO update

# Hyperscale Perf diagnostics

Log throttling has new set of [waits](#) (RBIO\_RG\*)

Page Server reads (Remote reads) reported in all the DMVs below

- [sys.dm\\_exec\\_requests](#)
- [sys.dm\\_exec\\_query\\_stats](#)
- [sys.dm\\_exec\\_procedure\\_stats](#)

Page Server reads added to Actual Execution Plan to distinguish remote reads.

RPBEX IO exposed in DMVS

- FileID 0 in [sys.dm\\_io\\_virtual\\_file\\_stats\(\)](#) accounts for RBPEX IO.
- All RBPEX IO is 8KB

Additional details: [Hyperscale performance diagnostics](#)

# Index Creation/Rebuild

- Tempdb is still a finite resource, size depends on compute size
- Consider resumable indexes for index creation/rebuild on very large tables
  - Persists index build state which is replicated, can be resumed.
  - CREATE INDEX is also resumable
  - New database scoped options:  
`ELEVATE_RESUMABLE= { OFF | WHEN_SUPPORTED | FAIL_UNSUPPORTED }`  
`ELEVATE_ONLINE = { OFF | WHEN_SUPPORTED | FAIL_UNSUPPORTED }`
- Consider MAXDOP hint for index create/rebuild
  - Particularly if doing concurrent index create/rebuild
- Consider partition level index rebuilds

\*Offline indexing is faster than online/resumable



# Thank You

## Denzil Ribeiro



@denzilribeiro



denzilr@microsoft.com



# Azure SQL for DBAs

How it Works, and How to Solve  
Common Problems

Dimitri Furman, Denzil Ribeiro  
Microsoft



# Agenda

1. DBA tasks in the Azure SQL world
2. Solving common performance and scale problems

# Traditional DBA Tasks

## Platform DBA

- High availability
- Perf tuning (resources)
- Backup/restore
- Business continuity/DR
- Monitoring
- Infrastructure design
- Installation, upgrades, patching

## Application DBA

- Database design and development
- Code deployment
- Perf tuning (query)
- Data lifecycle management
- Data security

How does this change in Azure SQL Database?

# High Availability

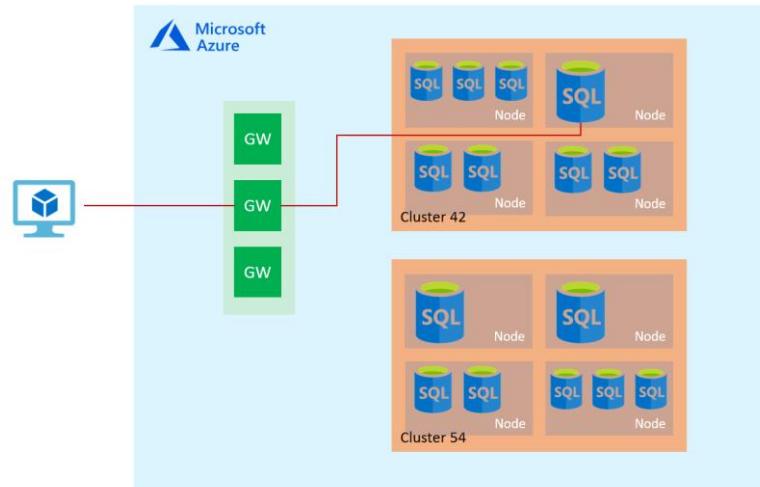
## Built-in, for any Azure SQL database

- Understand HA capabilities
- Which vary by service tier
  - Basic, Standard, General Purpose
  - Premium, Business Critical
  - Hyperscale

# High Availability

## How it works: all service tiers

- Highly redundant connectivity layer, multiple gateways
- Large pools of compute capacity (tenant rings), hundreds of nodes
- Managed by Service Fabric



# High Availability

## How it works: Basic, Standard, General Purpose

- Single compute replica (VM) running SQL Server
- Data is remote in Azure Storage
  - 99.99999999% durability guarantee over a year
- Standby VMs replace failed replicas
- Fast compute scaling in minutes, regardless of database size
  - Storage is re-attached to new replica
  - Failover takes up to a minute

# High Availability

## How it works: Premium, Business Critical

- Multiple compute replicas, typically 4
- Availability Groups with quorum commit
- Data is local on SSD storage
- Failed replicas are replaced from compute pool and reseeded
- Time to scale compute is proportional to database size
  - New replicas are built, data is copied at 1 GB/minute
  - Failover to a sync commit secondary takes <8 seconds on average
- Zone-redundant databases (optional)

DEMO

# Manual failovers in Azure SQL Database

Invoke-AzSqlDatabaseFailover  
Invoke-AzSqlElasticPoolFailover



# Resource Management

## Current hardware generations in Azure SQL

### Gen5

- Intel Broadwell CPU (E5-2673 v4)
- 5.1 GB/vCore memory ratio
- Up to 80 vCores (hyperthreaded)
- NVMe SSD
- Up to 4 TB of local storage
- Accelerated Networking is guaranteed

### Gen4

- Intel Haswell CPU (E5-2673 v3)
- 7 GB/vCore memory ratio
- Up to 24 vCores (physical)
- SAS SSD
- Up to 1 TB of local storage
- Accelerated Networking is not guaranteed

## New hardware generations in preview

### M128m

- 128 vCores (hyperthreaded)
- 3.8 TB memory
- Business Critical

### Fsv2

- 72 vCores (hyperthreaded)
- 3.4 GHz sustained turbo clock speed
- General Purpose

# Resource Management

## System resource governance in Azure SQL

- Database service objective: a set of resource limits
  - CPU, memory, storage size, IOPS, throughput, sessions, worker threads, etc.
  - [Single database limits](#) and [Elastic pool limits](#)
- Provides balanced Database-as-a-Service
  - Backup, HA, DR, predictable performance, multi-tenant fairness, etc.
- Investigating performance? Check resource usage first

DEMO

# Troubleshooting “Request limit reached”



# Backup and Restore

## Built-in, for any Azure SQL database

- Automatic full, differential, and log backups
- Configurable backup retention period, up to 35 days\*
- Optional long-term backup retention (LTR), up to 10 years\*
- Automated point-in-time restore
  - Restore time depends on service objective
- Fast backup and restore for VLDBs in Hyperscale

- \* Temporary limitations for Hyperscale
  - Fixed 7-day retention only
  - LTR is not yet supported

# Business Continuity/DR

## Built-in, for any Azure SQL database

- Geo-Restore
  - Backups copied to paired Azure region
  - Can be restored in any region

## Opt-in, for any Azure SQL database

- Geo-Replication
  - Database replicas in any Azure region
- Failover Groups
  - Geo-replication for a group of databases
  - Connection strings not affected by failovers

# Monitoring and alerting

## Built-in, for any Azure SQL database

- Resource metrics in Azure Monitor
- Query Store
  - Automatic tuning
  - Query performance insights
- Resource health
- DMVs

sys.dm_db_resource_stats
sys.resource_stats (in master)
sys.dm_db_wait_stats
sys.dm_os_wait_stats
sys.dm_exec_session_wait_stats
sys.dm_os_performance_counters
<b>sys.dm_os_job_object</b>
sys.dm_user_db_resource_governance
sys.dm_resource_governor_resource_pools
sys.dm_resource_governor_resource_pools_history_ex
sys.dm_resource_governor_workload_groups
<b>sys.dm_resource_governor_workload_groups_history_ex</b>

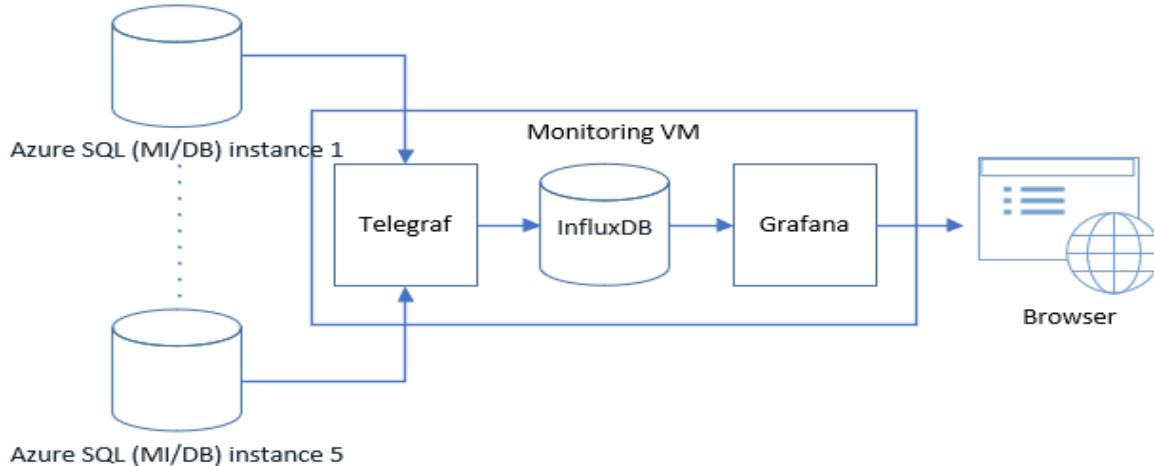
# Monitoring and alerting

## Opt-in, for any Azure SQL database

- Azure SQL Analytics, Intelligent Insights
  - Early detection, root cause analysis
  - Monitoring at scale
- Extended Events for in-depth troubleshooting
- Alert rules in Azure Monitor
  - Define an alert for multiple databases/managed instances
  - Alert on resource metrics exceeding thresholds
  - Alert on management activity events

# Additional monitoring options

- Third party tools
- [Query performance insight library](#)
- [Near-realtime monitoring for Azure SQL](#) with [telegraf](#)/InfluxDB/Grafana
  - Ability to monitor [Managed Instance](#), Azure SQL DB, SQL Server





# Solving performance and scale problems



# Data loading in Azure SQL Database

All service tiers have necessary resource limits

Cost/performance and workload specific decisions

Hyperscale: same log generation rate in any service objective

Full recovery model!

- On Heaps still use TABLOCK (#locks acquired)
- On Columnstore tables use batch size > 102,400

Partitioning very large tables is important

- Provides flexibility on index rebuilds, incremental stats

Business Critical has HADR replicas

# Log rate governance

## Why throttle log generation?

- Required for replicas to be up to date
- Allows frequent log backups to sustain recoverability SLAs

Applied to log generation itself to avoid storage throttling

Surfaced as [waits](#)

Wait Type	Notes
LOG_RATE_GOVERNOR	Database <a href="#">limiting</a>
POOL_LOG_RATE_GOVERNOR	Pool <a href="#">limiting</a>
INSTANCE_LOG_RATE_GOVERNOR	Instance level <a href="#">limiting</a>
HADR_THROTTLE_LOG_RATE_SEND_RECV_QUEUE_SIZE	Feedback control, availability group physical replication in Premium/Business Critical not keeping up
HADR_THROTTLE_LOG_RATE_LOG_SIZE	Feedback control, limiting rates to avoid an out of log space condition

DEMO

# Data loading in Azure SQL DB



# Other data loading tips

Load data in parallel improves throughput

- Reduce concurrent parallel loads on columnstore due to memory grants

Loading from a compute (ADF/Spark) offloads the “read” of the files to a compute engine rather than SQL DB

If loading by partition

- Load into staging table
- Build aligned indexes
- Switch into partitioned table

When comparing to a baseline, compare with Full recovery model

Load into Columnstore when feasible (> 102400 batch size)

For very large data warehouse loads ( >150 MB/sec), Azure SQL DW would be a better choice.

# Storage

## Local storage

General Purpose: used for tempdb only

Business Critical: used for *all* database files

## Remote storage

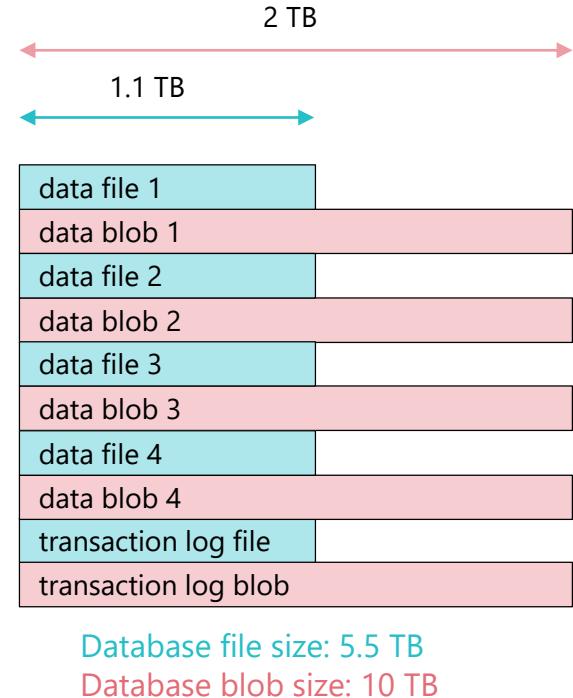
Database files are page blobs in Azure Premium Storage

IOPS/throughput is tied to blob size

5-10 ms latency, unless throttled

Blob size  $\geq$  database file size

Managed Instance storage limit is based on file size, not blob size



## Premium storage limits

DEMO

# Maximizing storage perf



# Table variables vs. Temp Tables

Functionality	Table variables	Temp Tables
Transactional behavior	Not transactional	Transactional
Indexes	Only in-line definition	Yes
Constraints	Only PK, Unique, Check	Yes
Auto Stats	No	Yes
Manual Stats	No	Yes
Compilation if used in same batch	TV compiled with the batch, estimate always 1	Deferred until first execution of statement using temp table

Compile deferred till first statement  
Compatibility LEVEL = 150

D E M O

# Troubleshooting high CPU due to bad estimates



# Networking perf recommendations

Use [redirect](#) mode

Collocate app & database in same region

Use [Accelerated Networking](#) on App VM when possible

Networking tools: [psping](#), [Latte for Windows](#), [Sockperf for Linux](#)

If possible, [batch data](#) together

- Using TVPs
- Using multi-row parameterized statements
- Using SQLBulkCopy

Where applicable use SET NOCOUNT ON

Recent versions of drivers and tools are recommended

D E M O

# Chatty applications





# Thank You



Dimitri Furman, Denzil Ribeiro  
Microsoft



# Azure SQL Database

Intelligent Database

Joe Sack  
Program Manager



# Intelligent Database landscape

Azure SQL Database full intelligence surfaced across:

## SQL Advanced Threat Protection

- Data Classification
- Vulnerability Assessment
- Threat Detection

Maximizing app performance...

# Maximizing app performance

## Monitor

- DMVs, Query Store, XEvents
- Query Performance Insight
- Azure SQL Analytics

## Understand

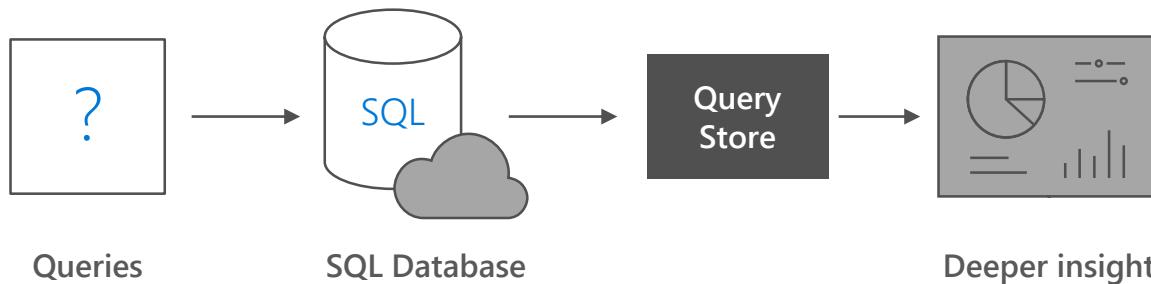
- Intelligent Insights
- Performance recommendations

## Tune

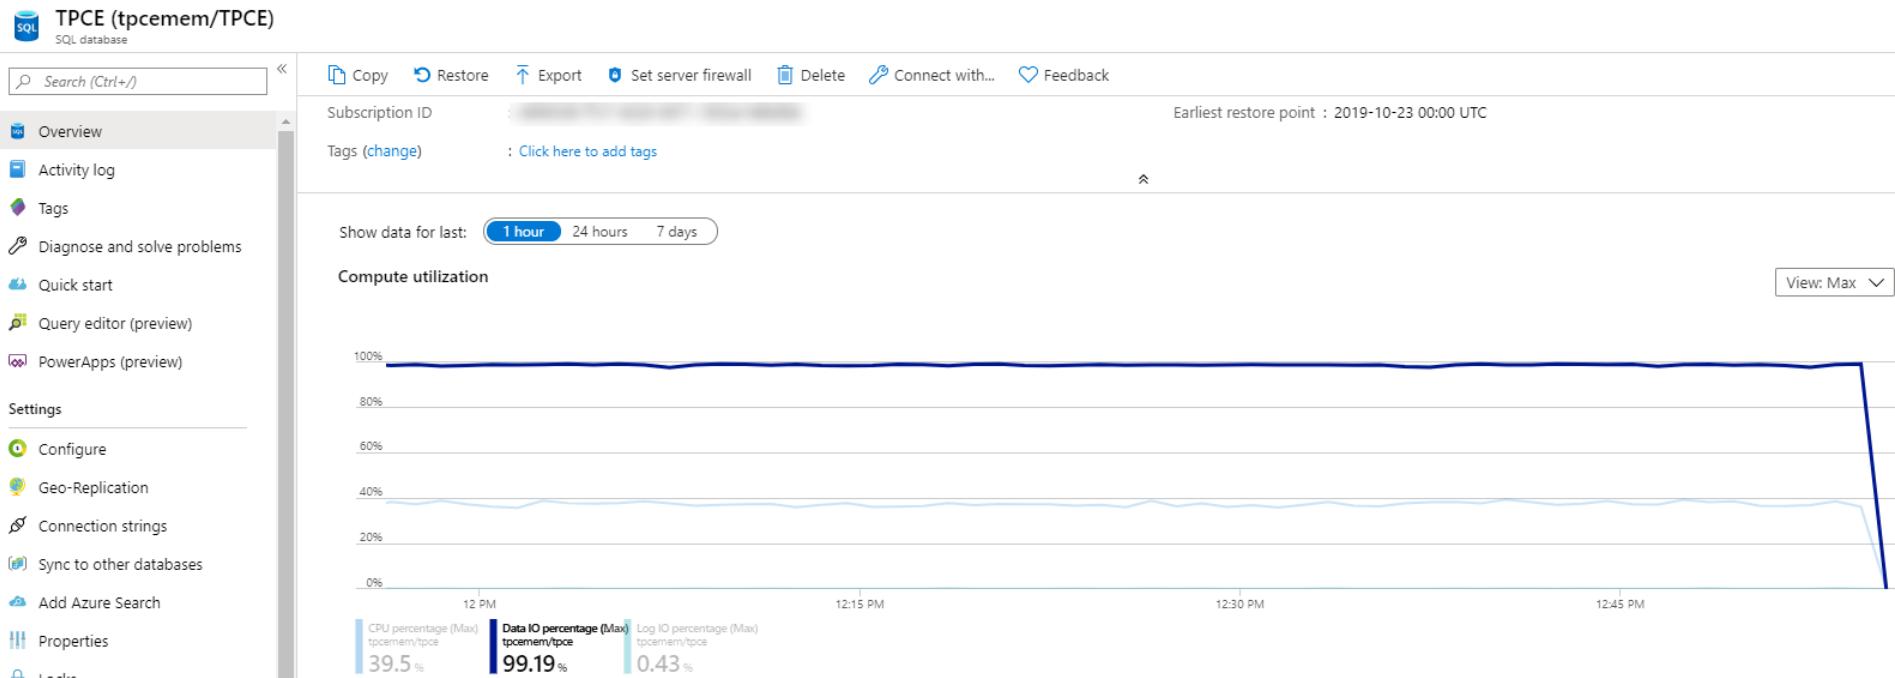
- Automatic tuning
- Intelligent QP

# Crucial default: Query Store

- “Flight data recorder” that simplifies query performance troubleshooting
- On *by default* in Azure SQL Database
- Foundation for performance analysis and tuning features



# Portal: Overview, Compute utilization



# Portal: Overview Notifications

Notifications (1) Database features (7)

All Alerts (0) Recommendations (1) Info (0)

 Database DTU consumption is high.  
Click here to review top queries.

Takes you to Query Performance Insight

# Query Performance Insight

TPCE (tpcemem/TPCE) - Query Performance Insight

Search (Ctrl+ /)

Reset settings Refresh Recommendations Getting started Feedback

Resource consuming queries Long running queries Custom

TOP 5 queries by: CPU Data IO Log IO

Aggregation type: SUM

Time period: LAST 24 HRS

Metrics comp... CPU FOR 46328 0 % CPU FOR 46339 0 % CPU FOR 46311 0 % CPU FOR 46308 0 % CPU FOR 46321 0 %

CPU DATA IO LOG IO

Click on a row below to get the details for the selected query.

QUERY ID	CPU[%]	DATA IO[%]	LOG IO[%]	DURATION[hh:mm:ss]	EXECUTIONS COUNT	#
46321	15.41	62.26	0.01	47:42:15.209	575545	<input checked="" type="checkbox"/>
46308	1.96	6.78	0	09:26:09.0	327394	<input checked="" type="checkbox"/>
46311	0.56	0.98	0	29:16:30.289	197040	<input checked="" type="checkbox"/>
46339	0.43	0.23	0	08:19:52.159	393774	<input checked="" type="checkbox"/>
46328	0.34	0	0	00:10:44.409	327394	<input checked="" type="checkbox"/>

PASS

# Query Performance Insight, Query Details

Home > tpcemem > TPCE (tpcemem/TPCE) - Query Performance Insight > Query details

## Query details

TPCE - Query ID 46321



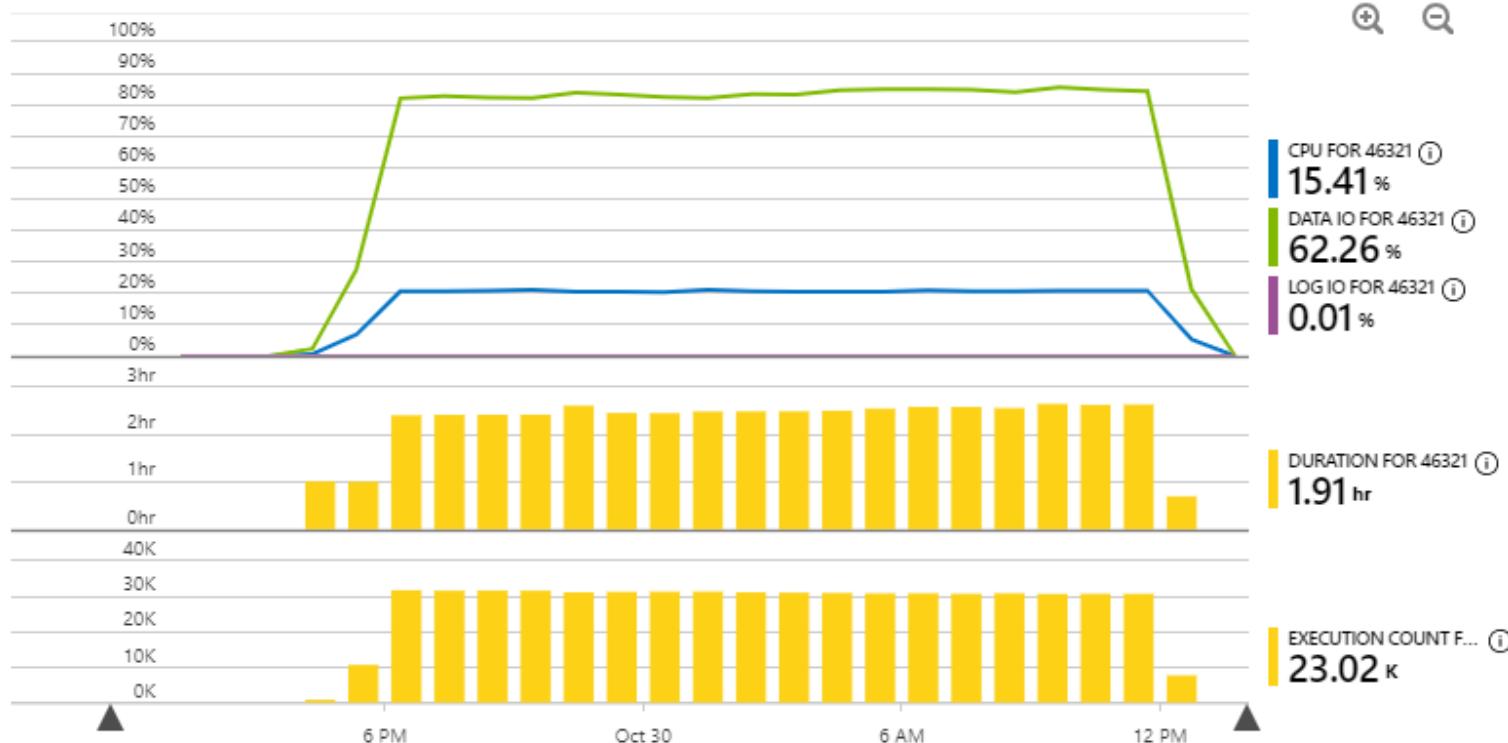
Settings Refresh Recommendations Query Text

**Query ID 46321:**

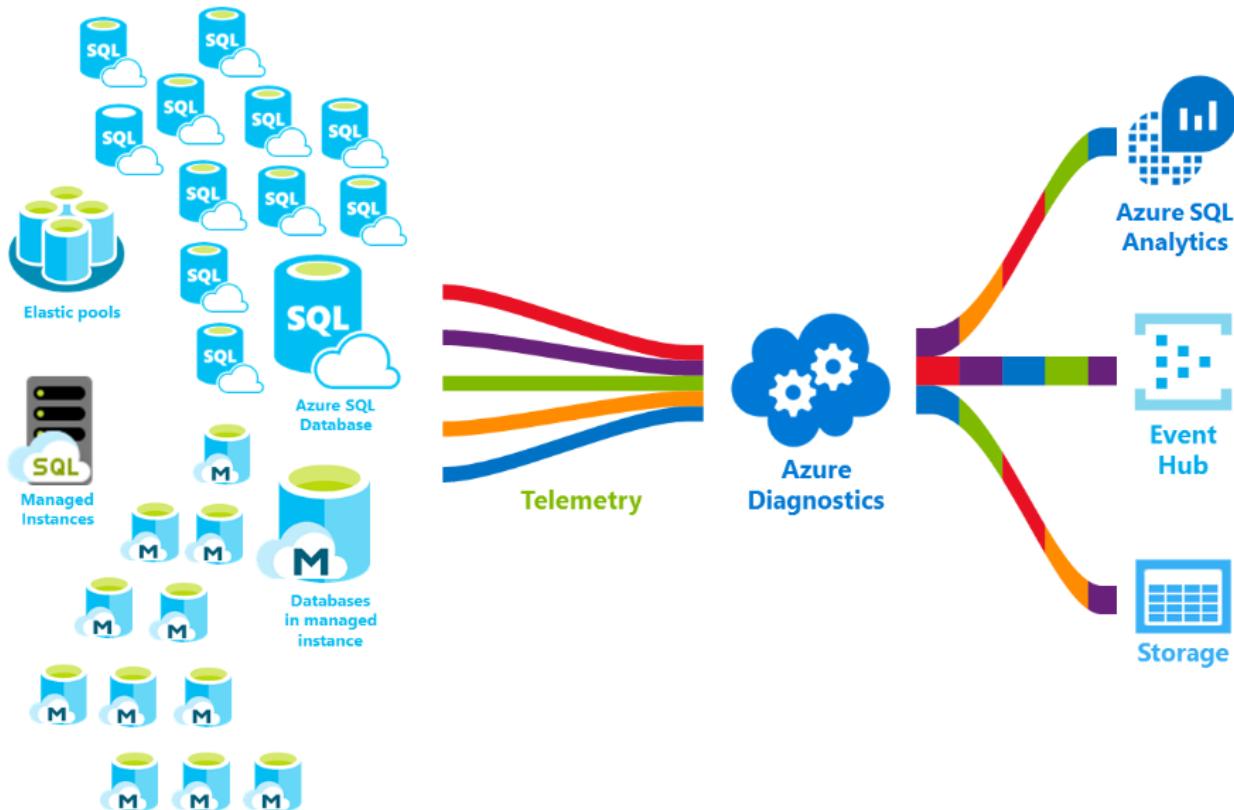
```
1 (@acct_id bigint)SELECT TOP 50
2           T_CHRG,
3           T_EXEC_NAME,
4           EX_NAME,
5           S_NAME,
6           ST_NAME,
7           T_S_SYMB,
8           T_DTS,
9           T_ID,
10          T_QTY,
```

# Query Performance Insight, Query Details

Details of Query ID 46321 (Aggregation type: sum) Last 24 hrs



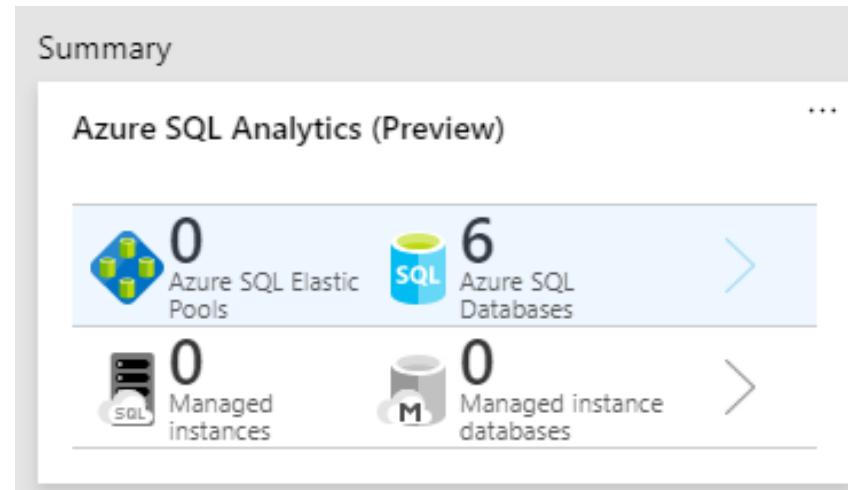
# Monitoring and Analysis at Scale



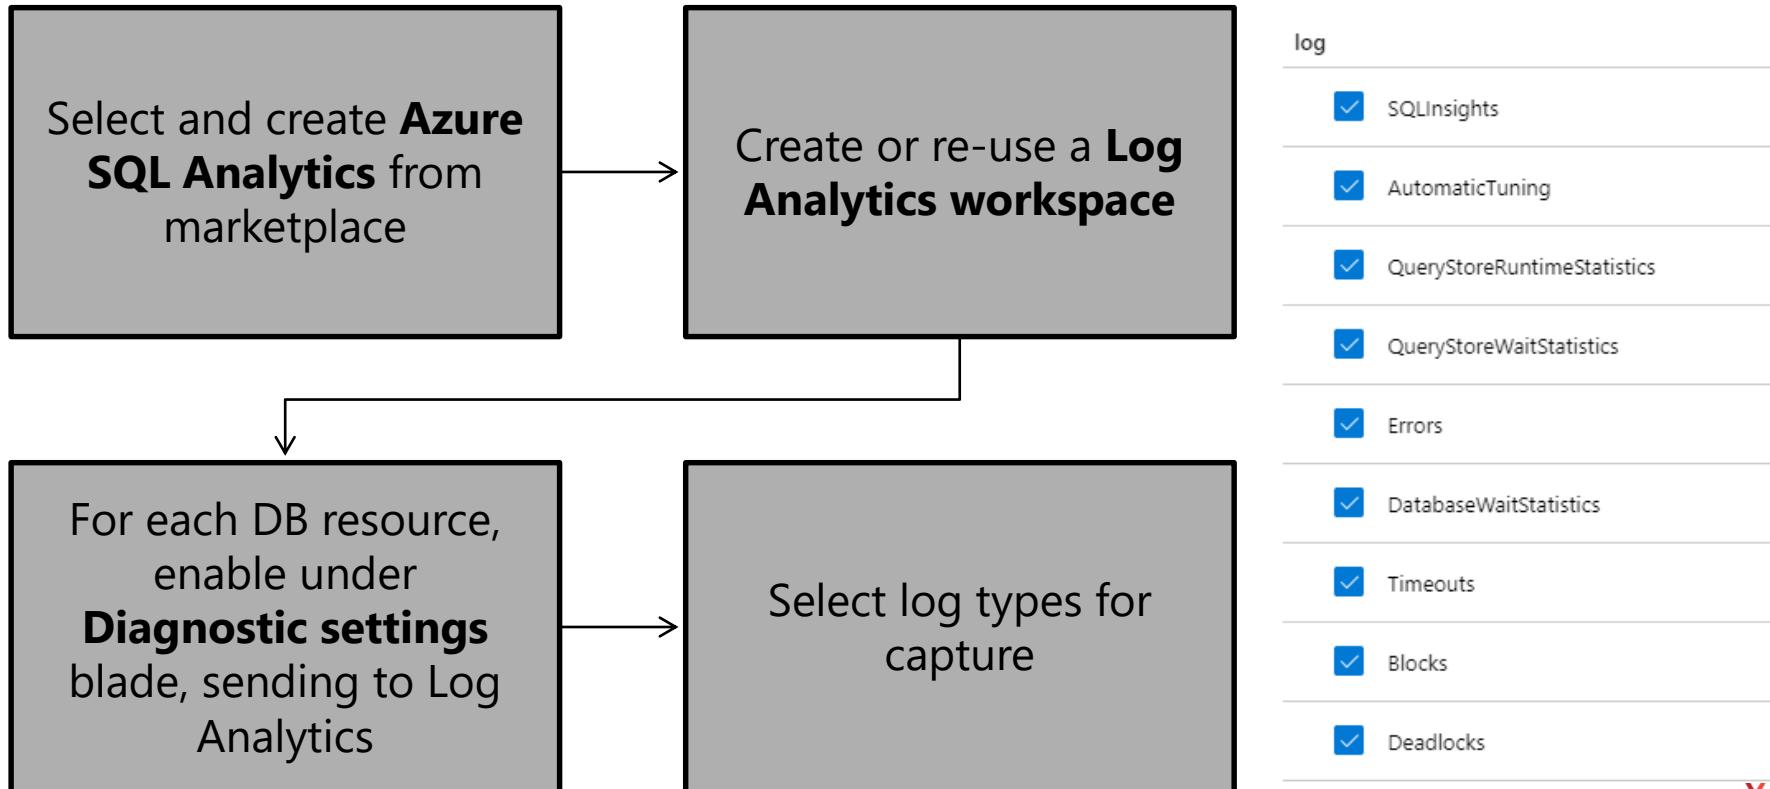
# Azure SQL Analytics

Monitor performance of Azure SQL databases, elastic pools, and Managed Instances at scale and across multiple subscriptions through a single pane of glass

- Resource usage
- Query performance
- SQL errors
- Timeouts
- Blocking
- Intelligent insights



# Enabling Azure SQL Analytics



# Azure SQL Analytics

## Azure SQL Analytics (Preview)

tpcEjosacktesting

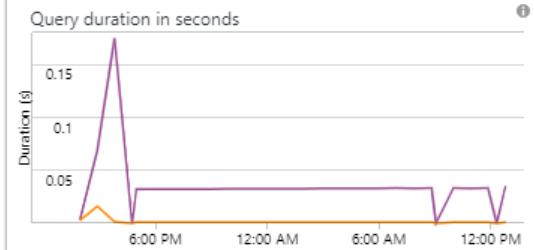
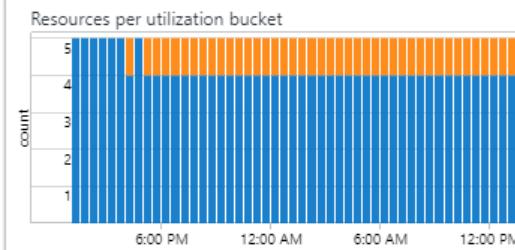
Refresh Logs

Last 24 hours

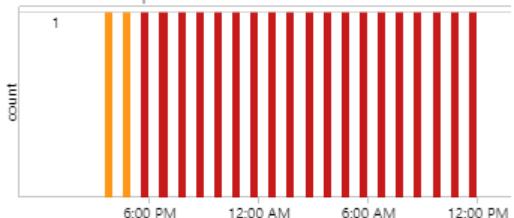
### TOP METRIC UTILIZATION PER DATABASE

DATABASE	METRIC	VALUE (%)
tpcemem_tpce	PHYSICAL MEMORY	99.3
tpce_tpce	STORAGE	74
tpce_tpce_medium	STORAGE	69
tpce_tpce_small_2019-08-08t00-00-00	STORAGE	65
tpce_wideworldimportersdw	STORAGE	1

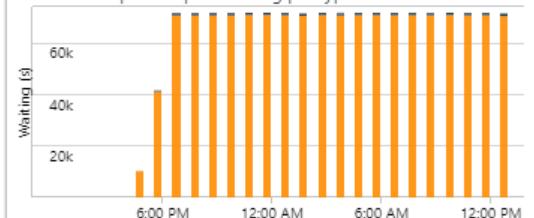
### DATABASE FLEET OVERVIEW



### Resources with performance issues



### Total time queries spent waiting per type

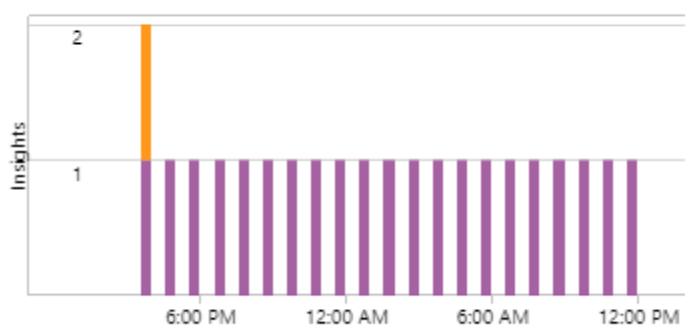


# Intelligent Insights

- Built-in intelligence that continuously monitors database usage
- Uses Artificial Intelligence to detect disruptive events that cause poor performance
  - Current hour vs. last 7 days
- Generates diagnostic log with intelligent assessment for disruptive events
- Performs root cause analysis and provides remediation where possible for performance improvements

# Intelligent Insights

## INSIGHTS SUMMARY



[View full insights report](#)

### Summary

**Start:** 2019-10-29T21:00:00Z

**Worst:** 2019-10-30T17:00:00Z

The database is hitting its IO limits. Average consumption has reached 98.7%.

**Last Update:** 2019-10-30T17:00:00Z

The database is hitting its IO limits. Average consumption has reached 98.7%.

### Detected patterns:

1. Critical SQL Errors

# Intelligent Insights patterns

Reaching  
resource limits

Workload  
increase

Memory pressure

Locking

Increased  
MAXDOP

Pagelatch  
contention

Missing Index

New Query

Increased wait  
stat

Tempdb  
contention

DTU shortage

Plan regression

Database scoped  
configuration  
change

Slow client

Pricing tier  
downgrade

# Diagnose and solve problems

Home > tpcemem > TPCE (tpcemem/TPCE) - Diagnose and solve problems

 TPCE (tpcemem/TPCE) - Diagnose and solve problems  
SQL database

Search (Ctrl+/  
Search for common problems, tools and more

Common problems

Explore the most common problems for your resource. Select Troubleshoot to run an automated troubleshooter, follow do-it-yourself troubleshooting steps, or expl

All categories ▾ No grouping ▾

 Health  
No health related issues found

Resource health dashboard | Did this resolve your issue? Yes No

 Provisioning  
**Updating the SLO is taking a long time**  
Learn more about scaling a resource and the factors that determine how long it takes to update SLO.  
**Troubleshoot**

 Performance  
**The workload for my database has increased**  
Learn more about resource limits and how to troubleshoot corresponding performance issues.  
**Troubleshoot**



# Automated Troubleshooter

Our automated troubleshooter will diagnose Performance issues for TPCE. You can also choose to manually troubleshoot your issue using our recommended steps.

## ^ Automated troubleshooter



### Troubleshooting "The workload for my database has increased" for TPCE

Our internal service telemetry detected high IO usage greater than 90% for more than 5 consecutive minutes on your Azure SQL DB database **TPCE** in server **tpcemem**.

#### Recommended Steps

This indicates possible application slowness, timeouts due to lack of IO resources to execute the requested workload during that specific period. Please see the recommended articles for more details.

#### Recommended Documents

- Consider batching to reduce IO consumption when doing large operations ↗
- Increase Azure Database tier to get more DTUs and hence accommodate for more IO in your environment ↗

[Rerun troubleshooter](#) | [Download report](#) | Did this resolve your issue? [Yes](#) [No](#)

# Automatic Tuning

One-click to enable | Prevent and mitigate issues | No app changes needed

## Tuning actions

Create missing indexes

Drop unused/duplicate indexes

Force last good plan

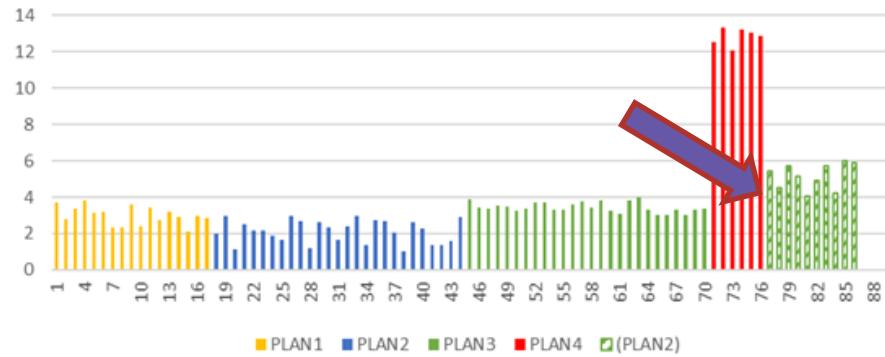
Configure the automatic tuning options 

Option	Desired state	Current state
 FORCE PLAN	<input checked="" type="button"/> ON <input type="button"/> OFF <input type="button"/> INHERIT	<b>ON</b> Forced by user
 CREATE INDEX	<input checked="" type="button"/> ON <input type="button"/> OFF <input type="button"/> INHERIT	<b>ON</b> Forced by user
 DROP INDEX	<input checked="" type="button"/> ON <input type="button"/> OFF <input type="button"/> INHERIT	<b>ON</b> Forced by user

# Plan Regression Correction

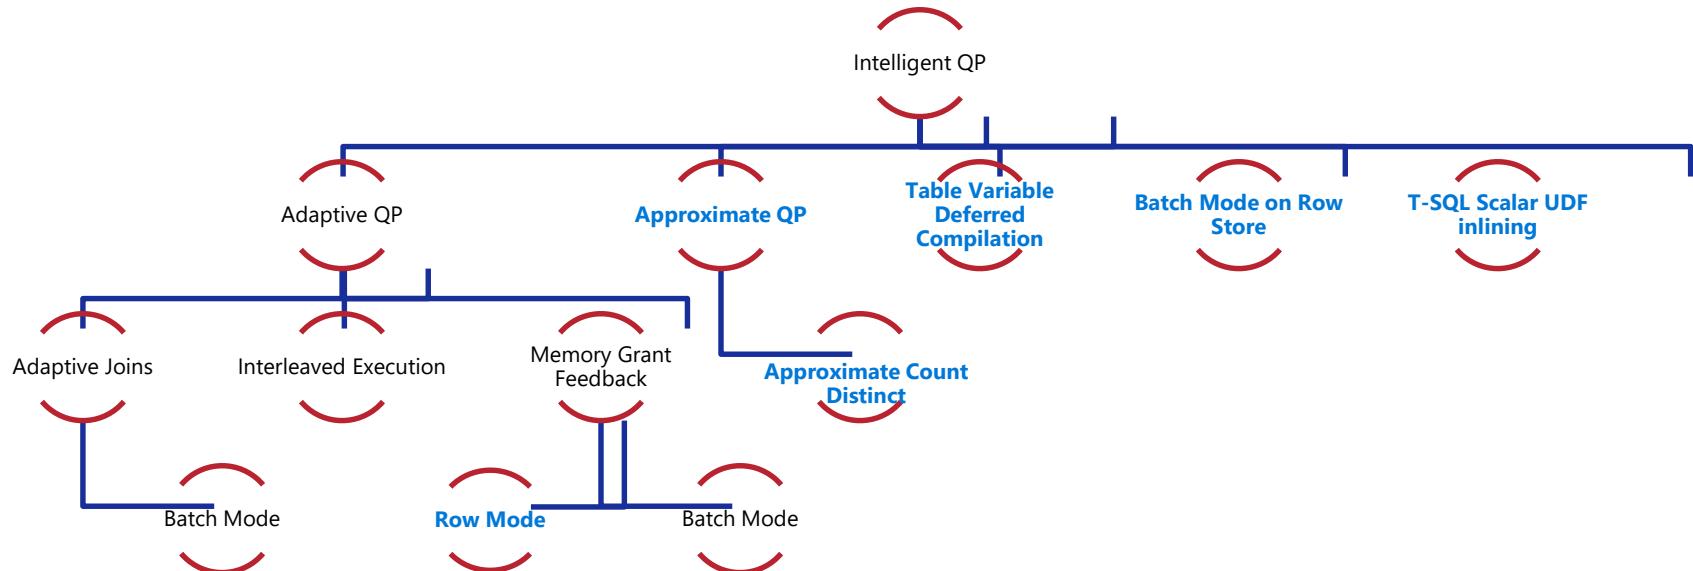
We can now detect and correct these scenarios without manual intervention  
Recommended actions surfaced via sys.dm\_db\_tuning\_recommendations  
We can now automatically switch to the last known good plan whenever the regression is detected

```
ALTER DATABASE CURRENT  
SET AUTOMATIC_TUNING (FORCE_LAST_GOOD_PLAN = ON);
```



# Intelligent Query Processing

Query processing features that automatically improve workload performance



Azure SQL Database

SQL Server 2017 **SQL Server 2019**  
PASS

# CLICK TO EDIT MASTER TITLE STYLE

Since SQL Server 2012 – we've bound these two features together

## Columnstore indexes

IO

Access only the data in columns that you need

Effective compression over traditional rowstore

## Batch Mode

CP  
U

Allows query operators to process data more efficiently by working on a batch of rows at a time

Built for analytical workload scale

# CLICK TO EDIT MASTER TITLE STYLE

- Click to edit Master text styles
  - Second level
    - Third level
      - Fourth level
        - Fifth level

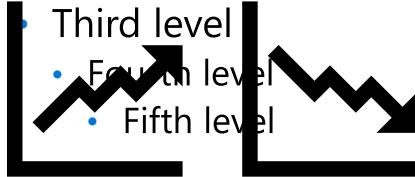
# When approximate is good enough...

Approximate Query Processing

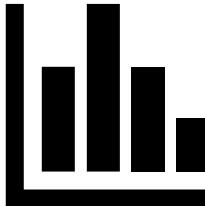
- Click to edit Master text styles

- Second level

- Third level
  - Fourth level
  - Fifth level



Dashboard scenarios against big data sets with many distinct values (for example, distinct orders counts over a time period) – and many concurrent users



Data science big data set exploration. Need to understand data distributions quickly and exact values are not paramount

# CLICK TO EDIT MASTER TITLE STYLE

- Click to edit Master text styles
  - Second level
    - Third level
      - Fourth level
        - Fifth level

DEMO

# Intelligent QP





# Q&A

